

Index

- A/B testing, 86
- Aaron Schwartz case, 68
- AB testing, 137
- academic data, 66
- accuracy, 215, 228
- activation function, 380
- AdaBoost, 364
- add-one discounting, 357
- agglomerative cluster trees, 338
- aggregation mechanisms, 83
- Akaike information criterion, 289, 335
- algorithm analysis, 397
- Amazon Turk, 67, 84
 - tasks assigned, 85
 - Turkers, 84
- American basketball players, 97
- analogies, 312
- anchoring, 82
- angular distance, 310
- Anscombe's Quartet, 159
- AOL, 64
- API, 65
- Apple iPhone sales, 34
- application program interfaces, 65
- area under the ROC curve, 219
- Aristotle, 326, 327
- Arrow's impossibility theorem, 84, 114
- artifacts, 69
- Ascombe quartet, 272
- asking interesting questions, 4
- associativity, 244
- autocorrelation, 46
- average link, 340
- Babbage, Charles, 57, 90
- backpropagation, 382
- Bacon, Kevin, 9
- bag of words, 14
- bagging, 362
- balanced training classes, 295
- bar charts, 179
 - best practices, 181
 - stacked, 181
- Barzun, Jacques, 5
- baseball encyclopedia, 5
- baseline models, 210
 - for classification, 210
 - for value prediction, 212
- Bayes' theorem, 150, 205, 299
- Baysian information criteria, 289
- bell-shaped, 123, 141
 - distribution, 101
- bias, 202, 417
 - lexicographic, 405
 - numerical, 405
 - temporal, 405
- bias–variance trade-off, 202
- big data, 391
 - algorithms, 397
 - bad data, 392
 - statistics, 392
- big data engineer, 4
- big oh analysis, 397
- binary relations, 320
- binary search, 398
- binomial distribution, 123

- bioinformatician, 50
- Blumenstock, Josh, 27
- Body Mass Index, 96, 177
- Bonferroni correction, 141
- boosting, 363, 364
 - algorithms, 364
- bootstrapping, 374
- Borda's method, 108
- box plots, 175
- Box, George, 201
- box-and-whisker plots, 176
- bubble plots, 179
- bupkis, 391
- Bush, George W., 326
- C-language, 59
- cache memory, 401
- canonical representation, 400
- canonization, 400
- CAPTCHAs, 89
- Carroll, Lewis, 423
- cartograms, 189
- Center for Disease Control, 204
- center of mass, 332
- centrality measures, 34
- centroids, 331
- character code unification, 74
- characteristic equation, 256
- characterizing distributions, 39
- chart types, 170
 - bar charts, 179
 - data maps, 187
 - dot and line plots, 174
 - histograms, 183
 - pie charts, 179
 - scatter plots, 177
 - tabular data, 170
- cicada, 46
- classification, 16, 210, 289
 - binary, 290, 314
 - multi-class, 297
 - regression, 290
- classification and regression trees, 357
- classifiers
 - balanced, 217
 - evaluating, 213
 - one-vs.-all, 298
 - perfect, 218
- Clinton, Bill, 326
- Clinton, Bill], 326
- closest pair of points, 398
- cloud computing services, 410
- cluster
 - conductance, 343
 - distance, 337
- clustering, 327, 373
 - agglomerative, 336
 - applications, 328
 - biological, 337
 - cut-based, 341
 - k-means, 330
 - single link, 339
 - visualization of, 337
- clusters
 - number of, 333
 - organization of, 337
- Clyde, 111, 112
- coefficient vector, 270
- Cohen's d, 136
- collaborative filtering, 9
- communication, 408, 416
- commutativity, 243
- company data, 64
- computer scientist, 2
- conditional probability, 30, 31
- Condorcet jury theorem, 84
- confusion matrix, 214, 220
- connected components, 324
- contingency table, 214
- convex, 280
- convex hull, 368
- coordination, 408
- correlation
 - analysis, 40
 - interpretation, 43
 - significance, 45
- correlation and causation, 45, 135
- cosine similarity, 309
- cross validation, 227
 - advantages, 227
- CrowdFlower, 67, 84, 86

crowdsourcing, 67, 80
 bad uses, 87
 crowdsourcing services, 84
 cryptographic hashing, 400
 CSV files, 62
 cumulative density function, 33,
 132, 186
 currency conversion, 75
 cut, 343

 damping factor, 325
 Darwin, Charles, 81
 data, 237, 391
 quantitative vs. categorical, 15
 big vs. little, 15
 cleaning, 69
 collecting, 64
 compatibility, 72
 errors, 69
 for evaluation, 225
 for testing, 225
 for training, 225
 logging, 68
 properties, 14
 scraping, 67
 structured, 14
 types, 14
 unstructured, 14
 visualizing, 155
 data analysis, 404
 data centrism, 2
 data cleaning, 376
 data errors, 417
 data formats, 61
 data hygiene
 evaluation, 225
 data munging, 57
 data parallelism, 409
 data partition, 404
 data processing, 10
 data reduction, 329
 data science, 1
 languages, 57
 models, 210
 data science television, 17
 data scientist, 2

data sources, 64
 data visualization, 155
 data-driven, 156
 models, 207
 de Méré, Chevalier, 29
 decision boundaries, 291
 decision tree classifiers, 299, 357
 decision trees, 357
 advantages, 358
 construction, 359
 ensembles of, 362
 deep learning, 202, 352, 377
 models, 209
 network, 378
 DeepWalk, 385
 degree of vertex, 325
 depth, 378
 derivative, 281
 partial, 282
 second, 281
 descriptive statistics, 34
 deterministic sampling algorithms,
 404
 developing scoring systems, 99
 dictionary maintenance, 399
 DiMaggio, 148
 dimension reduction, 277, 376
 dimensional egalitarianism, 308
 dimensions, 384
 dinosaur vertebra, 78
 directed acyclic graph, 110
 directed graph, 321
 discounting, 209, 356
 disk storage, 401
 distance methods, 303
 distance metrics, 304
 L_k , 305
 euclidean, 305
 manhattan distance, 305
 maximum component, 305
 distances, 319
 measuring, 303
 distributed file system, 396
 distributed processing, 407
 divide and conquer, 411
 DNA sequences, 402

dot product, 241
 duality, 268
 duplicate removal, 400

 E-step, 335
 edge cuts, 324
 edges, 319
 effect size, 136
 eigenvalues, 255

- computation, 256
- decomposition, 257
- properties of, 255

 Elizabeth II, 326
 Elo rankings, 104
 embedded graph, 323
 embedding, 321
 Emoji Dick, 86
 Engels, Friedrich, 391
 ensemble learning, 363
 entropy, 310
 equations

- determined, 251
- underdetermined, 251

 error, 202, 221

- absolute, 221
- detection, 155
- mean squared, 223, 331
- relative, 222
- residual, 269
- root mean squared, 223
- squared, 222
- statistics, 221

 errors vs. artifacts, 69
 ethical implications, 416
 Euclidean metric, 303
 evaluation

- environments, 224
- statistics, 214

 event, 28
 Excel, 59
 exclusive or, 361
 exercises, 23, 53, 90, 119, 151, 199, 234, 263, 301, 346, 388, 419
 expectation maximization, 335
 expected value, 28

experiment, 27
 exploratory data analysis, 155, 156

- visualization, 160

 F-score, 216
 Facebook, 21
 false negatives, 214
 false positives, 214
 fast Fourier transform, 47
 fault tolerance, 408
 feature engineering, 375
 feature scaling, 274

- sublinear, 275
- z-scores, 275

 features

- highly-correlated, 277

 Fermat, Pierre de, 30
 Feynman, Richard, 229
 FFT, 47
 filtering, 403
 financial market, 126
 financial unification, 75
 fit and complexity, 288
 FoldIt, 89
 football, 111

- American players, 97
- game prediction, 111

 forecasting

- time series, 212

 formulation, 354
 Freedom of Information Act, 12, 65
 frequency counting, 400
 frequency distributions, 184
 furthest link, 340

 Galton, Francis, 81
 games with a purpose, 88
 gamification, 88
 garbage in, garbage out, 3, 69
 Gates, Bill, 130
 Gaussian

- elimination, 250

 Gaussian distribution, 124
 Gaussian noise, 125
 General Sentiment, 20
 genius, 19

- geometric mean, 35
- geometric point sets, 238
- geometry, 240
- Gini impurity, 360
- Global Positioning System, 12
- gold standards, 99
- good scoring functions, 101
- Goodhart's law, 303
- Goodhart, Charles, 303
- Google
 - AlphaGo, 372
 - TensorFlow, 377
- Google Flu Trends, 394
- Google News, 219
- Google Ngrams, 10
- Google Scholar, 66
- government data, 65
- gradient boosted decision trees,
 - 359, 366
- gradient descent search, 281
- graph embeddings, 384
- graph theory, 323
- graphs, 238, 319, 321
 - cuts, 342
 - dense, 322
 - directed, 321
 - embedded, 323
 - labeled, 323
 - non-simple, 322
 - simple, 322
 - sparse, 322
 - topological, 323
 - undirected, 321
 - unlabeled, 323
 - unweighted, 322
 - weighted, 320, 322
- Gray, Dorian, 251
- grid indexes, 315
- grid search, 409
- Hadoop distributed file system, 414
- Hamming, Richard W., 1
- hash functions, 399
 - applications, 399
- hashing, 399
- heatmaps, 178
- hedge fund, 1
- hierarchy, 298
- higher dimensions, 307, 370
- histograms, 32, 183, 222
 - best practices, 186
 - bin size, 184
- Hitler, Adolf, 326
- HTML, 67
- hypothesis development, 328
- hypothesis driven, 156, 392
- Imagenet, 379
- IMDb, 7
- imputation
 - by interpolation, 78
 - by mean value, 77
 - by nearest neighbor, 78
 - by random value, 77
 - heuristic-based, 77
- independence, 30, 123, 354
- inflation rates, 76
- information gain, 360
- information theoretic entropy, 360
- infrastructure, 396
- inner product, 243
- institutional review board, 88
- Internet Movie Database, 7
- Internet of Things, 68
- inverse transform sampling, 132
- IPython, 61
- IQ testing, 89
- Jaccard distance, 341
- Jaccard similarity, 341
- Jackson, Michael [136], 262
- Java, 59
- Jesus, 326
- JSON, 63
- k-means clustering, 343, 409
- k-medoids algorithm, 332
- k-nearest neighbors, 313
- Kaggle, viii
- kd-trees, 316
- kernels, 371
- Kolmogorov-Smirnov test, 139

Kruskal's algorithm, 339
 Kullback-Leibler divergence, 311
 labeled graphs, 323
 Lang, Andrew, 267
 Laplace, 356
 Laplacian, 343
 large-scale question, 9
 latent Dirichlet allocation, 373
 learning rate, 283, 383
 learning to rank, 119
 least squares regression, 270
 lie factor, 164
 Lincoln, Abraham, 242
 line charts, 174

- advantages, 174
- best practices, 175

 line hatchings, 177
 linear algebra, 237

- power of, 237

 linear algebraic formulae

- interpretation, 238

 linear equation, 238
 linear programming, 369
 linear regression, 212, 267

- error, 269
- solving, 270

 linear support vector machines, 369
 linear systems, 250
 Linnaeus, Carl, 326
 live data, 395
 locality, 401
 locality sensitive hashing, 317
 logarithm, 47
 logistic classification

- issues, 295

 logistic function, 381
 logistic regression, 366
 logit, 381
 logit function, 106, 292
 loss function, 280, 294
 LU decomposition, 254
 lumps, 329
 M-step, 335
 machine learning, 351

algorithms, 375
 classifiers, 85
 models, 208
 main memory, 401
 major league baseball, 6
 MapReduce, 407, 410

- programming, 412

 MapReduce runtime system, 415
 matchings, 324
 Mathematica, 59, 61
 Matlab, 58
 matrix, 14, 237, 270, 373

- addition, 242
- adjacency, 246, 320
- covariance, 245, 257, 271
- determinant, 249, 254
- eigenvalues, 255
- eigenvectors, 255
- factoring, 252
- identity, 246, 248
- inversion, 248
- linear combinations of, 242
- multiplication, 243
- multiplicative inverse of, 249
- non-singular, 249
- permutation, 247
- rank, 251
- reasons for factoring, 252
- rotation, 248
- singular, 249
- transpose of, 242
- triangular, 254
- underdetermined, 256

 matrix multiplication, 243, 398

- applications, 244

 matrix operations

- visualizing, 241

 maximum margin separator, 367
 mean, 34, 83, 125, 132, 138, 212, 227, 403

- arithmetic, 35
- geometric, 35

 measurement error, 125
 median, 35, 83, 132, 212, 403
 mergesort, 398
 metadata, 7

method centrism, 2
 metric, 304
 identity, 304
 positivity, 304
 symmetry, 304
 triangle inequality, 304
 minima
 global, 284
 local, 284
 minimum spanning tree, 324, 339
 missing values, 76, 376
 mixture model, 333
 Moby Dick, 86
 mode, 36
 model-driven, 417
 modeling, 201, 328, 416
 philosophies of, 201
 principles for effectiveness, 203
 models
 ad hoc, 208
 baseline, 210
 blackbox, 206
 data science, 210
 data-driven, 207
 deep learning, 209
 descriptive, 206
 deterministic, 208
 evaluating, 212
 first-principle, 207
 flat, 209
 Google's forecasting, 204
 hierarchical, 209
 linear, 206
 live, 204
 machine learning, 208
 neural network, 206
 non-linear, 206
 overfit, 203
 simplifying, 286
 simulation, 229
 stochastic, 208
 taxonomy of, 205
 underfit, 202
 Moneyball, 5
 monkey, 215
 monotonic, 307

Monte Carlo
 sampling, 134
 simulations, 229
 Mosteller, Frederick, 121
 multiclass systems
 evaluating, 219
 multiedge, 322
 multinomial regression, 299
 multiplying probabilities, 48

 naive Bayes, 354, 363
 name unification, 74
 Napoleon, 326
 NASA, 73
 National Football League, 112
 natural language processing, 20
 nearest centroid, 340
 nearest neighbor, 339, 397
 nearest neighbor classification, 311
 advantages, 311
 nearest neighbors
 finding, 315
 negative class, 213
 Netflix prize, 9
 network
 depth, 379
 network methods, 303
 networks, 109, 238, 319, 378
 induced, 320
 learning, 379
 neural networks, 377
 new data set, 156
 New York, 277
 Nixon, Richard, 326
 NLP-based system, 21
 no free lunch theorem, 353
 node
 bias of, 381
 non-linear classifiers, 366
 non-linear functions
 fitting, 273
 non-linear support vector
 machines, 369
 non-linearity, 358, 377, 380
 norm, 287
 normal, 125

- normal distribution, 79, 109, 124, 141
 - implications, 126
- normality testing, 141
- normalization, 103, 376
- normalizing skewed distribution, 49
- norms, 309
- NoSQL databases, 415
- notebook environments, 59
- numerical conversions, 73
- Obama, Barack, 326
- Obama, Barack [91], 327
- Obama, Barack], 326
- Occam's razor, 201, 211, 286
- Occam, William of, 202
- Oh G-d, 177
- optimization
 - local, 284
- Oracle of Bacon, 9
- outlier, 118
 - detection, 78
- outlier detection, 329
- outliers
 - removing, 272
- overfitting, 202, 296
- overlap percentage, 137
- ownership, 417
- p-values, 145
- packing data, 402
- PageRank, 100, 111, 325
- pairwise correlations, 158
- parallel processing, 407
- parallelism, 406
- parameter fitting, 279
- parameter spaces
 - convex, 280
- partition function, 299
- Pascal's triangle, 123
- Pascal, Blaise, 30
- paths, 246
- Pearson correlation coefficient, 41, 136
- penalty function, 293
- performance of models, 39
- periodic table, 188
- Perl, 58
- permutation, 246
 - randomly generating, 147
 - tests, 145
- personal wealth, 130
- pie charts, 179
 - bad examples, 183
 - best practices, 181
- point spread, 112
- points
 - rotating, 248
- points vs. vectors, 309
- Poisson distribution, 127
- position evaluation function, 372
- positive class, 213
- power law distribution, 129
- power law function, 276
- precision, 3, 215, 221
- prefetching, 401
- principal components, 260
 - analysis, 260
- prior probability, 354
- privacy, 418
- probabilistic, 203
- probability, 27, 29, 354
- probability density function, 32, 132, 186
- probability distribution, 32
- probability of an event, 28
- probability of an outcome, 28
- probability vs. statistics, 29
- program flow graph, 322
- programming languages, 57
- protocol buffers, 63
- proxies, 99
- psychologists, 89
- Pubmed, 70
- pure partition, 360
- Pythagorean theorem, 306
- Python, 58, 67
- Quant Shop, 17
- R, 58
- Rand index, 341

- random access machine, 397
- random sampling, 403, 406
- random variable, 28
- ranking systems
 - class rank, 100
 - search results, 100
 - top sports teams, 100
 - university rankings, 100
- rankings, 95
 - digraph-based, 109
 - historical, 117
 - merging, 108
 - techniques, 104
- ratio, 48
- ray
 - unit, 241
- Reagan, Ronald, 326
- rearrangement operations, 238
- recall, 216, 221
- receiver-operator characteristic
 - curve, 218
- rectified linear units, 381
- rectifier, 381
- redundancy, 393
- regression, 16
 - application, 359
 - for classification, 290
 - LASSO, 287
 - logistic, 289, 292
 - ridge, 286
- regression models, 272
 - removing outliers, 272
- regularization, 286, 376
- reinforcement learning, 372
- Richter scale, 131
- right-sizing training data, 404
- road network, 322
- robustness, 3, 96, 359
- Roosevelt, Franklin D., 326
- Rota, Gian-Carlo, 237
- sabermetrics, 22
- sample space, 27
- sampling, 132, 403, 404
 - beyond one dimension, 133
 - by truncation, 404
- scalar
 - multiplication, 242
- scale invariant, 132
- scales
 - Likert, 298
 - ordinal, 297
- scaling constant, 293
- scatter plots, 98, 177
 - best practices, 177
 - three-dimensional, 179
- Schaumann, Jan, 351
- scientist, 2
- scores, 95
 - threshold, 218
- scores vs. rankings, 100
- scoring functions, 95
- security, 418
- self-loop, 322
- semi-supervised learning, 374
- Shakespeare, William, 326
- sharp, 215
- Sheep Market, 86
- shortest paths, 324
- signal to noise ratio, 37
- significance level, 139
- Silver, Nate, 203
- similarity graphs, 341
- similarity matrix, 342
- simple graph, 322
- single-command program, 224
- single-pass algorithm, 402
- singular value decomposition, 258
- sketching, 403
- skew, 413
- Skiena, Len, ix
- small evaluation set, 226
- social media, 392
 - analysis, 21
 - data, 394
- Social Network–movie, 105
- spam, 393
- spam filtering, 393
- Spearman rank correlation
 - coefficient, 42, 108
- Spears, Britney [566], 262
- spectral clustering, 343

spidering, 67
 splitters, 329
 sports performance, 38
 SQL databases, 63
 standard deviation, 36, 125, 132,
 138, 227, 403
 statistical analysis, 121, 230
 statistical distributions, 122
 statistical proxy, 7
 statistical significance, 135
 statistics, 29
 stochastic gradient descent, 285,
 296, 382
 stock market, 37, 79, 126
 Stony Brook University, 20
 stop words, 414
 storage hierarchy, 401
 streaming, 402
 summary statistics, 157, 159
 supervised learning, 372
 supervision, 372
 degrees of, 372
 support vector machines, 352, 366
 support vectors, 368
 Surowiecki, James, 82
 sweat equity, 66

 T-test, 137
 tangent line, 282
 target scaling, 274
 sublinear, 276
 taxi
 records from New York, 11
 tipping model, 286
 tipping rate, 13, 277
 taxi driver, 277
 terms of service, 68
 test statistic, 138
 text analysis, 253
 theory of relativity, 144
 Tikhonov regularization, 287
 time unification, 75
 Titanic, 183, 358
 top-k success rate, 219
 topic modeling, 373
 topological graph, 323

topological sorting, 110, 324
 transparency, 417
 tree, 298, 336
 trees
 agglomerative, 337
 true negatives, 214
 true positives, 214
 truncation, 405
 Tufte, Edward, 155, 162
 Twitter, 392, 404, 405

 U.S. presidential elections, 187, 203
 uncertainty, 175
 undirected graph, 321
 uniform distribution, 406
 uniform sampling, 405
 uninvertible, 400
 unit conversions, 72
 UNIX time, 75
 unlabeled graphs, 323
 unrepresentative participation, 393
 unsupervised learning, 372
 unweighted graph, 322
 urban transportation network, 11

 validation data, 95
 value prediction, 210
 value prediction models
 evaluating, 221
 variability measures, 36
 variance, 36, 202, 403
 interpretation, 37
 variation coefficient, 137
 variety, 394
 vectors, 240
 unit, 240
 velocity, 394
 veracity, 395
 vertex, 319
 vertices, 319
 visualization, 404
 chart types, 170
 critiquing, 189
 interactive, 195
 tools, 160
 visualization aesthetic, 162

chartjunk, 162, 165
colors, 163, 168
data-ink ratio, 162, 163
lie factor, 162, 164
repetition, 163, 169
scaling and labeling, 162, 167
volume, 394
Voronoi diagrams, 315
voting, 363
 with classifiers, 363
web crawling, 68
weighted average, 84

weighted graph, 322
Welch's t-statistic, 138
Wikipedia, 20, 79, 116, 326
wisdom, 19
wisdom of crowds, 81
Wolfram Alpha, 59, 144
word embeddings, 254, 383
word2vec, 384
XML, 62
Z-score, 103, 308, 376
Zipf's law, 131, 357