

Undergraduate Topics in Computer Science

Bit Message

Contents

Undergraduate Topiss in soniniter cienteo 0°0°0°0°0° B.

o
o
o

0.1 to Assembly

Language

0.2 A Concise Introduction

Second Edition

Ian Mackie, University of Sussex, Brighton, UK

0.3 Advisory Editors

Samson Abramsky © Department of Computer Science, University of Oxford, Oxford, UK

Chris Hankin (®), Department of Computing, Imperial College London, London, UK Mike Hinchey (1), Lero - The Irish Software Research Centre, University of Limerick, Limerick, Ireland

Dexter C. Kozen, Department of Computer Science, Cornell University, Ithaca, NY, USA

Andrew Pitts © , Department of Computer Science and Technology, University of Cambridge, Cambridge, UK

Hanne Riis Nielson (®), Department of Applied Mathematics and Computer Science, Technical University of Denmark, Kongens Lyngby, Denmark

Steven S. Skiena, Department of Computer Science, Stony Brook University, Stony Brook, NY, USA

Iain Stewart © , Department of Computer Science, Durham University, Durham, UK

'Undergraduate Topics in Computer Science' (UTiCS) delivers high-quality instructional content for undergraduates studying in all areas of computing and information science. From core foundational and theoretical material to final-year topics and applications, UTiCS books take a fresh, concise, and modern approach and are ideal for self-study or for a one- or two-semester course. The texts are all authored by established experts in their fields, reviewed by an international advisory board, and contain numerous examples and problems, many of which include fully worked solutions.

The UTiCS concept relies on high-quality, concise books in softback format, and generally a maximum of 275-300 pages. For undergraduate textbooks that are likely to be longer, more expository, Springer continues to offer the highly regarded Texts in Computer Science series, to which we refer potential authors.

More information about this series at <http://www.springer.com/series/7592>

James T. Streib

0.4 Guide to Assembly Language

A Concise Introduction

Second Edition

Springer

ISSN 1863-7310

ISSN 2197-1781 (electronic)

Undergraduate Topics in Computer Science

ISBN 978-3-030-35638-5

ISBN 978-3-030-35639-2 (eBook)

<https://doi.org/10.1007/978-3-030-35639-2>

1st edition: © Springer-Verlag London Limited 2011

2nd edition: © Springer Nature Switzerland AG 2020

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed.

The use of general descriptive names, registered names, trademarks, service marks, etc. in this publication does not imply, even in the absence of a specific statement, that such names are exempt from the relevant protective laws and regulations and therefore free for general use.

The publisher, the authors and the editors are safe to assume that the advice and information in this book are believed to be true and accurate at the date of publication. Neither the publisher nor the authors or the editors give a warranty, expressed or implied, with respect to the material contained herein or for any errors or omissions that may have been made. The publisher remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

This Springer imprint is published by the registered company Springer Nature Switzerland AG
The registered company address is: Gewerbestrasse 11, 6330 Cham, Switzerland

0.5 Purpose

The purpose of this text is to assist one in learning how to program in assembly language on an Intel processor using Microsoft Assembler (MASM) in a minimal amount of time. In addition, through programming the reader learns more about the computer architecture of the Intel processor and also the relationship between high-level languages and low-level languages.

0.6 Need

In the past, many departments have had two separate courses: one in assembly language programming (sometimes called computer systems) and a second course in computer organization and architecture. With today's crowded curriculums, there is sometimes just one course in the computer science curriculum in computer organization and architecture, where various aspects of both courses are included in the one course. The result might be that unfortunately there is not enough coverage concerning assembly language programming.

0.7 Importance of Assembly Language

Although the need for assembly language programmers has decreased, the need to understand assembly language has not, and the reasons why one ought to learn to program in assembly language include the following:

- Sometimes just reading about assembly language is not enough, and one must actually write assembly language code to understand it thoroughly (although the code does not have to be extremely complicated or tricky to gain this benefit).
- Although some high-level languages include low-level features, there are times when programming in assembly language can be more efficient in terms of both speed and memory.
- Programming in assembly language has the same benefits as programming in machine language, except it is easier. Further one can gain some first-hand

knowledge into the nature of computer systems, organization, and architecture from a software perspective.

Having knowledge of low-level programming concepts helps one understand how high-level languages are implemented and various related compiler concepts.

0.8 Comparison to Computer Organization and Other Assembly Language Textbooks

Many textbooks on computer organization have only a few sections or chapters dealing with assembly language, and as a result, they might not cover the aspects of assembly language thoroughly enough. Also, instead of discussing a real assembly language, they might just use a hypothetical assembly and machine language. Although this can be helpful in understanding some of the basic concepts, the student might neither see the relevance nor appreciate many of the important concepts of a real assembly language.

On the other hand, there are a number of assembly language texts that go into significant detail which can easily fill an entire semester and almost warrant a two-semester sequence. Unfortunately, some of the more comprehensive assembly language texts might not be the best choice for learning to program in assembly language due to the same reasons that make them excellent comprehensive texts.

This current text does not attempt to fill the needs of either of these two previous varieties of texts because it falls between the scopes of these two types of texts. The purpose of this text is to provide a concise introduction to the fundamentals of assembly language programming, and as a result, it can serve well as either a stand-alone text or a companion text to the current popular computer organization texts.

0.9 Features of This Text

The primary goal of this text is to get the student programming in assembly language as quickly as possible. Some of these features that make this possible include simplified register usage, simplified input/output using C-like statements, and the use of high-level control structures. All of these features help the reader begin programming quickly and reinforce many of the concepts learned in previous computer science courses. Also, many of the control structures are implemented without the use of high-level structures to allow readers to understand how they are actually implemented. Further, many of the assembly language code segments are preceded by C program code segments to help students see the relationships between high-level and low-level languages. Other notable features at the end of each chapter include the following:

- One or more complete programs illustrating many of the concepts introduced in that chapter.

- Chapter summaries, which by themselves do not substitute for reading a chapter, but after reading a chapter they serve as a nice review of important concepts for students preparing for a quiz or exam.
- Exercises composed of a variety of questions, from short answer to programming assignments. Items marked with an * have solutions in Appendix D.

0.10 Features New to the Second Edition

The second edition retains all of the features of the first edition. Of course, any known errors have been corrected and areas that could be clarified have been reworded. Features new to the second edition include the following:

- Many illustrations have had color added to them to help readers visualize important concepts.
- A new section has been added to Appendix B on floating-point number representation.
- A new chapter has been added on floating-point number processing.
- 32-bit processing has been retained throughout the text, and a new chapter has been added on 64-bit processing.
- A new section has been added to Chap. 12 on floating-point and 64-bit machine language.
- Additional exercises have been added to various chapters.

0.11 Brief Overview of the Chapters and Appendices

If this text is used in conjunction with another text in a computer organization course, then there is a potential for some duplication between the texts. For example, many texts in assembly language begin with an introduction to binary arithmetic, which of course is incredibly important in a low-level language. However, should this text be used in conjunction with a computer organization text, then many of those concepts will have already been introduced. As a result, this text begins at the outset to get students into programming quickly and introduces or reviews binary on an as-needed basis. However, should this text be used as a stand-alone text, then Appendix B introduces binary numbers, hexadecimal numbers, conversions, logic, arithmetic, and floating-point number representation, should the instructor or student wish to examine this material first. What follows is a brief overview of the chapters and the appendices:

- Chapter 1 provides an overview of assembly language and an introduction to the general-purpose registers.
- Chapter 2 introduces the reader to input/output in assembly language, specifically using the C programming language `scanf` and `printf` instructions.
- Chapter 3 explains basic arithmetic in assembly language, including addition, subtraction, multiplication, division, and operator precedence.
- Chapter 4 shows how to implement selection structures in assembly language, such as if-then, if-then-else, nested if structures, and the case (switch) structure.
- Chapter 5 continues with iteration structures, specifically the pre-test, post-test, and definite iterations loop structures, along with nested loops.
- Chapter 6 introduces the logic, shift, arithmetic shift, rotate, and stack instructions.
- Chapter 7 discusses procedures, introduces macros, and explains conditional assembly.
- Chapter 8 presents arrays, sequential searching, and the selection sort.
- Chapter 9 discusses strings, string instructions, arrays of strings, and comparisons of strings.
- Chapter 10 provides an introduction to the floating-point register stack, arithmetic, input/output, and instructions.
- Chapter 11 explains basic 64-bit input/output, storage, and processing.
- Chapter 12 introduces machine language from a discovery perspective and can serve as an introduction to some of the principles of computer organization or it might be used as a supplement to a companion computer organization text.
- Appendix A illustrates how to assemble programs using MASM.
- Appendix B provides an overview of binary and hexadecimal conversions, logic, arithmetic, and representation of floating-point numbers. The first three chapters of the text require limited use of binary and hexadecimal numbers, so one might not need to read this appendix until later in the text. However, Chap. 6 requires extensive use of binary numbers and logic. Depending on the reader's background, this appendix should be read prior to that chapter. If not covered elsewhere or it has been a while since one has studied numbering systems, this appendix can serve as a basic introduction or a good review, respectively. If one has had exposure to these topics in a previous course, concurrent course, or from another textbook in the same course, then this appendix can be skipped.

- Appendix C summarizes the assembly language instructions introduced in this text.
- Appendix D provides answers to selected exercises marked with an * that appear at the end of each chapter and at the end of Appendix B.
- The Glossary contains terms that are first introduced in italics in the text. The descriptions of terms in glossary should not be used in lieu of the complete descriptions in the text but rather they serve as a quick review and reminder of the basic meaning of various terms. Should a more complete description be needed, then index can guide the reader to the appropriate pages where the terms are discussed in more detail.

0.12 Scope

This text includes the necessary fundamentals of assembly language to allow it to be used as either a stand-alone text in a one-semester assembly language course or a companion text in a computer organization and architecture course. As with any text, decisions then must be made on what should be included, excluded, emphasized, and deemphasized. This text is no exception in that it does not include every idiosyncrasy of assembly language and thus it might not contain some of the favorite sub-topics of various instructors. Some of these might include 16-bit processing and Windows programming among others, but these, of course, can be supplemented at the instructor's discretion. However, what is gained is that readers should be able to write logically correct programs in a minimal amount of time, which is the original intent of this text. The Intel architecture is used because of its wide availability, and Microsoft Assembler (MASM) is used due to a number of high-level control structures that are available in that assembler. Note that Java is a registered trademark of Oracle and/or its affiliates, Intel and Pentium are trademarks of Intel Corporation, and Visual Studio, Visual C++, and Microsoft Assembler (MASM) are registered trademarks of Microsoft Corporation.

0.13 Audience

It is assumed that the reader of this book has completed a two-semester introductory course sequence in a high-level language such as C, C++, or Java. Although a student might be able to use this text only after a one-semester course, an additional semester of programming in a high-level language is usually preferred to allow for a better understanding of the material due to increased programming skills and knowledge of data structures, specifically stacks and linked lists.

0.14 Acknowledgements

The author wishes to acknowledge both his editor Wayne Wheeler and associate editor Simon Rees for their assistance; thank his reviewers Mark E. Bollman of Albion College, James W. Chaffee of the University of Iowa, and Takako Soma of Illinois College for their suggestions; and offer personal thanks to his wife Kimberly A. Streib and son Daniel M. Streib for their patience. On a special note, this second edition is written in remembrance of Curt M. White who served as a reviewer of the first edition.

0.15 Feedback

As with any work, the possibility of errors exists. Any comments, corrections, or suggestions are welcome and should be sent to the e-mail address listed below. In addition to copies of the complete programs at the end of each chapter, any significant corrections can also be found at the Web site listed below. Website: <http://www.jtstreib.com/GuideAssemblyLanguage.html>

Jacksonville, Illinois, USA

James T. Streib

September 2019

e-mail: james.streib@jtstreib.com

Contents

Chapter 1

Iteration Structures

As should be recalled from previous courses, there are many different types of iteration structures available to a programmer in a high-level programming language. Just as there are many structures in a high-level language, there are corresponding structures in assembly language, such as the pre-test, post-test, and fixed-iteration loop structures. Depending on the circumstances, one should use the best structure for the task at hand.

1.1 Pre-test Loop Structure

Probably the most versatile loop is the count-controlled pre-test while loop, where any number of tasks can be performed in the body of the loop. The basic structure of this loop can be found below in the C code segment:

```
i=1;
while(i<=3) {
    // body of loop
    i++;
}
```

The three parts of any loop are "initialization," "test," and "change." In the segment above, i is known as the loop control variable (LCV), where it is initialized to 1, it is then tested, and loops when i is less than or equal to 3, and it changes when it is incremented by 1. As with the if-then structure, MASM has directives that simplify the implementation of the while loop structure. The directives are the `.while` and `.endw` directives as shown below:

```
mov i,1
_while i<=3
; body of loop
inc i
_endw
```

The `.while` directive has the same limitations as the `.if` directive, where a comparison cannot be made between two memory locations. Also, unless addressing elements of an array or a string, the use of arithmetic expressions should be avoided. Further, regardless of the number of statements in the body of the loop, the structure must end with the `.endw` directive. Lastly, a register could be used in place of the variable i to help increase speed, but as will be seen shortly, there is another loop that could be used if speed is a concern.

As with the `if` structure, the `while` structure can be implemented using a compare statement and the appropriate jump statements. As before, one must be sure to use the opposite jump from the relational. Also as with the `if-then-else` structure, one must be careful to include an extra unconditional jump, but in this case the jump is back to the beginning of the loop:

```

                                mov i,1
                                ; while i<=3
while01: cmp i,3
        jg endw01
        ; body of loop
        inc i
        jmp while01
endw01: nop

```

Keeping the label scheme used with the `if` structure, numbers are used to avoid multiple label names with the same name when more than one loop is used in a program. As mentioned above, notice the inclusion of the unconditional `jmp while01` at the bottom of the loop, because without it, the loop would execute the body of the loop only once.

As an example of how the `while` loop could be used, some very small microprocessors do not have a multiplication or a division instruction as part of their instruction set. These processors are not designed to solve mathematical problems, but rather to control devices. Further, these types of processors have very little memory and are known as embedded systems. Although a processor might not have a multiplication instruction, it would have a way to perform iteration. If multiplication does need to be performed, one way is to implement it as repetitive addition. For the sake of convenience, assume that multiplication can only be between non-negative numbers, similar to the `mul` instruction in the Intel processor:

```

ans=0;
i=1;
while(i <= y) {
    ans=ans+x;
    i++;
}

```

In the C code above, should y be 0, then the loop will never be executed and the `ans` will be equal to 0. However, if x is 0, then it is possible that the loop

will iterate y times and redundantly add x to ans . How could this problem be solved? An if statement could be added as shown below:

```
ans=0;
if(x!=0) {
    i=1;
    while(i<=y) {
        ans=ans+x;
        i++;
    }
}
```

The above C code can be implemented in assembly language below illustrating the use of the `.while` directive and further illustrating the use of the `.if` directive:

```
mov ans,0 ; initialize ans to 0
.if x != 0
mov i,1 ; initialize i to 1
mov eax,y ; load eax with y for while
.while i<=eax
mov eax,ans ; load eax with ans
add eax,x ; add eax to ans
mov ans,eax ; store eax in ans
mov eax,y ; reload eax with y for while
inc i ; increment i by 1
.endw
.endif
```

Note in the `.while` directive, as with the `.if` directives in the previous chapter, one of the variables is moved into the `eax` register to be compared in the `.while` directive. Using `eax` might seem incorrect at first, since it is being used for addition in the body of the loop, but since `ans` is loaded into and subsequently stored back into `ans`, using `eax` is acceptable. Note that toward the bottom of the body of the loop, the value of y is copied back into `eax` for the subsequent times through the loop. Although another register could have been used, this method minimizes the number of registers used and on completion of the segment, all the values of the respective memory locations contain the same values as in the previous C code. Alternatively, instead of using i as a loop control variable, a register could be used as shown below. Since `ecx` is known as the counter register, this would be a good choice and would cause the loop to execute quicker at the expense of using one more register to implement the code. If the variable i needs to contain the corresponding final value as it would in the previous C code, then the value in `ecx` can simply be moved into the variable i at the end of the segment as illustrated below. Provided one does not branch out of the middle of the loop, which would result in unstructured code and should be avoided, the code below would work and should be acceptable:

```

mov ans,0 ; initialize ans to 0
.if x!= 0
mov ecx,1 ; initialize ecx to 1
.while ecx<=y
mov eax,ans ; load eax with ans
add eax,x ; add x to ans
mov ans,eax ; store eax in ans
inc ecx ; increment ecx by one
.endw
mov i,ecx ; store ecx in i
.endif

```

Although embedded processors might not have high-level directives, the solution to the problem is still the same and the above can be implemented with compares and jumps in their respective assembly languages. The result would be similar to implementing the code without directives in MASM, which is left as an exercise at the end of the chapter.

1.2 Post-test Loop Structures

The C programming language has a post-test loop structure called the do-while. The unique feature of post-test loops is that the body of the loop is executed at least one time, unlike the pre-test loop where the body of the loop might not be executed at all. It is because of this difference that the pre-test loop is sometimes used more often than the post-test loop, but the latter is helpful in various circumstances in some languages, such as filtering interactive input or reading files. In MASM, the post-test loop structure is implemented using the .repeat and .until directives. Given the following C do-while loop on the left, the corresponding assembly language appears on the right:

```

i=1; mov i,1
do { .repeat
    // body of loop ; body of loop
    i++; inc i
} while (i<=3); .until i>3

```

Note that instead of $i \leq 3$, the .until has $i > 3$. This is not a mistake. Whereas the do-while continues to loop while i is less than or equal to 3 and falls through when i is equal to 4, the repeat-until loops until i is greater than 3, where it also falls through to the next instruction when i is equal to 4. This is similar to other languages such as Pascal and VBA (Visual Basic for Applications), where the latter has both do-while and repeat-until instructions. One just has to be careful to use the exact opposite relational and must not forget to consider the case where the values are equal, which could result in a subsequent logic error. The implementation of this loop without MASM high-level directives using compares, jumps, and labels is shown below. As before, the relation of the jump is reversed from the one in the .until directive above:


```

mov i,1
repeat01: nop
        ; body of loop
        inc i
        cmp i,3
        jle repeat01
endrpt01: nop

```

Implementing the multiplication of the previous section using `.repeat .until` directives requires a little rethinking, since the body of a post-test loop structure is executed at least once. If the value of y is equal to 0, then the loop will execute once and the answer will be incorrect. As a result, there needs to be an if statement prior to either `do-while` in the C code on the left or the `.repeat . until` directives in the assembly code on the right. The requirement that there often needs to be an if statement prior to a post-test loop is one of the reasons why these types of loops are not usually the first choice when solving many problems:

```

ans=0;
if (y!=0) {
}

```

```

i=1; mov ecx,1
do \{ .repeat
    ans=ans+x; mov eax, ans
    i++; add eax,x
\} while (i<=y); mov ans, eax

```

```

mov ans,0
.if y!=0
inc ecx
.until ecx>y
mov i,ecx
.endif

```

However, if the value of x is checked to see if it is equal to 0 as done in the while loop in the previous section, then this is not much of an imposition, where the values of both x and y could be checked in the if statement using an "and" operation. As with the previous section, notice that the above code is implemented using the `ecx` register. The advantages and disadvantages of using `ecx` as a loop control variable will be made more apparent in the next section where it is not an option, but rather a requirement. Given that an if statement is often needed prior to the `.repeat`, the `.while` loop will tend to be used more often in this text.

1.3 Fixed-Iteration Loop Structures

As found in many high-level languages, there usually exists a fixed-iteration loop structure often called a for loop structure. Its primary advantage is that it is used when a loop needs to iterate only a fixed number of times. An example of such a loop is the for loop in C, where the braces are optional when there is only one statement:

```
for(i=1;i<=3;i++) {  
    // body of loop  
}
```

Typically most machine architectures have a specialized instruction to accomplish this task and it can often execute a little faster than the loops discussed in the previous two sections. In MASM, the directives that can be used for this task are the `.repeat` and `.untilcxz` directives. If one recalls when the general purpose registers were first introduced in Chap. 1, it was mentioned that the `ecx` register was sometimes used as a counter register. As with some other instructions, the `.repeat` and `.untilcxz` directives use the `ecx` register as a counter. Unlike using the separate compare and jump instructions in the previous two loop structures, the `.untilcxz` directive performs two tasks: it decrements the `ecx` register by 1 and then jumps to the `.repeat` directive when `ecx` is not equal to 0. In other words, it loops until the `ecx` register equals 0 (`cxz`). Unfortunately, unlike the for statement which is typically implemented as a pre-test loop, the `.repeat` and `.untilcxz` directives are implemented as a postfix loop structure, which means the body of the loop is executed at least once. However, if one is careful with the `.repeat` and `.untilcxz` directives, they can prove to be very useful. An understanding of how it works can be helpful later in Chap. 9 when learning how to manipulate strings. The above for loop can be implemented as follows:

```
mov ecx,3  
.repeat  
; body of the loop  
.untilcxz
```

First, the `ecx` register is loaded with the number of times the body of the loop should be executed. Then, each time the `.untilcxz` directive is executed, the value of the `ecx` register is decremented by 1 and then compared to 0. If the value is not equal to 0, the loop repeats. When the value is 0, the flow of control is passed onto the instruction immediately following the `.untilcxz` directive. One temptation that the beginning assembly language programmer has is to decrement the `ecx` register within the body of the loop. Just like with a for loop in many high-level languages, where the loop control variable should not be altered in the body of the loop, neither should the `ecx` variable be altered in the body of the `.repeat - .untilcxz` loop.

As with the other structures before, the `.repeat` and `.untilcxz` directives can be implemented using only assembly language instructions. In this case, it is the `loop` instruction which implements the `.repeat` and `.untilcxz` directives:

```

                                mov ecx,3
for01: nop
    ; body of the loop
    loop for01
endfor01: nop

```

The loop instruction works the same as the `.repeat` and `.untilcxz` directives, where the `ecx` register is loaded with the number of times to iterate and the loop instruction then decrements the `ecx` register by 1, branches to the label indicated in the operand field when `ecx` is not equal to 0, and falls through otherwise.

With both the `.repeat` and `.untilcxz` directives and the loop instruction, one has to be careful that the `ecx` register does not contain a 0 or a negative number. If one takes a moment and thinks of how the loop instruction works, the potential problem should be apparent. If `ecx` is initially 0, then what is the first thing the loop instruction does? It decrements the `ecx` register by 1, thus causing `ecx` to be a negative one. Since it is not 0, it branches back to the beginning of the loop and the process continues. Would this be an infinite loop? No, because the loop instruction would continue to decrement `ecx` until it hits `-2,147,483,648` and on the subsequent decrement, the negative number would turn to a positive `2,147,483,647` (see Appendix B). Eventually it would decrement that number back to 0. Although it would not be an infinite loop, it would loop over four billion times.

If the value is merely being assigned to `ecx` prior to the loop, this might not be as much of a problem. However, if, for example, the value of `ecx` is being input from a user, then the value `ecx` should be checked. An instruction that can help with this problem is the `jecxz` instruction that will jump to a label after the loop should the value of `ecx` be equal to 0. This is especially useful when using the loop instruction where labels are already being used. Although this works to prevent a situation where `ecx` might contain a 0, it does not check for negative values which can cause just as much havoc as a value of 0. If necessary, an if structure can be used to check for a non-positive value and the `.if` directive would work well when using the `.repeat` - `.untilcxz` directives. An example of each can be seen below, where the value of `ecx` is not assigned and it can be assumed to already have a value that needs to be checked:

```

; check for zero
; check for non-positive
jecxz endfor01
.if ecx > 0
for01:
nop
    .repeat
    ; body of the loop
    ; body of the loop
    loop for01
.untilcxz

```

```
endfor01:
nop
.endif
```

Another problem with this loop is that the `.repeat` directive can only be 128 bytes prior to the `.untilcxz` directive, or the label referenced in a loop instruction can also only be 128 bytes prior to the loop instruction. As will be discussed in Chap. 10, the instructions used in this text vary from 1 byte (such as the `inc eax` instruction) to 6 bytes (such as a `mov ebx, amount` instruction). Given that the loop instruction is 2 bytes long, there could be 126 one-byte instructions and in the worst case only 11 six-byte instructions could be in the body of a `.repeat - .untilcxz` loop. If the number of bytes is exceeded, the assembler will generate an error message indicating how many bytes the loop has exceeded this limit. Although this seems rather restrictive, in practice this does not occur too often, and if it does, a `.while` loop can always be used instead.

In spite of some of the above restrictions, the `.repeat-untilcxz` directives and loop instructions are very useful and can be used in a variety of situations. For example, the previous multiplication problem can be solved with this loop as well:

```
ans = 0;
if (y != 0)
    for(i=1; i<=y;i++)
        ans = ans + x;
```

```
mov ans,0
.if y != 0
mov ecx,y
.repeat
mov eax,ans
add eax, x
mov ans,eax
.untilcxz
.endif
```

Note that in the assembly code on the right, the final value of the `ecx` register is not moved into `i`. The reason is that unlike the previous two loops, where `i` or `ecx` started at 1 and ended up being one more than the variable `y`, here the `ecx` register starts at the variable `y` and counts down to zero. It is possible that the final value of `i` could be mathematically calculated to be equal to the correct value whether through normal termination or branching outside from somewhere within the loop. However, in many languages the final value of the loop control variable in a fixed-iteration loop structure is said to be indeterminate, and the result here is consistent with those languages.

1.4 Loops and Input/Output

If a code segment needs to be written to input, process, and output a fixed number of items, then a fixed-iteration loop is probably the best choice. Although a simple example, assume that the segment needs to input and sum exactly 10 integers and then output the sum:

```
sum=0;
for(i=1; i<=10; i++) {
    printf("%s","Enter an integer: ");
    scanf("%d",&num);
    sum=sum+num;
}
printf("\n%s%d\n", "The sum is ",sum);
return 0;
```

Assuming all the formats and variables are declared correctly, the partial equivalent in assembly is shown below. Note that the value of `ecx` is stored in a memory location called `temp` at the top of the loop and then the value of `ecx` is restored at the bottom of the loop. Recall from Chap. 2 that the `INVOKE` directive can destroy the `eax`, `ecx`, and `edx` registers. Since the `.repeat-untilcxz` directive uses the `ecx` register, care must be taken to save and restore its value. In the next chapter, the stack will be discussed and will be a convenient way to accomplish the same task:

```
        .data
msg1 byte "Enter an integer: ",0
msg2 byte "The sum is ",0
        .code
mov sum,0
mov ecx,10
.repeat
mov temp,ecx
INVOKE printf, ADDR msg1fmt, ADDR msg1
INVOKE scanf, ADDR in1fmt, ADDR num
mov eax,sum
add eax,num
mov sum,eax
mov ecx,temp
.untilcxz
INVOKE printf, ADDR msg2fmt, ADDR msg2, sum
```

Should more or fewer than 10 numbers need to be input and summed, the above code segments would be rather restrictive. In order to allow more versatility, a prompt and input for the number of integers could be added prior to the loop. Also, to help avoid the problems of entering a 0 or a negative number, the loop could be changed from a `.repeat - .untilcxz` to a `.while - .endw`. Again,

assuming the formats and data declarations are correct, the partial program is as follows:

```
.data
msg0 byte "Enter the number of integers to input: ",0
msg1 byte "Enter an integer: ",0
msg2 byte "The sum is ",0
.code
mov sum,0
INVOKE printf, ADDR msglfmt, ADDR msg0
INVOKE scanf, ADDR in1fmt, ADDR count
mov ecx,1
.while ecx<=count
mov temp,ecx
INVOKE printf, ADDR msglfmt, ADDR msg1
INVOKE scanf, ADDR in1fmt, ADDR num
mov eax,sum
add eax,num
mov sum,eax
mov ecx,temp
inc ecx
.endw
INVOKE printf, ADDR msg2fmt, ADDR msg2, sum
```

However, what if one did not want to enter the number or integers to be input and summed? A convenient way to solve this problem is to use what is known as a sentinel-controlled loop, or what is sometimes called an end-of-data loop (EOD). As commonly presented in a first semester computer science text, it contains two input statements, where the first one appears prior to the loop and is sometimes referred to as a priming read and the second one appears as the last statement in the body of the loop. The appearance of two input statements sometimes confuses the beginning programmer, but remembering the three parts of any loop, "initialization," "test," and "change," this priming read can be thought of as the initialization portion of the loop. Next, the test does not check a loop control variable (LCV), but rather the test is of the value input. Of course this is not to say that a counter cannot be added to the loop, where the counter may or may not be part of the control of the loop. Lastly, the second input statement appears as the last statement in the body of the loop, which serves as the change in the loop:

```
sum = 0;
printf("%s", "Enter an integer or a negative integer to stop: ");
scanf("%d",&num);
while (num > 0) {
    sum=sum+num;
    printf("%s", "Enter an integer or a negative integer to stop: ");
    scanf("%d",&num);
}
```

```

}
printf("\n%s%d\n", "The sum is ", sum);

```

Again assuming that all the formats and variables are declared properly, the partial assembly language equivalent is as follows:

```

.data
msg1 byte "Enter an integer or a negative integer to stop: ",0
msg2 byte "The sum is ",0
.code
mov sum,0
INVOKE printf, ADDR msg1fmt, ADDR msg1
INVOKE scanf, ADDR in1fmt, ADDR num
.while num >=0
mov eax,sum
add eax,num
mov sum,eax
INVOKE printf, ADDR msg1fmt, ADDR msg1
INVOKE scanf, ADDR in1fmt, ADDR num
. endw
INVOKE printf, ADDR msg2fmt, ADDR msg2, sum

```

It is possible to implement a sentinel-controlled loop using only one input statement, but this is not considered good programming practice. This is usually not a problem in many high-level languages because of the lack or discouraged use of a goto statement. If it is not a good method, then why present it here? The reason is that some older programs may have been written with this style of loop and should the code need to be debugged or modified, then knowledge of this type of loop might be helpful. At the same time, by understanding the loop, the disadvantages of such a structure can be understood and its use in the future limited.

This style of loop is written by having only one input statement in the body of the loop and then comparing the value to see if it is equal to the sentinel value, and if so branching out of the middle of the loop. The actual loop is often written to possibly loop an infinite number of times so that the only way out of the loop is the comparison from somewhere in the middle of the loop. Another and possibly more common way that these loops are written is that the loop itself is controlled by a loop control variable, which is then the default number of times to loop. Should the sentinel value be encountered prior to the default number of times, the branch is then taken to some point outside the loop. Using the equivalent of a goto statement, this can be anywhere else in the program, creating a very hard to follow program known as spaghetti code. In C, this effect can be minimized by using the break statement, which restricts the branch to the end of the current structure, not unlike with the switch statement discussed previously. The following C code segment loops infinitely until the sentinel value is detected by the if statement and the break is executed:

```

sum = 0;
while (1) {
    printf("%s", "Enter an integer or a negative integer to stop: ");
    scanf("%d", &num);
    if (num < 0)
        break;
    sum = sum + num;
}
printf("\n%s%d\n\n", "The sum is ", sum);

```

Since the above code does not exit from either the beginning or the end of the loop but rather from the middle, some purists would say that the code is unstructured, while others would say that since the break statement limits the branch to the end of the while statement, it is an acceptable branch. In the purist's defense, it is possible that the if might not appear as the first statement immediately after the scanf and it might be tempting for a programmer to add code prior to the if statement. This can possibly introduce a potential logic error that when the loop terminates, some processing has occurred that should not have prior to the input of a sentinel value. Although it might be counter-argued that was the intended reason for placing code prior to the if statement anyway, it does introduce the possibility of misplaced code and subsequent unintended logic errors, especially by beginning programmers. Unlike the break statement in C which only branches to the end of the loop structure, in assembly the jump instructions can branch anywhere in the program. So, the endl: label in the code segment below does not necessarily need to be located at the end of the loop but placed anywhere in the program, even pages away from the loop itself which is not recommended. The above C code can be implemented as shown in the partial assembly language code segment below:

```

.data
msg1 byte "Enter an integer or a negative integer to stop: ",0
msg2 byte "The sum is ",0
.code
mov sum,0
.while 1
INVOKE printf, ADDR msg1fmt, ADDR msg1
INVOKE scanf, ADDR in1fmt, ADDR num
cmp num,0
jl endl
mov eax,sum
add eax,num
mov sum,eax
.endw
endl:
nop
INVOKE printf, ADDR msg2fmt, ADDR msg2, sum

```


As mentioned above, the result of this approach of using the single input is some-what controversial and will not be used in this text. However, if the instructor of the class says it is okay to use this method or if it is used extensively at work, then hopefully the reader has learned the potential dangers of this method and will thus use it sparingly and carefully.

1.5 Nested Loops

As one might recall from a first or a second computer science course, nested loops are helpful when accessing a two-dimensional array or with various sorting algorithms, such as the selection, bubble, or insertion sorts. In this section, the equivalent of nested while and for loop structures will be introduced and the actual application of nested loops will be deferred until Chap. 8.

As might be suspected, the implementation of a nested while loop is not much more difficult than the implementation of a nested if statement. The most important thing that should be remembered is to be sure to use different loop control variables for each of the loops as shown in the C code segment below to the left and the assembly code using high-level directives to the right:

```
i = 1; mov i,1
while(i<=2) { .while i <= 2
    j=1;
    while(j<=3) {
        // body of nested loop ; body of nested loop
        j++; inc j
    }
    i++; inc i
} .endw
```

Of course, the code gets a little more complicated without the benefit of high-level directives. One must pay particularly close attention to the conditional jumps to make sure that the relational is reversed properly and also that the unconditional jumps branch to the appropriate location as shown below:

```
    mov i,1
while02: cmp i,2
    jg endwhile02
    mov j,1
while03: cmp j,3
    jg endwhile03
    ; body of nested loop
    inc j
    jmp while03
endwhile03: nop
    inc i
    jmp while02
```

endwhile02: nop

Often when dealing with two-dimensional arrays and sorting algorithms, when there are a fixed number of times to loop in both the outer and inner loops, a fixed-iteration loop structure is understandably used for convenience and speed as shown below:

```
for (i = 1; i <= 2; i++)
    for (j = 1; j <= 3; j++) {
        // body of nested loop
    }
```

As learned previously, the above can be implemented using the .repeat and .untilcxz directives in assembly language. However, care must be taken when writing the code or it can be implemented incorrectly as shown below:

```
; *** Caution: Incorrectly implemented code ***
mov ecx,2
.repeat
mov ecx,3
.repeat
; body of nested loop
.untilcxz
.untilcxz
```

What is wrong with the above code? Although it is syntactically correct, there is a logic error. Unlike the while loops above that can use two different variables, or alternatively could use two different registers, the .repeat and .untilcxz directives can work with only one register. The result is that the value of ecx is 0 upon completion of the nested loop, which causes the outer loop to never terminate. Could another register be used with the .repeat - .untilcxz directives? Unfortunately no, because as should be recalled, the underlying loop instruction that implements the .untilcxz directive works only with the ecx register. The question then is how can the above problem be solved? One way is to store the

value of ecx in a memory location prior to the inner loop and then restore the value of ecx when the inner loop is complete as follows, where it can be assumed that tempecx is declared as a temporary memory location:

```
; *** Note: Correctly implemented code ***
mov ecx,2
.repeat
mov tempecx,ecx
mov ecx,3
.repeat
; body of nested loop
.untilcxz
mov ecx,tempecx
.untilcxz
```

There is of course another way to save and restore the value of `ecx`. Again, a stack is a convenient way to save and restore registers, and this will be discussed in the next chapter. Could the above be implemented without using the `.repeat` and `.untilcxz` directives? Of course, by using the loop instruction and that is left as an exercise at the end of this chapter.

1.6 Complete Program: Implementing the Power Function

The selection structures of the previous chapter and the iteration structures of this chapter can obviously be combined. As an example of a complete program, consider the implementation of the power function (x^n), where an iterative definition of the power function is as follows:

x^n = If $x < 0$ or $n < 0$, then negative message

Else if $x = 0$ and $n = 0$, then undefined message Else if $n = 0$, then 1

Otherwise $1 * x * x * \dots * x$ (n times)

For the purposes of this program, it will not calculate the case where either x or n is negative and should both x and n be 0, the result is undefined. In each case, an appropriate error message is output. The following C program implements the above definition:

```
#include <stdio.h>
int main() {
    int x,n,i,ans;
    printf("%s","Enter x: ");
    scanf("%d",&x);
    printf("%s","Enter n: ");
    scanf("%d",&n);
    if(x<0 || n<0)
        printf("\n%s\n\n","Error: Negative x and/or y");
    else
        if(x==0 && n==0)
            printf("\n%s\n\n","Error: Undefined answer");
        else {
            i=1;
            ans=1;
            while(i<=n) {
                ans=ans*x;
                i++;
            }
            printf("\n%s%d\n\n","The answer is: ",ans);
        }
    return 0;
}
```

The above can be implemented in assembly language using directives and as follows:

```

        . 68
        .model flat,c
        .stack 100h
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
printf PROTO arg1:Ptr Byte, printlist:VARARG
.data
in1fmt byte "%d",0
msg1fmt byte "%s",0
msg3fmt byte "%s%d",0Ah,0Ah,0
errfmt byte "%S",0Ah,0Ah,0
errmsg1 byte 0Ah,"Error: Negative x and/or y",0
errmsg2 byte 0Ah,"Error: Undefined answer",0

        msg1 byte "Enter x: ",0
        msg2 byte "Enter n: ",0
        msg3 byte 0Ah,"The answer is: ",0
x
n
ans
sdword ?
i sdword ?
. code
main proc
    INVOKE printf, ADDR msg1fmt, ADDR msg1
    INVOKE scanf, ADDR in1fmt, ADDR x
    INVOKE printf, ADDR msg1fmt, ADDR msg2
    INVOKE scanf, ADDR in1fmt, ADDR n
    .if x<0 || n<0
    INVOKE printf, ADDR errfmt, ADDR errmsg1
    .else
    .if x==0 && n==0
    INVOKE printf, ADDR errfmt, ADDR errmsg2
    .else
    mov ecx,1
    mov ans,1
    .while ecx <= n
    mov eax,ans
    imul x
    mov ans,eax
    inc ecx
    .endw
    mov i,ecx
    INVOKE printf, ADDR msg3fmt, ADDR msg3, ans
    .endif
main endp

```

```

    .endif
    ret
main endp
end

```

The implementation of the above MASM code is fairly straightforward and follows the corresponding C program. Note that `ecx` is used for loop control, but the value of `i` is updated upon completion of the loop to reflect the logic of the corresponding C code.

1.7 Summary

- The `.while - .end` directives implement a pre-test loop structure.
- The `.repeat - .until` and `.repeat - .untilcxz` directives are both post-test loop structures.
- The `.repeat - .untilcxz` directives are a fixed-iteration loop structure.
- The loop instruction underlies the `.repeat - .untilcxz` directives.
- As with the `.if` directive, the `.while - .end` and `.repeat - .until` directives cannot compare memory to memory due to the underlying `cmp` instruction.
- Be extra careful to initialize `ecx` to a positive number (not zero or negative) when using the loop instruction or the `.repeat - .untilcxz` directives. The `jecxz` instruction or an `.if` directive, respectively, can be helpful in avoiding this problem.
- When using either the loop instruction or the `.repeat - .untilcxz` directives, it is a good idea to not alter the contents of the `ecx` register in the body of the loop.
- When nesting `.repeat - .untilcxz` directives or loop instructions, be careful to save and restore the `ecx` register just before and after the inner loop.
- The beginning of the `.repeat - .until` directives or loop instruction cannot be more than 128 bytes away.

1.8 Exercises (Items Marked with an * Have Solutions in Appendix D)

1. Given the following assembly language statements, indicate whether they are syntactically correct or incorrect. If incorrect, indicate what is wrong with the statement:

*A.

```
.for i=1 ; i<=3 ; i++
;body of loop
.endfor
```

B.

```
mov i, 1
while i <= x
;body of loop
inc i
.endw
```

*C.

```
mov i,0
.repeat
; body of loop
add i,2
.until i>10
```

D.

```
mov edx,3
.repeat
;body of loop
.until edx
```

E.

```
mov ebx,2
.do
;body of loop
.while ebx>0
```

2. Implement the last code segment in Sect. 5.1 without using directives and using only conditional and unconditional jumps.
3. Given the following while loops implemented using conditional and unconditional jumps, indicate how many times the body of each loop will be executed:

*A.

1.8. EXERCISES (ITEMS MARKED WITH AN * HAVE SOLUTIONS IN APPENDIX D) 31

```
mov i,2
while04: cmp i,8
jge endwhile04
; body of loop
add i,2
jmp while04
endwhile04: nop
```

B.

```
mov k,0
repeat05: nop
; body of loop
add k,3
cmp k,3
jl repeat05
endrepeat05: nop
```

C.

```
mov j,1
while06:
cmp j,0
jg endwhile06
; body of loop
inc j
jmp while06
endwhile06: nop
```

4. Implement the `.repeat` and `.until` directive at the end of Sect. 5.2 using only compare and jump instructions, along with the appropriate label names.
5. Implement unsigned divide (similar to the `div` instruction) using repetitive subtraction, with your choice (or your instructor's choice) of any of the following (start with the dividend in `eax` and the divisor in `ebx`, then place the quotient in `eax` and the remainder in `edx`. Note: Do not worry about division by zero or negative numbers):
 - *A. `.while`
 - B. `.repeat - .until`
 - C. `.repeat - .untilcxz`
6. Implement the following C segment using the `.repeat - .untilcxz` directives. What if the value of `n` is 0 or negative? Does your code segment still work properly? How can this problem be rectified?

```

sum = 0;
for (i=1; i<=n; i++)
    sum = sum + i;

```

7. Implement the following do-while loop first using the .repeat - .until directives and then using only compares, and conditional and unconditional jumps:

```

i=10;
sum=0;
do {
    sum=sum+i;
    i=i-2;
} while i>0;

```

8. Implement the last code segment in Sect. 5.5 using the loop instruction instead of .repeat and .untilcxz directives.
9. Given the factorial function ($n!$) defined iteratively as follows:

If $n = 0$ or $n = 1$, then 1
 If $n = 2$, then $1 * 2 = 2$
 If $n = 3$, then $1 * 2 * 3 = 6$
 If $n = 4$, then $1 * 2 * 3 * 4 = 24$
 etc.

Implement the above function iteratively with your choice (or your instructor's choice) of any of the following:

- A. .while
- B. .repeat - .until
- C. .repeat - .untilcxz

10. Given the Fibonacci sequence defined iteratively as follows:

if $n = 0$, then 0
 if $n = 1$, then 1
 if $n = 2$, then $0 + 1 = 1$
 if $n = 3$, then $1 + 1 = 2$
 if $n = 4$, then $1 + 2 = 3$
 etc.

Implement the above function iteratively with your choice (or your instructor's choice) of any of the following:

- A. .while
- B. .repeat - .until
- C. .repeat - .untilcxz

1.9 Introduction

As introduced in most first semester computer science courses and previously discussed in Chap. 4, various relationals in an if statement can be connected

via the use of logical operators such as "and" (`&&`), "or" (`||`), and "not" (`!`), where these operators in assembly language work with comparisons between variables, registers, and literals. However, sometimes it is necessary to not just compare the contents of variables or registers but check the individual bits within a memory location or a register. These types of operations are known as bit-wise operations. An example of this is when interfacing with an external device, when often only a single bit is needed to be checked or set on the external device.

As may or may not have been learned in a previous course, one of the reasons why the C-like languages are very popular is that they have some capabilities to manipulate individual bits. Instead of having to learn a particular low-level language for a particular processor, basic bit-wise operations can be done in a high-level language that is transferable from processor to processor, provided there is a C or C++ compiler for that particular processor. Of course, assembly language has these same capabilities by using logic, shifting, and rotating instructions for manipulating the contents of registers and memory locations, as well as built-in instructions for manipulating a stack.

Although previous exposure to both bit-wise manipulations (such as in a course in C or C++) and binary arithmetic (such as in a course in computer organization) is helpful, it is not a requirement for this text since this material is contained in Appendix B. Should one not have the above previous experience, then Appendix B is recommended reading prior to starting this chapter.

1.10 Logic Instructions

There are many times in low-level programming that individual bits need to be set, cleared, tested, or toggled in a register or a memory location. In order to do so, the use of the logic operations "or", "and", and "exclusive-or" can be very useful. As also shown in Appendix B, Table 6.1 can be helpful in summarizing which logic operation is used under which circumstances.

So how are the above logic operations to set, clear, test, and toggle bits implemented in assembly language? As before, it is helpful to start with similar code in a language like C. As mentioned previously, one of the advantages of the C-like languages is their ability to perform some logic operations. To help introduce this topic, what happens if the second ampersand (`&`) is accidentally left off when performing an "and" (`&&`) operation in an if statement? Depending on the compiler or the level of the warning messages set in the compiler, either a warning would be issued or there might be some unintended results. The reason for this is that using only one of the symbols causes a bit-wise logic operation to be performed which might not be what was originally intended. However, this is precisely how logic operations are performed in a programming language like C.

For example, to test if a particular bit is a 1, a single ampersand logical "and" operator (`&`) would be used instead of the double ampersand logical "and" operator (`&&`). In the code segment below, the variable `flag` contains various

bits that are set as a result of some previous operation. The variable maskit is what is known as a mask that has a particular bit set to 1 that will be used to test and others at 0 to filter or clear all the other bits that do not need to be tested in this instance:

```
if(flag & maskit)
    count++;
```

The above is not trying to determine whether both flag and maskit are true, but rather assuming that flag equals 01101110 in binary and maskit equals 00000100 in binary, the above is a bit-wise & operation between flag and maskit, where the result is equal to 00000100 as shown below:

Table 1.1: Logic operations

Operation	Logic
Set	Or
Clear	And
Test	And
Toggle	Xor

```
01101110
00000100
&
00000100
```

Since anything that is not zero is assumed to be true, the then section of the if statement is performed. On the other hand, should flag equal 01101010 in binary, then the result of the bit-wise & operator would be 00000000, where zero is interpreted to be false, and the then portion would not be executed. As one might suspect, given the high-level directives in MASM, the above is very easy to implement as shown below:

```
mov eax, flag
.if eax & maskit
inc count
.endif
```

In the above code segment, why do the contents of maskit need to be moved to a register prior to the .if directive? As indicated in the past two chapters, many assembly languages do not have the benefit of having high-level directives, so it is necessary to use the logic instructions that are part of the instruction set. In the case of the Intel processor, these are and, or, xor, and not as shown in Table 6.2.

Further, just like the compare and arithmetic instructions, these logic instructions cannot have two operands that are memory locations. Given the above, the previous code can be implemented without directives as follows:

```

if01: mov eax, flag
      and eax, maskit
      jz endif01
then01: inc count
endif01: nop

```

Table 6.2 Logic instructions

And instructions	Or instructions	Xor instructions	Not instructions
and reg, reg	or reg, reg	xor reg, reg	not reg
and reg, imm	or reg, imm	xor reg, imm	not mem
and reg, mem	or reg, mem	xor reg, mem	
and mem, reg	or mem, reg	xor mem, reg	
and mem, imm	or mem, imm	xor mem, imm	

Note the use of the *jz* instruction instead of *jne*, where the *jz* instruction was introduced in Chap. 4. The reason for this is that variables such as *flag* and *maskit* are typically unsigned numbers and are not compared to determine if one is greater or less than the other. As before, care must be taken to reverse the jump to allow the if-then logic to work correctly. Is there a way to use both the logic instructions and the high-level directives? Yes, the logic instruction can be used followed by the *.if* directive as shown below. The advantage of this format is that one can gain familiarity with the use of the actual logic instructions and still have relatively clean code using high-level directives:

```

mov eax, flag
and eax, maskit
.if !ZERO?
inc count
.endif

```

As with the *cmp* and other arithmetic instructions, the result of the *and* operation sets various bits in the *eflags* register. These individual bits can be accessed via the logical operators introduced in Chap. 4. In particular, the *ZERO?* operator is of interest here which returns true should the zero flag be set. Thus in the above code segment, if the result of the *and* operation is 1, indicating a bit is set, then *ZERO?* would be false and *! ZERO?* would be true allowing the increment of the variable *count*.

It should be noted that instead of using a variable such as *maskit*, an immediate value in any of the above examples could be used, thus avoiding the need to move the mask into a register prior to the *and* instruction. The disadvantage of this method is that if the mask has to be used many times, the chance of making an error in one of the instances is greater. However, if the mask is going to be used only once, then the literal method is acceptable. Of course, do not forget to use the letter *b* after the literal to indicate a binary number, otherwise a decimal number will be assumed and a logic error would probably occur.

Although it is possible to use decimal numbers with logic operations, they are rarely used in these situations because the exact bit pattern cannot be seen by other programmers. If the number is larger than 8 bits, then a hexadecimal number followed by an *h* can be used. Further, it is always a good idea to show all the bit positions used to help any reader of the code understand how many bits are being compared. Below are three different ways of using an immediate value as a mask:

```
.if flag & 00000010b
inc count
ifo2:
.endif
theno2:
endifo2:
nop
.endif
```

Lastly, it should be noticed that in the process of testing an individual bit, the other bits are cleared to zero. In this sense the and instruction is useful in not only testing bits but clearing bits as well. But what if when testing bits, one does not want to clear the other bits? Although on some processors this is not a possibility, this can be accomplished in MASM with the test instruction which is illustrated at the end of the next section.

Although probably not seen much in a first-year computer science sequence, a particular bit can be set in a memory location in the C programming language using the bit-wise or (*|*) operator:

```
flag = flag | maskit;
```

The same can be accomplished in assembly language using the or instruction. As in previous chapters, notice the high-level comment prior to the code segment below:

```
; flag = flag | maskit
mov eax,flag
or eax,maskit
mov flag,eax
```

As application of the bit-wise or, consider the changing of an uppercase character to a lowercase character. In looking at the ASCII table in Section B.8, notice that the bit pattern for an uppercase character always has bit 5 (sixth from the right) equal to 0, whereas the bit pattern for lowercase characters always has bit 5 equal to 1. For example, the letter "S" is equal to 53 in hexadecimal or 01010011 in binary and the letter "s" is equal to 73 in hexadecimal or 01110011 in binary. In order to convert an uppercase character to a lowercase character, an or instruction using a mask of 00100000 in binary could be used to set bit 5. Assuming the variable letter is declared as a byte and already contains a letter, the following instruction would work:

or letter,00100000b

To convert a lower-case letter "s" to an upper-case letter "S", bit 5 is cleared to zero. This is accomplished using an and operation as follows:

and letter,11011111b

Although C does not have a logical exclusive-or operation that could be used between two relationals in an if or a while statement, it does have a bit-wise exclusive-or operation (^). As mentioned above, the "xor" operation can be used for toggling a bit. For example, if one wanted to toggle bit 1 from a 0 to a 1 or from a 1 to a 0, the following instruction would accomplish that task:

```
flag = flag ^ maskit;
```

which can be written in assembly language using the xor instruction as

```
; flag = flag ^ maskit
mov eax,flag
xor eax,maskit
mov flag,eax
```

On occasion, when examining code previously written by someone else, one might see something similar to the following instruction:

```
xor eax, eax
```

At first, it may seem a bit strange to see a logic operation with the same register for both operands. If this was done with any of the other logic instructions, it would not accomplish anything. For example, with the or operation, 0 or 0 is 0 and 1 or 1 is 1. However, it is in this second case that there is a difference with the xor instruction where 1 xor 1 is 0. In other words, all the bit positions with a 0 remain a 0 and all the bits with a 1 become a 0, thus clearing all the bit positions in the register to 0. Logic instructions are some of the fastest instructions in most processor architectures and using xor eax, eax is usually faster than using mov eax, 0. This is one of those tricks used by experienced assembly language programmers to speed up the execution of a program, but unless it is used in a critical location such as within a loop or nested loops, the speed gained is negligible compared to the loss of readability for inexperienced assembly language programmers.

1.11 Logical Shift Instructions

Sometimes if there are more than one bit to test, set, or toggle in a register or a memory location, it is easier to move the bit patterns instead of having multiple if statements. This can be accomplished by using a shift or a rotate instruction. The C programming language has the ability to shift bits to the left or the right in

a memory location by using the $\ll$ or $\gg$ operators, respectively. For example, if the memory location `num` contained a 2 and the following instruction was executed, the contents of `num` would then be a 16:

```
num = num << 3;
```

Assuming only 8 bits for the sake of convenience, the 2 in `num` would be represented as a 00000010 in binary. Then shifting the bits three places to the left would cause `num` to contain a 00010000, which is a 16 in base 10. Once the bit in

question is in the correct location, the previous logical operators could be applied to the memory location. However, a question arises as to whether to move the bits in the mask or move the bits in either the register or the memory location to be tested. Which should be moved is largely up to the application, the preference of the instructor, or the preference of the programmer. There are, however, some guidelines that can help one make a choice. If only part of the register or the memory location needs to be checked and it is moved, then the original contents of the register or the memory location will be altered. Of course, if it is no longer needed, then this is not a concern, but if the original contents are needed again in the future or if there is a chance that they might be needed, it might be better to move the mask instead.


However, if the mask is relatively complicated, where more than a single bit is set, then it might be better to move the register or the memory location. Also, by shifting the data, the mask can be kept as immediate data and the actual code might be a little cleaner. One technique is to save the original contents in a temporary location, shift the data, and then the original data in the temporary location can always be restored back into the original location later. (An alternative to using a temporary memory location is to use a stack, which will be discussed later in this chapter.) A second option is to move the original to a temporary memory location or register and then shift the temporary location, which would preserve the original contents from alteration. A third technique to be discussed later is to rotate the bits back to the original location so that subsequent logic can access the data in its original format. All three of these methods have their advantages under various circumstances, but for the time being the first alternative will be used to illustrate how data must be saved and restored.

Two very helpful instructions are the logical shift instructions, which shift the contents of a register or a memory location to either the left or the right of a specified number of bits. The instructions are known as logical instructions, because they do not assume the presence of a sign bit. These two instructions are listed in Table 6.3.

Note that on older 8086/8088 processors, the only immediate number in an operand that could be used was 1 so that any other number would first need to be loaded into the `cl` register. Although any number can now be used on newer processors, on occasion, some programmers who originally programmed these older processors might carry on that tradition. These instructions move each bit the

Table 1.2: Shift instructions

Shift left instructions	Shift right instructions
shl reg,cl	shr reg,cl
shl reg,imm	shr reg,imm
shl mem,cl	shr mem,cl
shl mem,imm	shr mem,imm

Figure 1.1: Initial contents of `al` register

number of positions indicated to the left or the right, accordingly. Assuming an 8-bit register is being used, the `shl` instruction moves the contents of the leftmost bit (bit 7) into the carry flag (`CF`), moves the contents of the other 7 bits to the left one bit, and then moves a 0 into the rightmost bit (bit 0). For example, assume that the `al` register contained the bit pattern in Fig. 6.1 and that the content of the carry flag on the left was unknown, as indicated by the question mark.

After the execution of the instruction `shl al, 1`, the 0 in bit position 7 on the left would be moved into the carry flag, the contents of bit position 6 would be moved into bit position 7, the contents of bit position 5 would be moved into bit position 6, and so on, where finally the contents of bit position 0 would be filled with a 0 as shown in Fig. 6.2.

What happens to the previous contents of the carry flag? Some might say it disappears into thin air or some "old timers" might say that its contents are moved to the bit bucket. With respect to the latter, the old timers might sometimes further inform beginning assembly language programmers that when the bit bucket gets full, it needs to be emptied. However, be aware there is no such thing as a bit bucket nor does it need to be emptied. It is just an expression to mean that bit is no longer accessible and more importantly it is merely a way to have a little fun at the expense of beginning assembly language programmers!

The shifting of bits in the reverse direction is also possible, where if a `shr al, 1` were executed on the original contents of `al` in Fig. 6.1, the contents of `al` would appear as shown in Fig. 6.3. The contents of the carry flag would go into the bit bucket (see how convenient the terminology is?), the contents of bit position 0 would move into the carry flag, the contents of bit position 1 would move into

bit position 0, and so on, where a 0 would be placed into bit position 7.




Figure 1.2: Results in al register after shl instruction




Figure 1.3: Results in al register after shr instruction

The carry flag is now drawn on the right side of the al register just for the sake of convenience. For the most part, one usually does not worry about the carry flag when doing logical shifting. In a number of other processors, the end bit usually goes directly to the bit bucket.

Given the original contents of the al register in Fig. 6.1, what would happen if instead of a `shl al,1` the instruction `shl al, 8` was executed? If each bit was shifted to the left eight times, bit 0 would end up in the carry flag, the other 7 bits would be in the bit bucket, and all 8 bits of the register would be filled with zeros. This same concept could be expanded to 32 bits and on rare occasions, one might see code such as this to clear a 32-bit register such as `eax` to zero. However, shift instructions are traditionally some of the slower instructions on most processors and this method can be slower than using either the `mov` instruction or the `xor` instruction mentioned in the last section. If however each of the bits in a register needed to be processed one at a time and the register also needed to be cleared to zeros, then clearing a register to zero is a nice by-product of using the `shl` instruction in a loop as discussed below.

As an example of processing each bit individually and using only a byte instead of a double word to save space, assume that each bit in the al register

represents a device that is connected to the processor. Further, if a bit is a 1 or a 0, it would indicate whether the device is turned on or off, respectively. For this task, also assume that the original data needs to be retained. Given these assumptions, how could one determine how many devices are turned on?

In order to solve this problem, there are a number of questions that need to be answered. First, a loop is obviously needed, but which loop should be used? Since there are a fixed number of bits in a byte, the equivalent of a for loop structure appears to be the best choice, which in MASM is the `.repeat-.untilcxz` loop. Another question is should the mask be shifted or should the register be shifted? As mentioned previously, it often depends on whether the original data needs to be retained and since it was mentioned above that it needs to be retained, it might be simpler to shift the mask. However, as also mentioned previously, shifting the data means that the mask can be kept as immediate data. This makes the actual code a little simpler even though the original data will need to be saved and restored in a temporary memory location. Lastly, should the `shl` or the `shr` instruction be used? Often it really does not make a difference, unless one is only processing the bits on one side of the register or the other. By default it is probably best to process the bits in the order in which the bit positions are numbered, which is from right to left. In the following code segment, assume that the memory location `temp` is declared as an unsigned byte:

```
mov count,0 ; initialize count to zero
mov ecx,8 ; initialize loop counter to zero
mov temp,al ; save al in temp
.repeat
mov ah,al ; mov data in al to ah for testing
and ah,00000001b ; test bit position zero
.if !ZERO? ; is the bit set?
inc count ; yes, count it
.endif
shr al,1 ; shift al right one bit position
.untilcxz
mov al,temp ; restore al from temp
```

Note that the first thing prior to the loop is that the data in `al` is saved in memory location `temp` declared as a byte and the last thing after the loop is that the data in `temp` is restored to the `al` register. Further, notice that there appears to be an extra move instruction at the top of the body of the loop. The reason for this is that since there are bits set to 0 in the mask and as a result the `and` instruction will destroy other bits in the `al` register, the data needs to be moved to another register. Since the upper 8 bits of the `ax` register are not being used, the `ah` register is a good choice. At the bottom of the body of the loop, the `al` register is shifted one bit to the right and then back at the top of the loop the `al` register is again moved into the `ah` register so that the next bit can be tested.

However, would it not be nice to have a method of checking a particular bit without having to destroy the other bits around it? Luckily, the designers of the Intel processor have designed just such an instruction. It is called the test instruction, and instead of performing an actual and operation, it performs what is known as an implied and operation. This means it performs the and operation and sets the eflags register, but it does not actually alter the corresponding register or memory location. Although in the example above the contents of the al register will still be altered due to the shift instruction, the contents will not be destroyed each time through the loop by the test instruction and the above code can be rewritten without the extra mov instruction as shown below:

```
mov count,0 ; initialize count to zero
mov ecx,8 ; initialize loop counter to zero
mov temp,al ; save al in temp
.repeat
test al,00000001b ; test bit position zero
.if !ZERO? ; is the bit set?
inc count ; yes, count it
.endif
shr al,1 ; shift al right one bit position
.untilcxz
mov al,temp ; restore al from temp
```

1.12 Arithmetic Shift Instructions

Besides the logical shift instructions, there are also the arithmetic shift instructions called sal and sar, which stand for shift arithmetic left and shift arithmetic right, respectively. The arithmetic shifts have the same operand formats as their logical counterparts as shown in Table 6.4. Unlike the logical instructions, the arithmetic shifts assume that the leftmost bit in a register or memory location is a sign bit. With the sal instruction, this does not make any difference and it performs the same as the shl instruction because when shifting to the left, the leftmost bit is moved into the carry flag whether the leftmost bit is or is not a sign bit. Using the original data from Fig. 6.1, the result of the sh1 in Fig. 6.2 is the same as the sal in Fig. 6.4.

Although the sal instruction works the same as the shl instruction, this is not the case with the sar instruction. With the sar, the leftmost bit is copied to the bit to the right as with the slr, but instead of bringing in a 0 into the leftmost position, the leftmost position is copied into itself, thus preserving the sign bit as shown in Fig. 6.5.

The result is that the arithmetic shifts can be used for performing arithmetic. Again for the sake of simplicity in the following examples, 8-bit registers and memory locations will be used instead of 32-bit ones. Assume that the memory location number contains a 5 as 00000101 in binary. If this number was shifted to the left one bit, the result would be 00001010, which is the number 10 and

Shift arithmetic left	Shift arithmetic right
sal reg, cl	sar reg, cl
sal reg, imm	sar reg, imm
sal mem, cl	sar mem, cl
sal mem, imm	sar mem, imm

Table 1.3: Arithmetic shift instructions




Figure 1.4: Results in al register after sal instruction




Figure 1.5: Results in al register after sar instruction

is twice the number 5. Likewise, shifting number to another bit position results in 00010100, which is the number 20. In other words, for every bit position shifted to the left, the number in the register is effectively multiplied by a power of two. Likewise, shifting the number 00010100 two bits to the right results in the number 00000101 and is the equivalent of dividing by 4.

However, what about negative numbers? For example, when shifting a -4, 11111100 one bit to the right using a shr instruction, the result would be a 01111110, which is a positive 125 and is clearly incorrect. This is the reason for the arithmetic sar instruction, which would cause the sign bit to be copied not only to the right but also back into the leftmost bit position. This would result in a 11111110, a -2 which is correct. If the original number was a -5, 11111011, then a sar al, 1 would result in a 11111101, which is a -2 demonstrating that it is

the equivalent of integer division, where the remainder would be in the carry flag.

When multiplying, choosing either the `sal` or the `shl` instruction is not really a problem, because to multiply 1111100 by 2 the result would be 1111000 in binary, which is a -8. The reason why one would use a `sal` instead of a `shl` in this instance is to indicate to others that the purpose of the shift is the arithmetic operation of multiplication. Although a good comment is always in order, this is in a sense an example of self-documenting code.

However, should a number be shifted too many times to the left, a negative number would eventually become a positive number. Although the numbers used in the examples in this text will be smaller, as with the arithmetic instruction counterparts from Chap. 3, one should always be careful with the possibility of overflow and underflow.

So given the above and for further practice, how would the following C statement be implemented using a shift instruction?

```
product = num1 * 8;
```

Assuming 32-bit words and using the `sal` instruction, it could be implemented as follows:

```
; product = num1 * 8;
mov eax,num1 ; load eax with num1
sal eax,3    ; multiply by 8
mov product,eax ; store eax in product
```

Notice that `num1` above is not shifted, but rather it is first moved to the `eax` register and then it is shifted. As in Chap. 2, `num1` appears to the right of the assignment symbol in the original high-level code implementation and it should not be altered. Another common mistake made by beginners is putting the multiplier in the second operand of the shift instruction instead of the number of positions to be shifted. For example, if the number 8 was accidentally used as the second operand in the shift instruction in the above example, it would cause `num1` to be multiplied by the number 256 which is clearly incorrect. As another example to help reinforce these concepts, how would one implement the following using shift instructions?

```
answer = amount / 4;
```

With positive numbers the choice of `slr` or `sar` does not really make a difference. For example, if `amount` in the above example contained a 32 (00100000 in binary), then after shifting to the right 2 times to divide by 4, the result would be 8 (00001000 in binary). Since the sign bit is 0, it would not make a difference whether a 0 was shifted into bit position 7 by a `slr` instruction or it was copied onto itself by a `sar` instruction.

However, it cannot always be known whether the number in a memory location such as `amount` is positive or negative. Unlike above with multiplication,

the choice of which shift to use, logical or arithmetic, is of critical importance with division when dealing with negative numbers or the possibility of negative numbers. Consider the following implementation:

```
; *** Caution: Possible incorrect code ***  
; answer = amount / 4  
mov eax,amount  
shr eax,2    ; divide by 4  
mov answer,eax
```

Again assuming 8-bit words for convenience, if in the above code segment, amount contains a -8, or 11111000 in binary, then shifting the eax register to the right two positions would result in a 00111110 in binary, or 62 in decimal, which is clearly not correct. However, if the above code segment is rewritten with the appropriate instruction as follows:

```
; *** Note: Correctly implemented code ***  
; answer = amount / 4  
mov eax,amount  
sar eax,2 ; divide by 4  
mov answer,eax
```

Then when 11111000 in binary is shifted to the right arithmetically, the result is that eax contains 11111110 in binary, because in addition to being moved the right, the sign bit is copied back into bit 7 and answer is a -2 as it should be.

The result is that when using shift instructions to perform multiplication and division, it is best to use the arithmetic versions to not only help alert other programmers that the shift is being done for the purposes of arithmetic but also avoid a potential logic error in the event that the quotient in the division operation is a negative number. If the code using shifts is not as clear as using imul and idiv, why would one want to use this method? The answer is that it can be faster and more convenient than are its multiplication and division counterparts, especially when multiplying and dividing by multiples of two, respectively.

1.13 Rotate Instructions

There are many cases with shifting where the unused bits are not needed and their disappearance into the bit bucket is not a problem, especially with multiplication and division. However, there are other cases where it might be convenient to keep the unused bits. The instructions that help in these cases are known as rotate instructions. The rotate instructions are similar to the logical shift instructions in that the end bit goes into the carry flag, and the previous contents of the carry flag go into the proverbial bit bucket. However, instead of zeros being inserted at the other end as with the logical shift instructions, with the rotate instructions the bits from the one end are carried around and inserted into the other end as will be shown shortly.

The format of the two rotate instructions can be found in Table 6.5, where `rol` means rotate left and `ror` means rotate right. Although there are two other rotate instructions, `rcr` and `rer` that rotate out of the carry flag, they will not be considered here. Using the same initial drawing from Fig. 6.1 used previously with the shift instructions and as repeated in Fig. 6.6, a `rol al, 1` instruction would have the results shown in Fig. 6.7.

Similarly, rotating the drawing in Fig. 6.6 to the right using `ror al, 1` would work as shown in Fig. 6.8.

Table 1.4: Rotate instructions

Rotate left instructions	Rotate right instructions
<code>rol reg,cl</code>	<code>ror reg,cl</code>
<code>rol reg,imm</code>	<code>ror reg,imm</code>
<code>rol mem,cl</code>	<code>ror mem,cl</code>
<code>rol mem,imm</code>	<code>ror mem,imm</code>




Figure 1.6: Initial contents of al Register

Fig. 6.8 Results in al register after `ror` instruction The advantage of the rotate instructions is that if the bits are rotated the exact number of times as there are bits in a register or a memory location, the register or the memory location is returned back to its original state. The advantage of this is that there is no need to save or restore the register or the memory location prior to testing it and the same would apply if one were to rotate the mask instead of the data.

The example of testing the 8 bits shown previously is now redone below using a rotate instruction instead of a shift instruction:

```
mov count,0 ; initialize count to zero
mov ecx,8 ; initialize loop counter to zero
.repeat
test al,00000001b ; test bit position zero
.if !ZERO? ; is the bit set?
```




Figure 1.7: Results in al register after rol instruction

```
inc count ; yes, count it
.endif
rol al,1 ; shift al left one bit position
.untilcxz
```

Again, the advantage here of using a rotate instruction instead of a shift instruction is that the contents of the al register need not be saved and restored. The only danger is that sometimes when only part of the register or the memory location needs to be processed, one might forget to rotate the rest of the register or the memory location back to its original location, or inadvertently rotate the wrong number of times, which could lead to a logic error later in the program. When in doubt, one can always save and restore the register or the memory location, whether either a shift or a rotate instruction is used.

1.14 Stack Operations

If one took the second semester computer science course that is usually required in a computer science major or minor, there is a good chance that one has been exposed to the data structure called the stack and the related methods or functions, push and pop. A stack is a LIFO (last in first out) structure, where the last item pushed onto the stack is the first one popped off the stack. As should be recalled, there are number of useful applications for stacks. Some of those applications include reversing data, matching, number conversions, evaluation of expressions, and implementing recursion. Given the usefulness of a stack, most processors include built-in stack instructions and this is true with the Intel processor as well.

In order to use the stack instructions, one must be sure to reserve memory space for the stack itself. As introduced in Chap. 1, this is accomplished by using the `.stack` directive which indicates how much memory should be reserved, where typically 100 hexadecimal bytes, or in other words 256 decimal bytes, is usually sufficient as shown below:

```
.stack 100h
```

Although there are other variations of the push and pop instruction, only the two simplest versions are introduced at the present time. As would be expected, the instructions push and pop are used to put data on top of the stack and remove data from the top of the stack, respectively. Note that 16, 32, or 64-bit registers and memory locations work with the push and pop instructions. However, if the al register needs to be pushed on the stack, then either the entire ax, eax, or rax register would need to be used. The format of these instructions can be found in Table 6.6, where obviously it is possible to push an immediate value onto a stack, but it is not possible to pop a value off the stack and put it into an immediate value.

The use of push and pop instructions is typically a good way to save and restore values. On some processors the use of a stack can be faster than using a temporary memory location. However, on the Intel processor it tends to be a little slower when saving and restoring a memory location but is about the same speed when saving and restoring a register. If there is not much of a difference and in some cases using the stack might be a little slower, what is the advantage of using the stack to save and restore values over a temporary memory location? The benefit of using the

Table 1.5: Push and pop instructions

Push instructions	Pop instructions
push reg	pop reg
push mem	pop mem
push imm	

stack is primarily convenience. Since memory for the stack has already been allocated using the .stack directive, extra temporary memory locations do not need to be declared. Further, the stack is always available and the names of various temporary memory locations do not need to be remembered. As indicated in the last section, it is often useful to use a stack to hold the original contents of a register or a memory location prior to manipulation of the bit pattern. Instead of moving the original bit pattern into a memory location, it can be pushed onto the stack prior to the loop and then restored to its original pattern after the loop as shown below. Although this might not be necessary when using a rotate instruction as mentioned in the last section, sometimes one might accidentally not rotate the register or the memory location the correct number of times. So saving or restoring the original pattern is inexpensive insurance and the push and pop instructions make this easy. For example, the following uses the shift instruction to demonstrate the use of the push and pop instructions. As mentioned previously, note that although only the al register needs to be pushed on and popped off the stack, the eax register is used here, where the ax register would have worked as well:

```
push eax
```



```

mov count,0 ; initialize count to zero
mov ecx,8 ; initialize loop counter to zero
.repeat
test al,00000001b ; test bit position zero
.if !ZERO? ; is the bit set?
inc count ; yes, count it
.endif
shr al,1 ; shift al right one bit position
.untilcxz
pop eax

```

If for some reason more than one register needs to be saved and restored, the order of the pushes and pops must be taken into consideration. For example, what if following arithmetic statement needs to be evaluated?

$$w = x / y - z;$$

Assuming that the previous contents of the registers used to evaluate the expression should not be altered, they would need to be saved and restored. Using the stack, Fig. 6.9 illustrates how this could be accomplished.

It should be noted that both the `eax` and `edx` registers are pushed onto the stack. Although it is obvious that the `eax` register is used by the `sub` and `mov` instructions, do not forget that the `cdq` instruction extends the sign bit through the `edx` register and that the `idiv` instruction leaves the remainder in the `edx` register. As a result, both registers need to be pushed onto the stack. As can be seen, the first register popped off the stack is the `edx` register. Recall that a stack is a LIFO structure, where the last item pushed onto the stack is the `edx` register, so it is the first item that needs to be popped off the stack.

A possible problem when using the stack extensively is that one might forget which items were pushed onto the stack and in what order. This can result in some difficulty in debugging logic errors. To help avoid these errors, one possibility is to avoid overusing the stack. Further, when using the stack it is a good idea to keep the associated pushes and pops relatively close to one another so that the connections between the two can easily be seen. It can be quite confusing to see a push in a middle of a code segment only to find the corresponding pop many pages further away in the code. Yet another technique that can help is to use scoping lines to help match push instructions with the corresponding pop instructions to insure that the item being popped off the stack is the correct one. Scoping lines can be drawn on a program listing by hand to help with creating or debugging code and are illustrated in the code segment in Fig. 6.9.

1.15 Swapping Using Registers, the Stack, and the `xchg` Instruction

As another example of using the stack, assume that two values need to be swapped, such as is done in a number of sorting algorithms. The typical high-level code




Figure 1.8: Saving and restoring multiple registers

is as follows:

```
temp = num1;
num1 = num2;
num2 = temp;
```

This could be implemented on a line-by-line basis in assembly language using registers, but that would require two instructions for each line of code as follows:

```
mov eax,num1
mov temp,eax
mov eax,num2
```

```

mov num1,eax
mov eax,temp
mov num2,eax

```

Clearly, the above seems inefficient. Instead, a register such as `edx` could be used instead of the temporary memory location `temp`. The `edx` register is chosen as the temporary register to make sure that the `ecx` register is free to be used for loop control and `ebx` is free to be used as an index register, both of which will be discussed further in Chap. 8. However, the middle high-level instruction would still need to use another register such as `eax` to enable the transfer between the two memory locations:

```

mov edx,num1
mov eax, num2
mov num1,eax
mov num2,edx

```

The above use of the registers can be rearranged as follows to help in readability:

```

mov eax,num1
mov edx,num2
mov num1,edx
mov num2,eax

```

Instead of using registers and `mov` instructions, another possibility is to use the stack. Not only is a stack a nice way to save and restore values, it can also be useful in swapping two values. An advantage over the `mov` instructions is that the stack does not need to use any of the general purpose registers which free them up for other uses. The following code segment swaps the values in `num1` and `num2`:

```

push num1
push num2
pop num1
pop num2

```

Is the order of the `pop` instructions above correct? Yes, since the purpose of the above code is not to save and restore the contents of `num1` and `num2` but rather to swap their contents. In other words, since the last value pushed onto the stack is the contents of `num2`, it is the first value popped off the stack. Instead of being popped back into `num2`, it is popped into `num1`. The same happens with the value originally in `num1`, thus swapping the two values.

Yet another method of swapping two values is to use the exchange (`xchg`) instruction. When swapping two registers, it is faster than the two methods previously presented. The format of this instruction is given in Table 6.7. Of course the one instruction that would directly allow the swapping of two memory locations is noticeably absent. Like many previous instructions, memory

Xchg instructions

```
xchg reg,reg  
xchg reg,mem  
xchg mem,reg
```

Table 1.6: Exchange instructions

to memory exchanges are not possible. However, the above does allow an exchange between two memory locations to occur in only three instructions, instead of the four needed in the previous three examples:

```
mov eax, num1  
xchg eax, num2  
mov num1, eax
```

This is accomplished by moving one of the two values into a register, swapping the register with the other memory location, and then moving the contents of the register back into the original memory location. Assuming that num1 originally contains a 5 and num 2 originally contains a 7, the three diagrams in Fig. 6.10 illustrate the three instructions in the above code segment.

With respect to register usage when swapping to memory locations, the stack does not use any general purpose registers, the xchg instruction uses one register and mov instructions require two registers. With respect to speed, using mov instructions is the fastest, the xchg instruction is just a little slower, and using the stack is the slowest. The result is that the xchg instruction is a nice compromise between the other two methods in terms of both register usage and speed. Also, given the convenience of the xchg instruction, it will usually be the method of choice in subsequent examples.

1.16 Complete Program: Simulating an OCR Machine

As alluded to in many of the preceding sections, computers use individual bits in a register or a memory location to indicate the status or to control various parts of the CPU or peripheral devices. One such machine might be an optical character

recognition (OCR) device that reads typed or handwritten characters from a piece of paper. On larger machines, they use a transport device that can handle more than a single piece of paper at one time, similar to a copying machine, where a memory location might be used to indicate the status of the paper in the transport. For the following simulation, the memory location used will be called the document status byte (DSB). There are a variety of problems that can happen with any sheet of paper in a transport as indicated in Table 6.8.

As can be seen, the bit number indicates the corresponding location of the error bit in the DSB, remembering that the low-order bit in a byte is located on


Figure 1.9: Swapping using the `xchg` instruction

Table 1.7: Simulated error messages

Bit	Message	Meaning
0	Short document	The document just read is shorter than anticipated
1	Long document	The document just read is longer than anticipated
2	Close feed	Current document is too close to the preceding document
3	Multiple feed	Two documents were detected at the same time
4	Excessive skew	The document is skewed (crooked) in the transport
5	Document misfeed	The document fails to feed into the transport
6	Document jam	The document jammed in the transport
7	Unspecified error	An unknown/unspecified error occurred

the right. It is also possible that more than one of the above conditions could occur at the same time. For example, if two documents are overlapped, it might cause both a multiple feed and a long document error condition.

To help create appropriate input to test the program, it would be inconvenient to input the bit patterns in decimal (base 10). Instead it would be easier to input in binary or hexadecimal, where fortunately C has the ability to input hexadecimal numbers. This is accomplished by using the letter *x* instead of the letter *d* in the format string as shown below. Further, in writing this program, one might be tempted to use nested if statements in order to make the code more readable and efficient. However, recall from above that it is possible to have more than one error condition and the use of nested if statements would rule out this possibility. It might be possible to use the equivalent of the switch statement but leave the break statements out of the code to allow more than one case to be tested. However, not using break statements in a switch statement might be considered by some to be unstructured. Further, given that there is no high-level directive equivalent of the case structure, it would have to be implemented using only low-level code. As has been seen previously, sometimes the code can get rather ugly using a lot of jump statements. Instead the approach taken here is the use of non-nested high-level if directives:

```

    INVOKE printf, ADDR msg1fmt, ADDR msg1
    INVOKE scanf, ADDR in1fmt, ADDR dsb
    INVOKE printf, ADDR msg2fmt, ADDR msg2, dsb
    .while dsb<=offh
    test dsb,000000001b
    .if !ZERO? ; if bit 0 = 1 then
    INVOKE printf, ADDR msglfmt, ADDR msgshort
    .endif
    test dsb,000000010b
    .if !ZERO? ; if bit 1 = 1 then
    INVOKE printf, ADDR msglfmt, ADDR msglong
    .endif
    test dsb,000000100b
    .if !ZERO? ; if bit 2 = 1 then
    INVOKE printf, ADDR msglfmt, ADDR msgclose
    .endif
    test dsb,000001000b
    .if !ZERO? ; if bit 3 = 1 then
    INVOKE printf, ADDR msglfmt, ADDR msgmult
    .endif
    test dsb,000010000b
    .if !ZERO? ; if bit 4 = 1 then
    INVOKE printf, ADDR msglfmt, ADDR msgskew
    .endif
    test dsb,000100000b
    .if !ZERO? ; if bit 5 = 1 then

```

```

        .686
        .model flat,c
        .stack 100h
scanf    PROTO arg2:Ptr Byte, inputlist:VARARG
printf   PROTO arg1:Ptr Byte, printlist:VARARG
        .data
msg1fmt  byte "%s",0
in1fmt   byte "%x",0
msg2fmt  byte "%s%x",0Ah,0Ah,0
msg1     byte 0Ah,"Enter a hexadecimal number: ",0
msg2     byte "The hexadecimal number is: ",0
msgshort byte "SHORT DOCUMENT",0Ah,0
msglong  byte "LONG DOCUMENT",0Ah,0
msgclose byte "CLOSE FEED",0Ah,0
msgmult  byte "MULTIPLE FEED",0Ah,0
msgskew  byte "EXCESSIVE SKEW",0Ah,0
msgfeed  byte "DOCUMENT MISFEED",0Ah,0
msgjam   byte "DOCUMENT JAM",0Ah,0
msgerror byte "UNSPECIFIED ERROR",0Ah,0
dsb      dword ?
        .code
main     proc

```

```

    INVOKE printf, ADDR msg1fmt,ADDR msgfeed
    .endif
    test dsb,01000000b
    .if !ZERO? ; if bit 6 = 1 then
    INVOKE printf, ADDR msg1fmt,ADDR msgjam
    .endif
    test dsb,10000000b
    .if !ZERO? ; if bit 7 = 1 then
    INVOKE printf, ADDR msg1fmt,ADDR msgerror
    .endif
    INVOKE printf, ADDR msg1fmt,ADDR msg1
    INVOKE scanf, ADDR in1fmt,ADDR dsb
    INVOKE printf, ADDR msg2fmt,ADDR msg2, dsb
    .endw
    ret
main endp
end

```

As can be seen in the while loop, any bit combination that is less than or equal to 0FFh, or 1111111b is allowed, where the *h* and *b* stand for hexadecimal and binary, respectively. Note that when a hex number begins with a letter, it has to be preceded with a 0 so that the assembler does not confuse it with a variable name. Once a number of 100 h or greater is entered, the loop stops, as can be seen in the sample input/output below.

1.17 Sample Input/Output

```

Enter a hexadecimal number: 1
The hexadecimal number is: 1
SHORT DOCUMENT
Enter a hexadecimal number: 2
The hexadecimal number is: 2
LONG DOCUMENT
Enter a hexadecimal number: 3
The hexadecimal number is: 3
SHORT DOCUMENT
LONG DOCUMENT
Enter a hexadecimal number: ff
The hexadecimal number is: ff
SHORT DOCUMENT
LONG DOCUMENT
CLOSE FEED
MULTIPLE FEED
EXCESSIVE SKEW
DOCUMENT MISFEED

```


DOCUMENT JAM
UNSPECIFIED ERROR
Enter a hexadecimal number: 100
The hexadecimal number is: 100
Press any key to continue . . .

1.18 Summary

- The inclusive-or includes the case when both operands are true and the result is true, whereas the exclusive-or excludes this case and the result is false when both operands are true.
- To set, clear or test, and toggle bits, use the or, and, and xor instructions, respectively.
- If data is needed later, be sure to save the data when using the shl and shr instructions.
- As a by-product of other tasks, a register or a memory location can be cleared to zero using the shl and shr instructions. However, the shift instructions can be slower than the mov or xor instructions and the latter two are usually a better choice.
- To multiply or divide by powers of two, use sal and sar, respectively, to communicate to others that arithmetic is being performed and to insure that negative numbers are handled properly with division.
- If a bit pattern is rotated exactly the same number of bits that are in a register or a memory location, then the bit pattern does not need to be saved and restored.
- When saving and restoring data using push and pop instructions, be sure to remember that the last one pushed on the stack should be the first one popped off the stack (LIFO).
- The use of scoping lines when using push and pop instructions can be helpful when creating or debugging code.
- Data in memory can be swapped using only mov instructions which use more registers and are faster compared to using the push and pop instructions which do not use any general purpose registers. Using the xch instruction along with the appropriate mov instructions is a good compromise in terms of register usage and speed.

1.19 Exercises (Items Marked with an * Have Solutions in Appendix D)

1. Given the following assembly language statements, indicate whether they are syntactically correct or incorrect. If incorrect, indicate what is wrong with the statement:
 - *A. or eax, ebx
 - B. xor al, ah
 - *C. rotate al, 1
 - D. shr ax, 2
 - *E. sar eax, 3
 - F. xchg dog, cat
 - G. ror ecx, 1
 - H. lol dx, 8
 - I. shift 2, ax
2. Given the following C arithmetic instructions, implement them using arithmetic shift instructions, where possible:
 - A. $\text{answer} = \text{num} - \text{total} / 32;$
 - *B. $\text{result} = (\text{amount} + \text{number}) * 4;$
 - C. $x = y * 8 + z / 2;$
 - D. $a = a / 16 - b * 6;$
3. Write a code segment that takes the contents of eax, ebx, ecx, and edx, and puts them in the reverse order of edx, ecx, ebx, and eax using only the push and pop instructions. In other words, eax should contain the contents of edx and vice versa, etc.
4. Assume that a status register in a processor indicates the current state of a photocopying machine according to the following table. For each bit, output an appropriate message indicating the status of the machine. Note that although there can be more than one bit set at one time, only one error message can be generated, where bit 0 has the highest priority, followed by bit 1, etc. At the discretion of the instructor, implement using high-level directives, without high-level directives, or a combination as shown in the text.

1.20 Bit Message

- 0 Paper jam
 - 1 Paper misfeed
 - 2 Paper tray empty
 - 3 Toner low
 - 4 Toner empty
5. Similar to the program in Sect. 6.8, write a program to simulate a security alarm system according to the following table, where it is possible that any of

the first three high-priority items could happen at the same time. Although the last three items can also occur at the same time, the program should check and output messages for them only when none of the higher priority first three items have occurred.

- 0 Fire alarm
- 1 Carbon monoxide
- 2 Power outage
- 3 Gate unlocked
- 4 Door open
- 5 Window open

1.21 Procedures and Macros

This chapter will first show the reader how procedures are implemented in assembly language. The implementation of macros is introduced next which is probably a new topic to most readers. Both procedures and macros are tools that allow programmers to save time by not having to rekey the same code over and over again, but there are important differences between the two mechanisms. The chapter then continues with the introduction of conditional assembly which can be a difficult concept for new assembly language programmers. Lastly, this chapter shows the beginning of the implementation of what might be called a macro calculator which simulates a one register (accumulator) computer.

1.22 Procedures

Most readers are probably familiar with procedures from a previous programming class. Depending on what language was used in that class, procedures may have also been called subprocedures, subprograms, subroutines, functions, or methods. The most generic of these terms is subprograms, which encompass all the others. Functions and many methods can or should return only a single value, whereas procedures, subprocedures, and subroutines are designed to return anywhere from zero to many values. In assembly language, subprograms are called procedures and belong to this last group. Although there are ways to make it possible to utilize parameters, the simplest way to communicate between a program and a procedure is to use either global variables or registers.

The instruction used to invoke a procedure is the call instruction. The call instruction has one operand that specifies the name of the procedure to be invoked. Upon return from the procedure, execution will continue with the instruction after the call instruction. An example is given below, where pname is a placeholder for the procedure name:

```
call pname
```

Although the actual procedure can be placed in a number of locations in the program, probably the most convenient place is after the `endp` statement in the main program and prior to the `end` statement. The first line of the procedure contains the name of the procedure in the label field, represented by the word `pname`, followed by the `proc` directive in the opcode field. Next comes the body of the procedure, followed by the return instruction `ret`, followed by the `endp` directive, which has the name of the procedure in the label field as shown below:

```
pname proc
    ; body of the procedure
    ret
pname endp
```

The `proc` and `endp` directives indicate to the assembler the beginning and the end of the procedure, respectively. The `ret` instruction, during the execution of the procedure, indicates when to return to the calling program. Unlike many high-level language, the `ret` instruction does not return a value to the calling program but rather indicates that the execution of the program should return to the calling program. One of the most common errors made by beginning assembly language programmers is forgetting to include the `ret` instruction, allowing the execution of the program to continue past the end of the procedure and possibly into another procedure following the current procedure. Although there can be more than one `ret` instruction in a procedure, like in many high-level languages, it is recommended to include only one return statement in a procedure. This helps to keep the program structured with only one entry point and one exit point. Further, almost any procedure can be rewritten to contain only one `ret` instruction. For example, given the following procedure with two `ret` instructions:

```
sample1 proc
    .if eax == 0
    mov edx,1
    ret
    .else
    mov edx,0
    ret
    .endif
sample1 endp
```

it can be rewritten to utilize only one `ret` instruction:

```
sample1 proc
    .if eax == 0
    mov edx,1
    .else
    mov edx, 0
    .endif
```

```
ret
sample1 endp
```

The result is cleaner code that is less prone to logic errors during modification. Also, it is usually best to be sure that the `ret` instruction is the last statement in a procedure prior to the `endp` directive. Can you determine what is wrong with the following procedure that is supposed to add all the registers together and return the value in `eax`?

```
; *** Caution: Contains a logic error ***
sample2 proc
    add eax, ebx
    add eax, ecx
    ret
    add eax, edx
sample2 endp
```

Yes, the value in the `edx` register is never added to the `eax` register, and `eax` only contains the sum of `eax`, `ebx`, and `ecx` upon return to the calling program. The `add eax, edx` instruction in the above procedure is sometimes referred to as "dead code," because although it takes up space in memory, it is never executed. In larger programs, whole sections of code might never be executed if the code is located incorrectly and it might create a difficult situation to debug. The correct procedure is given below:

```
; *** Note: Correctly implemented code ***
sample2 proc
    add eax, ebx
    add eax, ecx
    add eax, edx
    ret
sample2 endp
```

What if one wanted to implement the multiplication algorithm from Chap. 5 in two different locations? The code could be written twice in two different sections of the program, but instead of writing the code twice, it would be much easier to put the logic in a procedure and then call the procedure from two different locations in the main program:

```
call mult
.
call mult
```

Then after the main program, the code for the `mult` procedure could be written. Recall this algorithm from the end of Sect. 5.1 which used the `.` while directive. The variables `x` and `y` could still contain the two values to be multiplied, but instead of using the variables `i` and `ans`, the following procedure uses the `ecx` and `eax` registers, respectively, along with other minor changes:

```

mult proc
    mov eax,0 ; initialize eax to 0
    .if x != 0
        mov ecx,1 ; initialize i to 1
        .while ecx<=y
            add eax,x ; add x to eax
            inc ecx ; increment i by 1
        .endw
    .endif
    ret
mult endp

```

Is there a potential problem with the above procedure? Since `ecx` is being used as a temporary variable to implement the loop, the contents of `ecx` will be destroyed. If the main program is not using the `ecx` register, this would not be a problem. However, if `ecx` is being used to hold various values, such as a counter for another loop in the calling program, this routine could cause problems. It could be difficult to debug if a programmer did not know that `ecx` is being used by the procedure. Although it is fairly obvious in this case, it can be difficult to notice in larger procedures.

One solution is to document the procedure carefully and include a comment right at the beginning of the procedure indicating which registers are destroyed by the procedure to warn potential users of the procedure. The responsibility for saving the contents of the affected register then lies with the programmer of the calling program. Although for some small, seldom-used procedures, this effectively solves the problem, it is still possible that the programmer using the procedure might miss the warning. Further, if the procedure is going to be called many times, then the calling program needs to save and restore the affected registers many times, thus wasting memory. Also, the possibility of forgetting to save and restore the registers at some point in time is increased.

When writing procedures, it is often a good idea to have the procedure take the responsibility of saving and restoring any registers being destroyed. This saves memory, since there is only one copy of the code and also lessens the chance for error by the calling program. What is the best way to accomplish this task? Although a temporary variable could be used, this is an excellent situation to use the stack as discussed in Chap. 6. The following multiplication procedure includes the pushing and popping of the `ecx` register:

```

mult proc
    push ecx ; save ecx
    mov eax,0 ; initialize eax to 0
    .if x != 0
        mov ecx,1 ; initialize i to 1
        .while ecx<=y
            add eax,x ; add x to eax
            inc ecx ; increment i by 1

```

```

.endw
.endif
pop ecx ; restore ecx
ret
mult endp

```

Although the act of calling and returning a procedure is a little slower than straight line code, it does save memory because the code needs to be written only once. Of course, the memory saving is compounded as the size of the procedure and the number of calls increase.

1.23 Complete Program: Implementing the Power Function in a Procedure

To illustrate a complete example, consider the problem of calculating x^n from Sect. 5.6. Instead of having the code to calculate x^n in the main program, it could be placed in a procedure. The procedure can then be invoked more than one time from the main program without having to duplicate the code each time. For the sake of simplicity both in the C program and more importantly in the subsequent assembly language program, power is implemented as a procedure (void function) and x , n , and ans are implemented as global variables. In addition to outputting a message in the case of an error, the procedure also returns a -1 in the variable ans :

```

#include <stdio.h>
int x,n,ans;
int main() {
    void power();
    printf("%s","Enter x: ");
    scanf("%d",&x);
    printf("%s","Enter n: ");
    scanf("%d",&n);
    power();

    printf("\n%s%d\n\n","The answer is: ",ans);
    return 0;
}
void power() {
    int i;
    ans=-1;
    if(x<0 || n<0)
        printf("\n%s\n","Error: Negative x and/or y");
    else
        if(x==0 && n==0)
            printf("\n%s\n","Error: Undefined answer");

```

```

    else {
        i=1;
        ans=1;
        while(i<=n) {
            ans=ans*x;
            i++;
        }
    }
}

```

As mentioned previously, global variables are used for x , n , and ans both in the C program above and in the assembly language below. Since i is declared as a local variable in the C code above and is not needed in the main program, ecx is used as the loop control variable in the assembly language procedure below:

```

        .686
        .model flat,c
        .stack 100h
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
printf PROTO arg1:Ptr Byte, printlist:VARARG
        .data
in1fmt byte "%d",0
msg1fmt byte "%s",0
msg3fmt byte "%s%d",0Ah,0Ah,0
errfmt byte "%s",0Ah,0
errmsg1 byte 0Ah,"Error: Negative x and/or y",0
errmsg2 byte 0Ah,"Error: Undefined answer",0
msg1 byte "Enter x: ",0
msg2 byte "Enter n: ",0
msg3 byte 0Ah,"The answer is: ",0
x sdword ?
n sdword ?

ans sdword ?
        .code
main proc
    INVOKE printf, ADDR msg1fmt, ADDR msg1
    INVOKE scanf, ADDR in1fmt, ADDR x
    INVOKE printf, ADDR msg1fmt, ADDR msg2
    INVOKE scanf, ADDR in1fmt, ADDR n
    call power
    INVOKE printf, ADDR msg3fmt, ADDR msg3, ans
    ret
main endp
power proc
    push eax ; save registers
    push ecx

```


1.23. COMPLETE PROGRAM: IMPLEMENTING THE POWER FUNCTION IN A PROCEDURE⁶⁵

```
push edx
mov ans,-1 ; default value for ans
.if x<0 || n<0
INVOKE printf, ADDR errfmt, ADDR errmsg1
.else
.if x==0 && n==0
INVOKE printf, ADDR errfmt, ADDR errmsg2
.else
mov ecx,1 ; initialize ecx loop counter
mov ans,1 ; initialize ans
.while ecx <= n
mov eax,ans ; load eax with ans
imul x ; multiply eax by x
mov ans,eax ; store eax in ans
inc ecx ; increment ecx loop counter
.endw
.endif
.endif
pop edx ; restore registers
pop ecx
pop eax
ret
power endp
end
```

Could the assembly language procedure above use registers instead of global variables to communicate back and forth between the procedure and the main program? Yes, but in the procedure above, x and y are checked to see if they are negative in the `.if` directive. Recall from Chap. 4 that the default in high-level directives is unsigned data unless a memory location declared as `sdword` is used.

Also, `INVOKE` directives are being used in the procedure to output error messages and remember from Chap. 2 that they destroy the contents of the `eax`, `ecx`, and `edx` registers. The result is that for smaller and simpler procedures the use of registers is probably the preferred method, but in instances like this, the use of global variables might be the better choice.

Note that the `eax`, `ecx`, and `edx` registers are saved at the beginning and restored at the end of the procedure. This is done not only because of the `INVOKE` directives but also because even if the procedure did not perform any output, the three registers should be saved and restored. It is obvious that `eax` is used in the `mov` instructions and the contents of `ecx` are destroyed when it is used for loop control. However, the `edx` register does not appear in the procedure, so why should it be saved and restored? Again, look carefully at the code and recall what happens with the `imul` instruction. The `imul` instruction extends the sign of the `eax` register into the `edx` register and destroys the contents of `edx`, so it should be saved and restored also. Even if the main program that called

the procedure does not use the `eax`, `ecx` and `edx` registers, the procedure should save and restore them so that the procedure could easily be used by other programs that might use these registers. Lastly, as discussed in Chap. 6, be careful to insure that the `pop` instructions are in the correct order to properly restore the three registers.

1.24 Saving and Restoring Registers

Note that if registers are not being used to communicate back to the main program, it is possible to save and restore all the registers, whether they were altered or not. Although this might be an easy way to avoid having to think about whether a register is altered, it is a sloppy method and does not help other programmers understand what is happening in the procedure. In other words, even though it might be simpler for the person writing the code, it is not necessarily easier for subsequent people reading and modifying the code. By saving and restoring only the registers that are altered, it helps others understand which registers are being altered and also helps make the code more self-documenting. Of course whether code appears to be self-documenting or not, documentation is always a good idea to supplement any code written.

However, what if a routine was altering all the registers and does not use a register to communicate back from a procedure to the main program? It could get a little messy trying to push and pop all the registers and further, there could be a chance for a logic error if the `pop` instructions were accidentally written in the wrong order. Luckily it is possible to save the four general purpose registers (`eax`, `ebx`, `ecx`, and `edx`) along with the `esi`, `edi`, `ebp`, and `esp` registers with only one instruction. The `pushad` instruction pushes contents of all the above registers onto the stack and the `popad` instruction subsequently pops the values from the stack and puts them back into their respective register locations.

As an example, consider a procedure that outputs blank lines a variable number of times. It would be nice that it does not destroy the contents of any registers and also not use any global variables, both of which would make it a very portable procedure that could be used in many different programs. As is known, using the `INVOKE` directive to output the blank lines can cause the contents of the `eax`, `ecx`, and `edx` registers to be altered. Since global variables will not be used, the `ebx` register could be used to communicate to the procedure how many blank lines need to be output. At the same time, since `ebx` would not be destroyed by the `INVOKE` directive, it would also make a good candidate for a loop control variable. Given these circumstances in this example, all four registers should probably be saved and restored. In the first example, each of the four general purpose registers are saved and restored individually:

```
blankln proc
    push eax
    push ebx
    push ecx
```

```
    push edx
    .repeat
    INVOKE printf, ADDR blnkfmt
    dec ebx
    .until ebx<=0
pop edx
pop ecx
pop ebx
pop eax
ret
blankln endp
```

Note the order of the push and pop instructions above to properly save and restore the contents of the four registers. In the example below, the four general purpose registers are saved all at once using the pushad and popad instructions:

```
blankln proc
    pushad
    .repeat
    INVOKE printf, ADDR blnkfmt
    dec ebx
    .until ebx<=0
    popad
    ret
blankln endp
```

As can be seen, clearly the second procedure is much cleaner than the first procedure. Since the above procedures do not return a value via a register and all four general purpose registers need to be saved and restored, this is a good example of when the pushad and popad instructions should be used. However, since most of the time the majority of the procedures in this text will be returning a value, the previously described method of saving and restoring registers individually will be used more frequently.

1.25 Macros

Another method to avoid having to write the same code again and again is the macro. However, be forewarned that even though it executes faster than a procedure, it tends to waste memory. Although most readers have probably not encountered macros in previous programming classes, they might have encountered the term in working with application software packages. Probably the most common occurrence of macros is in spreadsheet packages, where one can record a macro that consists of a series of steps performed by the user. Although a macro in assembler can contain a series of instructions, it is different from a macro in a spreadsheet, where the instructions are not recorded but rather need to be written by the programmer.

Like procedures, macros can be declared in many different places. Whereas procedures are usually declared after the main program, macros are usually written and located prior to the main program, just after the `.code` directive. This first writing of the macro is sometimes called the macro definition and does not take up any memory in the executable program. Further, a macro has a similar structure to a procedure, where on the first line the name of the macro is in the label field, indicated by `mname` below, and the macro directive is in the opcode field. The body of the macro follows, which is then followed by the `endm` directive which unlike an `endp` does not repeat the name of the macro in the label field as shown below:

```
mname macro
    ; body of the macro
endm
```

However, a question here might be, should there be a `ret` instruction just prior to the `endm` directive above? The answer is no, because there is not a set of instructions in a common place that are branched to and executed as with a procedure. Rather, a set of instructions exist only in the source code file (`.asm`), a copy of the instructions is inserted into the listing file (`.lst`), and the machine language equivalent is inserted into the execution file (`.exe`), wherever the macro is invoked.

Instead of a call instruction as with a procedure, a macro invocation is done by just using the name of the macro, `mname` in this instance, in the opcode field of the invoking program as demonstrated below:

```
.
mname
.
.
mname
.
```

Thus unlike a procedure, where the flow of executions jumps to the procedure and then returns back to the calling program, a copy of the macro is inserted into the program at each point where the macro is invoked. The result is that macros are faster because there is no calling or returning. Further, it is possible that if the macro is never invoked, it will never take up any memory because the code exists only in the source code file (`.asm`). However, usually macros tend to take up more memory due to the copying of the instructions at every location it is invoked and this is especially true when macros are large and/or invoked many times.

Consider an example of swapping two memory locations `num1` and `num2`, where a macro could be written as follows:

```
swap macro
    mov ebx,num1 ; copy num1 into ebx
```

```
xchg ebx,num2 ; exchange ebx and num2
mov num1,ebx ; copy ebx into num1
endm
```

As mentioned above, notice the lack of a `ret` instruction in the macro declaration. The invoking of a macro is done by just specifying the macro name in the opcode field in the calling program, as shown below, where it can be assumed that the programmer wanted to swap `num1` and `num2`, and then turn around and swap them back to their original locations:

```
.
swap
.
.
swap
```

A common mistake made by beginning assembly language programmers is that they include the `call` instruction when trying to invoke the macro, where this should be avoided and would cause a syntax error. A question that many readers might have at this point is: Why haven't `push` and `pop` been included in the macro

definition above? Couldn't there be the same problems with the calling program as with procedures, since the contents of the `ebx` register are being destroyed by the macro? The answer to the second question is yes, the same problem still exists. But in answer to the first question, the reason why the `push` and `pop` instructions have not been included is because they take up memory. One must remember that every time a macro is invoked, another copy of the macro is inserted into the code in the invoking program. As a result, many times the contents of registers are not saved in macros in order to save memory and a programmer needs to be extra careful whenever invoking a macro. This can be especially confusing, because unlike the procedure which is invoked via the `call` instruction, a macro is invoked only by using the name of the macro. Many times the name of a macro almost appears to be like the name of an instruction, so the user of the macro might be lulled into believing it is an instruction and forget about the hidden instructions in the macro that might destroy the contents of the registers. Many a programmer has accidentally made this mistake and spent much time trying to subsequently debug a program.

To help illustrate how macros take up memory and remind programmers that they are not instructions, it is sometimes useful to examine the assembly listing file (`.lst`) to see what is known as the macro expansion and all the instructions from the macro that are inserted into the code. For example the above calling program that invokes the macro `swap` twice would look as follows in the `.lst` file:

In the above listing, both the relative memory address and machine language equivalent in hexadecimal can be seen to the left. However, since the `.lst` file can sometimes get rather messy and have a cluttered appearance when macros are expanded, the addresses and machine code have been removed to

```

swap
00000000 8B 1D 00000046 R 1 mov ebx,num1 ; copy num1 into ebx
00000006 87 1D 0000004A R 1 xchg ebx,num2 ; exchange ebx and num2
0000000C 89 1D 00000046 R 1 mov num1,ebx ; copy ebx into num1
swap
00000012 8B 1D 00000046 R 1 mov ebx,num1 ; copy num1 into ebx
00000018 87 1D 0000004A R 1 xchg ebx,num2 ; exchange ebx and num2
0000001E 89 1D 00000046 R 1 mov num1,ebx ; copy ebx into num1

```

make this example easier to read in both the listing below and many subsequent listings (for more information on machine language, see Chap. 12):

```

swap
    mov ebx,num1 ; copy num1 into ebx
    xchg ebx,num2 ; exchange ebx and num2
    mov num1,ebx ; copy ebx into num1
swap
    mov ebx,num1 ; copy num1 into ebx
    xchg ebx,num2 ; exchange ebx and num2
    mov num1,ebx ; copy ebx into num1

```

A nice feature that should be noticed above is that any comments placed in the macro definition also appear in the macro expansion. Although the macro expansion might cause the program to appear more cluttered and also waste more paper when printing out the listing of the program, the possibility of avoiding errors during program development might well be worth it. It can also be especially helpful during the debugging process when trying to track down pesky logic errors.

A very useful feature when using macros is the ability to use arguments and parameters. Recall from high-level languages that the calling program sends arguments to procedures which correspond to parameters in the procedure. However, parameters in macros are different from many of the parameters that one may have encountered in various high-level languages. Depending on which language the reader has used previously, it should be recalled that reference and value parameters are used in C++ and that only value parameters are used in Java. Reference parameters refer to their corresponding arguments via an address and value parameters copy the values from the corresponding arguments. If the reader has had an upper-level course in programming languages, then name parameters might be familiar. Although name parameters are not used by very many modern programming languages, they were used in the past in languages such as Algol in the 1960s. For those who have not encountered them before, name parameters are essentially substitution parameters, where the names of the arguments are merely substituted in place of the parameter names.

For example, in the previous swap macro, what if one wanted to swap the contents of any two memory locations instead of just num1 and num2? The

swap macro could be rewritten as follows, where p1 and p2 are the two name parameters:

```
swap macro p1,p2
mov ebx,p1 ; copy p1 into ebx
xchg ebx,p2 ; exchange ebx and p2
mov p1,ebx ; copy ebx into p1
endm
```

Now, the above macro would work with any two memory locations as arguments. For example,

```
swap num1,num2
```

```
•
.
```

```
    swap x,y
```

Although at first glance the macro invocations look somewhat like memory to memory instructions, they are not. In order to understand how the above macros work it is best to look at the macro expansions. In the first swap, the code would look similar to the previous example without parameters because the argument names are the same as the previously used variable names. However, the second swap looks different because different memory locations are used in the arguments:

```
swap num1,num2
    mov ebx,num1 ; copy p1 into ebx
    xchg ebx,num2 ; exchange ebx and p2
    mov num1,ebx ; copy ebx into p1
swap x,y
    mov ebx,x ; copy p1 into ebx
    xchg ebx,y ; exchange ebx and p2
    mov x,ebx ; copy ebx into p1
```

The above example shows the versatility of parameters, where different arguments can be used with each invocation. Although the comments from the macro definition are included in the expansion, unfortunately the comments refer to the parameter names rather than the argument names. Should one want to document the macro definition using comments, but not see the comments in the macro expansion, double semicolons (; ;) should be used prior to the comment in the macro definition instead of the single semicolons (;) as used above.

What if a programmer left one or both arguments blank when trying to invoke the swap macro? With this particular macro it is not much of a problem, because in this instance a syntax error would occur from the instructions themselves in the macro that are missing a required operand. However, there are a few instructions, such as the imul instruction, that have optional operands and

not requiring an argument might cause a syntax error to be missed. As a result, it is good programming practice to indicate whether or not the arguments are required. This can be accomplished by using the `: REQ` statement in the parameter list. Should a required argument not be included, a syntax error would be generated regardless of the instructions used in the macro. To require both arguments in the swap macro, the resulting macro definition would look as follows:

```
swap macro p1:REQ,p2:REQ
mov ebx,p1 ;; copy p1 into ebx
xchg ebx,p2 ;; exchange ebx and p2
mov p1,ebx ;; copy ebx into p1
endm
```

Note also the comments above have been changed to use double semicolons (`;;`) so that they do not appear in subsequent macro expansions. Although the problem of a missing argument is solved, what would happen if an incorrect argument was used? For example, what would happen if an immediate value was inadvertently used as one of the arguments instead, as in `swap num,1`? The result is that this would cause a syntax error in the second instruction of the macro expansion in the `.lst` file, because immediate values cannot be exchanged as illustrated below:

```
swap num1,1
    mov ebx,num1
    xchg ebx,1
error A2070: invalid instruction operands
    mov num1,ebx
```

In another potential problem, what if registers were used instead of memory locations? Would this cause a problem, especially if one of the registers was `ebx` as in the second example below?

```
swap eax, ecx
•
-
swap ebx, ecx
```

The above invocations would generate the following macro expansions. As an aside, note that comments are not generated in the macro expansion due to the use of the double semicolons (`;;`) in the previous macro definition:

```
swap eax,ecx
    mov ebx,eax
    xchg ebx,ecx
    mov eax,ebx
-
swap ebx,ecx
```



```
mov ebx,ebx
xchg ebx,ecx
mov ebx,ebx
```

Although both of the above work, they are very redundant. In the first example, the value of `eax` is placed into `ebx`, then swapped with `ecx`, and then the value originally in `ecx` is placed into `eax`, where just a simple `xchg eax, ecx` would have sufficed instead. The second example is even more redundant where the value of `ebx` is moved into itself, then `ebx` is exchanged with `ecx`, and then `ebx` is moved again back into `ebx`. Although both of these expansions are redundant, they are syntactically correct and logically harmless. But what would happen if the two registers in the last example were reversed, as in `swap ecx, ebx`? The answer can be found in the following macro expansion:

```
swap ecx,ebx
mov ebx,ecx
xchg ebx,ebx
mov ecx,ebx
```

As can be seen, the contents of `ebx` are wiped out by the contents of `ecx` in the first line of the macro expansion, then `ebx` is swapped with itself, and then `ecx` is reloaded with the results that were originally in `ecx`. The result is that both `ebx` and `ecx` would now contain the contents of `ecx`. Would this produce an error message? No, this is not a syntax or an execution error but rather a logic error and points to a problem with using parameters with macros. As a result, where some problems produced syntax errors and some produced redundant code, this last one is the most serious. Programmers must be very careful when using macros and be sure to understand how they work before invoking them. It might be argued that a mistake like this is solely the responsibility of the programmer using the macro and let the user beware. However, can the problem be fixed? Yes, where a possible solution to this problem will be addressed in Sect. 7.6.

1.26 Conditional Assembly

Conditional assembly can be a confusing topic to beginning assembly language programmers. It uses what looks like `if` statements so that it seems like it is altering the flow of control during the execution of the program, but it is not the same as the selection structures learned in Chap. 4. Instead, the key to understanding conditional assembly is from its name, where it is "conditional assembly," not "conditional execution." Specifically, conditional assembly controls the assembler, not the flow of execution. Instead of having one or two possible routes for the execution to follow as with an `.if` statement, conditional assembly tells the assembler whether to put in a possible instruction or a possible set of instructions into the program as opposed to other possible instructions or no instructions at all.

There are a number of ways that conditional assembly can work, and this section will look at a few of the more commonly used methods. Table 7.1 lists the conditional assembly directives used in this and subsequent sections. Probably the best way to illustrate the concept and the directives is through an example.

Although somewhat simplistic, the first method to be examined is whether or not there is an argument in a macro invocation. In other words, instead of causing an error from a `:REQ` or subsequent instruction in the macro expansion, whether caused by an intentional or an unintentional missing argument, alternative code can be generated.

For example, suppose one wanted to create a macro called `addacc`, that when invoked without an argument, the default is to add the number 1 to the `eax` register. Since the most efficient way to do that is to use the `inc` instruction, that is

Table 7.1 Conditional assembly directives

Directive	Meaning
<code>if</code>	If (can use EQ, NE, LT, LE, GT, GE, OR, AND)
<code>ifb</code>	If blank
<code>ifnb</code>	If not blank
<code>ifidn</code>	If identical
<code>ifidni</code>	If identical case insensitive
<code>ifdif</code>	If different
<code>ifdifi</code>	If different case insensitive

what the macro will use. However, when a number, a register, or a memory location is used as an argument, the macro will then add that number, register, or memory location to the `eax` register. Clearly the `inc` instruction would not work in this second instance, so an `add` instruction must be used. How would the conditional assembly work in this example? If the argument is blank, then the `inc` instruction would be inserted into the code, otherwise the `add` instruction with the appropriate argument would be inserted into the code. The macro would be written as follows:

```
addacc macro parm
    ifb <parm>
        inc eax
    else
        add eax,parm
    endif
endm
```

Note that `:REQ` is not used in the parameter list, because a blank argument is one of the options. The first directive in the macro is `ifb`, which stands for "if blank." Again, it is not an instruction, nor is it like one of the directives from Chap. 4 which generate the instructions like `cmp` and `je`. Rather as the assembler is inputting statement after statement from the assembly language source file (`.asm`), it checks to see if the argument from the macro invocation statement

is blank. If it is blank, it inserts the `inc eax` instruction into the `.lst` file and the equivalent machine language instruction into the `.exe` file. Otherwise the `mov eax, parm` instruction after the `else` directive is inserted with the corresponding argument in place of the parameter `parm`. Lastly, the `endif` indicates the end of the `ifb`. It should be carefully noted that unlike the selection statements introduced in Chap. 4, none of these three directives have a decimal prior to the directive.

The following contains four different invocations of the preceding macro definition:

```
addacc
addacc 5
addacc edx
addacc num
```

The complete resulting code with the macro expansions is given below to show how it would actually look in the `.lst` file. Although it appears that there are a number of instructions, the only real lines of executable code that are generated are the ones that have addresses and machine code off to the left in hexadecimal (again see Chap. 12):

The above code segment does save a little memory under some circumstances because as mentioned back in Chap. 3 and will be demonstrated in Chap. 12, the `inc` instruction takes up less memory. Although it is getting a little ahead here in the text, it is interesting to point out that depending on the argument, a different machine language instruction is generated based on the operand. To illustrate, the `inc eax` instructions generates at relative memory location `0000000A` a hexadecimal `40` machine language instruction that is only 1 byte long and the `add eax, num` instruction generates at relative memory location `00000010` a hexadecimal `030500000003` A machine language instruction that is 6 bytes long (see Chap. 12). As before, since the above is rather cluttered, by eliminating the other lines of source code, the following shows a copy of each macro invocation followed by only the assembly instruction that would be generated and executed:

```
addacc
    inc eax
addacc 5
    add eax,5
addacc edx
    add eax,edx
addacc num
    add eax,num
```

Although the above cleaned up code segment seems relatively simple, it illustrates how different code can be used in place of other code when using conditional assembly, even though the source code containing the macro definition

	addacc
	1 ifb <>
0000000A 40	1 inc eax
	1 else
	1 add eax,
	1 endif
	addacc 5
	1 ifb <5>
	1 inc eax
	1 else
0000000B 83 C0 05	1 add eax,5
	1 endif
	addacc edx
	1 ifb <edx>
	1 inc eax
	1 else
0000000E 03 C2	1 add eax,edx
	1 endif
	addacc num
	1 ifb <num>
	1 inc eax
	1 else
00000010 03 05 0000003A R	1 add eax,num
	1 endif

looks as though there might be more instructions. Again, there is only one actual assembly language instruction generated for each invocation in the above example.

1.27 Swap Macro Revisited Using Conditional Assembly

Returning to the swap example in Sect. 7.4, what would happen if invocations such as the ones on the left in Table 7.2 were used to invoke the previous definition repeated on the right?

The result would be that the following code would be generated:

Table 7.2 Macro invocations and definition

Invocations	Definition
<pre>swap num1, num1 swap eax, eax</pre>	<pre>swap macro p1:REQ, p2:REQ mov ebx, p1 ;; copy p1 into ebx xchg ebx, p2 ;; exchange ebx and p2 mov p1, ebx ;; copy ebx into p1 endm</pre>

```
swap num1, num1
    mov ebx, num1
    xchg ebx, num1
    mov num1, ebx
swap eax, eax
    mov ebx, eax
    xchg ebx, eax
    mov eax, ebx
```

As has happened before, redundant code is generated in both cases, but is there a possible solution to this problem? The answer is yes and it can be solved by using conditional assembly. As per Table 7.1, an `ifidn` (if identical) checks to see if the two arguments are equal using case sensitivity and an `ifidni` does the same thing but is case insensitive. The directive `ifdif` (if different) uses case sensitivity to check to see if the arguments are different and `ifdif` does the same thing with case insensitivity.

The statement that can help in this instance is `ifidni` directive. For example, if the arguments are the same, then there is no need to swap the contents and thus no code needs to be generated as shown in the following macro definition:

```
swap macro p1:REQ, p2:REQ
ifdif <p1>, <p2>
    mov ebx, p1
    xchg ebx, p2
    mov p1, ebx
endif
endm
```

The invocation of `swap` using various different scenarios is shown below:

```
swap num1, num2
swap num1, num1
swap eax, ecx
swap eax, eax
```

The resulting macro expansions are as follows:

```
swap num1, num2
    mov ebx, num1
    xchg ebx, num2
```

```

    mov num1, ebx
swap num1,num1
swap eax,ecx
    mov ebx, eax
    xchg ebx,ecx
    mov eax, ebx
swap eax,eax

```

Note that code is not generated in the second and fourth examples because of the `ifidni` directive and thus the redundant code need not be generated. The swapping of two memory locations is carried out easily in the first example and although the third example is still redundant, there is no harm done. As can be seen, this helps clean up much of the redundant code but still allows for some redundant code as in the third case above. However, what about the more serious problem at the end of Sect. 7.4 when the `ebx` register was used as the second argument resulting in a logic error?

In addition to seeing if the two arguments are the same or different, the above directives can be used to see if the parameters are equal to a particular register. In other words, by putting a particular register in brackets in one of the two positions of the `ifidni`, and the name of the parameter in brackets in the other position, it can then compare the parameter to the particular register and generate the appropriate code.

```

swap macro p1:REQ,p2:REQ
    ifidni <ebx>,<p2>
        xchg p1,ebx
    else
        mov ebx,p1
        xchg ebx,p2
        mov p1,ebx
    endif
endm

```

In the code above, `<p2>` is compared to see if it is identical to `<ebx>`, and if so, different code can be generated, otherwise the original code is executed in the `else` section. The following invocations

-
- swap ebx,eax
-
- swap eax,ebx

show how the code is different when the `ebx` is in either of the argument positions:

```

swap ebx,eax
    mov ebx,ebx

```

```

    xchg ebx,eax
    mov ebx,ebx
swap eax,ebx
    xchg eax,ebx

```

Again, some of the code is still redundant, but the danger of the logic error has now been eliminated. Can this last bit of redundancy be eliminated? Yes, by nesting another set of if, else, and endif directives after the outer else directive. Even an elseif directive could be used as follows:

```

swap macro p1:REQ,p2:REQ
    ifidni <ebx>,<p2>
        xchg p1,ebx
    elseifidni <p1>,<ebx>
        xchg ebx,p2
    else
        mov ebx,p1
        xchg ebx,p2
        mov p1,ebx
    endif
endm

```

Using the same two swap invocations from above, the following code is generated:

```

swap ebx,eax
    xchg ebx,eax
swap eax,ebx
    xchg eax,ebx

```

If a macro is going to be used only a few times by knowledgeable programmers, does one need to create such an elaborate macro using conditional assembly? The answer is probably no, because the time and effort needed to create the macro is most likely not worth the savings in terms of memory and execution speed. However, if the macro will be used a lot and by programmers with a variety of different skill levels, then the savings not only in terms of memory and execution speed might be worth it, but also the savings in terms of not having to correct syntax, redundancy, and logic errors may prove to be well worth the additional effort too.

1.28 Power Function Macro Using Conditional Assembly

In another example, what if one wanted to implement the power function in a macro? Recall from Chap. 6 the definition of the power function, which is reiterated below. For the sake of convenience, instead of outputting error messages, only a flag such as a -1 could be returned to indicate error had occurred:

```

x
  Else if  $x=0$  and  $n=0$  then -1
  Else if  $n=0$  , then 1
  Otherwise  $1*x*x*\ldots*x$  (  $n$  times)

```

Although the above definition would ideally be better implemented as a procedure because only one copy of it would be needed, it is a good example that can help illustrate some more important concepts concerning macros and conditional assembly. Of course, a macro could be written without conditional assembly, but then every time the code was generated, both the selection statements and loop would need to be inserted whether they were needed or not, which would be a waste of memory.

However, by using conditional assembly, not all of the if statements need to be included every time the macro is invoked. Further, in some of the cases a loop is not needed and the code for the loop would not need to be generated. For example, when n is 0 , then the answer is 1 , when n is 1 the answer is x , and when x is either 0 or 1, the answer is x . Thus, the above definition can be further revised as follows to better reflect how the macro could be written. Only when both x and n are greater than 1 would the code for the loop need to be generated. Below is the modified definition used in this example:

```

 $x^n =$  $ If  $x < 0$  or  $n < 0$ , then -1
  Else if  $x = 0$  and  $n = 0$ , then -1
  Else if  $x = 0$  or  $x = 1$ , then  $x$ 
  Else if  $n = 0$ , then 1
  Else if  $n = 1$ , then  $x$  ,
  Otherwise  $1 * x * x * \ldots * x$  (  $n$  times)

```

In addition to the directives used previously (ifb, inb, idif, etc.), it is possible to use just a simple if directive that can use the equivalent of relationals (eq, ne, lt, gt, le, ge) and logic (and, or) as shown in Table 7.1. The implementation of the above definition can then be accomplished as below. However, note that the simple if directive cannot contain memory locations or registers, but only constants. Also, it should be noted to use parentheses when using the logical operators as shown below. For this example, a constant for x will be passed as the first argument and a constant for the exponent n as the second argument, and the answer returned in the eax register:

```

power macro x:REQ,n:REQ
if (x lt 0) or (n lt 0)
mov eax,-1
elseif (x eq 0) and (n eq 0)
mov eax,-1
elseif (x eq 0) or (x eq 1)
mov eax,x
elseif n eq 0
mov eax,1

```



```

elseif n eq 1
mov eax,x
else
mov eax,x
mov ebx,eax
mov ecx,n
dec ecx
.repeat
imul ebx
.untilcxz
endif
endm

```

When either x or n is less than 0, or if both x and n are equal to 0, the `eax` register is set to -1. Should x be equal to either 0 or 1, the result is x , if only n is 0, the answer is 1, or if n is 1, then the answer is x , where in all of these cases the loop does not need to be executed nor does it even need to be generated. Lastly, notice that the loop is implemented using a `.repeat-.untilcxz` directive instead of a while loop. Since the case when n is equal to one is already handled, the loop only needs to

iterate $n - 1$ times. Although the above code segment looks rather large, remember that not all of the code is used in each invocation. Only the last case is the largest and since it is generated only when x and n are greater than 1, it does not need to be generated for all of the other cases. Note that all the registers are used, so care must be used when invoking this macro. The following sample invocations test seven different cases:

```

power 2,-1
power 0,0
power 0,2
power 1,2
power 2,0
power 3,1
power 2,3

```

Given the invocations above, the following macro expansions would be generated in each of the seven cases. Note that it is in only the last case that the loop is actually generated, where the `.repeat-.untilcxz` directive is implemented as a loop instruction:

```

power 2,-1
    mov eax,-1
power 0,0
    mov eax,-1
power 0,2
    mov eax,0
power 1,2
    mov eax,1

```

```
power 2,0
    mov eax,1
power 3,1
    mov eax,3
power 2,3
    mov eax,2
    mov ebx,eax
    mov ecx,3
    dec ecx
@C0001:
    imul ebx
    loop @C0001
```

1.29 Complete Program: Implementing a Macro Calculator

A nice way to illustrate the use of macros is to create what could be thought of as a macro calculator that simulates a one register (accumulator) computer. In a sense, the macros created and subsequently invoked almost appear to be a new set of instructions that can be used by the programmer. Although it somewhat looks like a new assembly language has been created and it is even sometimes mistakenly called an assembler, it is not really a new assembly language because that would require a separate program to assemble the instructions into the corresponding machine language. Instead, macros are used to create what looks like new instructions for the programmer and those macros in turn use existing assembly language instructions. Even though it is not really an assembler, it is still interesting to invoke the macros that appear to be like instructions that come from another hypothetical assembly language.

To implement this macro calculator, it is assumed that there is only one register called the accumulator. The `eax` register naturally can assume the role of this accumulator. As far as the macro calculator instruction set is concerned, there are no other registers. However, that does not mean that the other registers cannot be used on occasion as necessary to implement some of the other instructions. In a sense, the other registers are hidden from the macro calculator programmer. For those who are taking or have had a computer organization course, this is not unlike many processors that have registers that are not directly accessible by the programmers, such as the MAR (memory address register) and MDR (memory data register) in the CPU. The instructions for this macro calculator are given in Table 7.3.

Table 7.3 Macro calculator instructions

Instruction	Implemented as	Description
INACC	proc	Prompt for and input an integer into the accumulator
OUTACC	proc	Output message and integer in the accumulator
LOADACC	macro	Load the accumulator with the operand
STOREACC	macro	Store accumulator in the operand
ADDACC	macro	Add operand to the accumulator
SUBACC	macro	Subtract operand from the accumulator
MULTACC	macro	Multiply accumulator by the operand (iterative)
DIVACC	macro	Divide accumulator by the operand (iterative)

As mentioned above, the only register that should be modified is the `eax` register, which serves as the accumulator for the macro calculator. When implementing the various macros, care should be taken not to alter any of the other registers needlessly. For example, with the `MULTACC` macro, it will be necessary to alter at least one other register. As stated previously, macros typically do not save and restore the registers because that can take up time and memory. However, the purpose of this program is not necessarily to be efficient but rather to simulate a one accumulator machine where the only register that should be altered is `eax` and this will provide further practice in using the stack.

As mentioned previously in Chap. 6, a possible solution to saving and restoring registers is to use the `pushad` and `popad` instructions, but many of the above macros return a value via the `eax` register. Also, many of the macros will only alter a single register, so this method would almost be overkill. Another solution is that instead of trying to determine which registers are indeed altered, just save and restore all the other registers not being used to return values to insure that none of them will accidentally be altered. Although this solution would work and is sometimes employed by some programmers, it is again overkill and also a sloppy solution. Instead, it is best to save and restore only those registers that are indeed altered, which helps other programmers examining the macros to understand the code and also helps cut down on the number of instructions necessary to implement the macro when saving and restoring registers.

In looking at a few of the above macros, the `LOADACC` macro simply loads the contents of the specified memory location into the accumulator. Obviously, this is like the `mov` instruction and it forms the body of the macro. The `ADDACC` macro is similarly just the implementation of the `add` instruction. In both cases, no other registers are altered, so there is no reason to save and restore any other registers. The implementation of each can be found below:

```
LOADACC macro operand
mov eax,operand
endm
ADDACC macro operand
add eax,operand
endm
```

The `MULTACC` macro could obviously be implemented using an `imul` instruction, but for the sake of practice the iteration method will be used. Again,

since there is more than one instruction, it might be better to have it implemented as a procedure, but to allow for additional practice, a macro will be used. The algorithm presented here is somewhat similar to the algorithm presented in Sect. 5.1; however, the one presented here stresses some different concepts. Although there are some inefficiencies in the following algorithm, a more interesting concern here is the multiplying by a negative number which will allow another demonstration of conditional assembly.

If the multiplier in the operand is positive, there is no problem, because the loop will repetitively add the value of the multiplicand in the accumulator (eax) and whether the value in the accumulator is positive or negative does not matter. But if the multiplier in the operand is negative, it needs to be made positive in order to loop the correct number of times. Then if the value of the multiplicand in the accumulator is positive, the answer will need to be made negative and if the value of the multiplicand in the accumulator is negative, the answer will need to be made positive, because a negative number multiplied by a negative is positive. Table 7.4 illustrates these four possibilities.

Using the simple conditional if directive, the following can use only immediate integer values for the operand:

Table 7.4 Four possibilities

Accumulator	Operand	Iteration	Answer	Answer corrected
2	3	$2 + 2 + 2$	$= 6$	
-2	3	$-2 + -2 + -2$	$= -6$	
2	-3	$2 + 2 + 2$	$= 6$	$-(6) = -6$
-2	-3	$-2 + -2 + -2$	$= -6$	$-(-6) = 6$

```

MULTACC macro operand
push ebx ;; save ebx and ecx
push ecx
mov ebx,eax ;; mov eax to ebx
mov eax,0 ;; clear accumulator to zero
mov ecx,operand ;; load ecx with operand
if operand LT 0 ;; if operand is negative
neg ecx ; make ecx positive for loop
endif
.while ecx >0
add eax,ebx ;; repetitively add
dec ecx ;; decrement ecx
.endw
if operand LT 0 ;; if operand is negative
neg eax ; negate accumulator, eax
endif
pop ecx ;; restore ecx and ebx
pop ebx
endm

```

Since the `eax` register is serving as the accumulator and it contains the value that needs to be returned as a result of the multiplication operation, note that the `ebx` and `ecx` registers are obviously saved and restored and that the `eax` register is not. Also notice the negation of `ecx` prior to the loop and the negation of `eax` after the loop are done using conditional assembly, where only if the operand is negative will these instructions be generated in the macro expansion. If the value of the operand is 0, then the loop will not iterate, but if the value in `eax` is 0, then the loop will iterate redundantly. Can this be solved using conditional assembly? Yes, and this is left as an exercise for the reader at the end of the chapter. Obviously there is no single instruction to output the contents of the accumulator. Since the implementation of the output is really a call to a procedure via the `INVOKE` directive and to incorporate some procedures in this example, the `OUTACC` is implemented below as a procedure. Of course the format statements and the temporary memory location `temp` will need to be defined as is shown in the program skeleton shortly:

```
OUTACC proc
    push eax ; save eax, ecx, and edx
    push ecx
    push edx
    mov temp, eax
    INVOKE printf, ADDR msg1fmt, ADDR msg1, temp
    pop edx ; restore eax, ecx, and edx
    pop ecx
    pop eax
    ret
OUTACC endp
```

Why does the above code save and restore the contents of the `eax`, `ecx`, and `edx` registers? Again remember from Chap. 2 that the `INVOKE` directive destroys the contents of these three registers. Although the `pushad` and `popad` might have been able to be used here, only three registers are used, the code is more self-documenting, and since the above is implemented as a procedure, space is not as much of a concern.

The following is the skeleton of the program which loads the accumulator with the value 1, adds the number 2 to the accumulator, adds the contents of memory location three which contains a 3, multiplies the accumulator by 4, and then multiplies the accumulator by a -3. Lastly, it outputs the contents of the accumulator:

```
. 368
.model flat, c
.stack 100h
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
printf PROTO arg1:Ptr Byte, printlist:VARARG
.data
msg1fmt byte 0Ah, "%s%d", 0Ah, 0Ah, 0
```

```

msg1 byte "The contents of the accumulator are: ",0
temp sdword ?
three sdword 3
. code
LOADACC macro operand
mov eax,operand ;; load eax with the operand
endm

ADDACC macro operand
add eax,operand ;; add to eax the operand
endm

MULTACC macro operand
push ebx ;; save ebx and ecx
push ecx
mov ebx,eax ;; mov eax to ebx
mov eax,0 ;; clear accumulator to zero
mov ecx,operand ;; load ecx with operand
if operand LT 0 ;; if operand is negative
neg ecx ; make ecx positive for loop
endif
.while ecx >0
add eax,ebx ;; repetitively add
dec ecx ;; decrement ecx
.endw
if operand LT 0 ;; if operand is negative
neg eax ; negate accumulator, eax
endif
pop ecx ;; restore ecx and ebx
pop ebx
endm

main proc
LOADACC 1
ADDACC 2
ADDACC three
MULTACC 4
MULTACC -3
CALL OUTACC
ret
main endp

OUTACC proc
push eax ; save eax, ecx, and edx
push ecx
push edx
mov temp,eax
INVOKE printf, ADDR msg1fmt, ADDR msg1, temp
pop edx ; restore eax, ecx, and edx

```

1.29. COMPLETE PROGRAM: IMPLEMENTING A MACRO CALCULATOR⁸⁷

```
    pop ecx
    pop eax
    ret
OUTACC endp
end
```

What is interesting is that the main program above only contains the following macro invocations and procedure call. As alluded to at the beginning of this section, they almost look like a new assembly language has been created:

```
LOADACC 1
ADDACC 2
ADDACC three
MULTACC 4
MULTACC -3
CALL OUTACC
```

When the above macros are expanded, the following code is generated:

```
LOADACC 1
    mov eax,1
ADDACC 2
    add eax,2
ADDACC three
    add eax,three
MULTACC 4
    push ebx
    push ecx
    mov ebx,eax
    mov eax,0
    mov ecx,4
    jmp @C0001
@C0002:
    add eax,ebx
    dec ecx
@C0001:
    cmp ecx, 0
    ja @C0002
    pop ecx
    pop ebx
MULTACC -3
    push ebx
    push ecx
    mov ebx,eax
    mov eax,0
    mov ecx,-3
    neg ecx ; make ecx positive for loop
```

```
    jmp @C0004
@C0005:
    add eax,ebx
    dec ecx
@C0004:
    cmp ecx, 0
    ja @C0005
    neg eax ; negate accumulator, eax
    pop ecx
    pop ebx
    CALL OUTACC
```

When this program is assembled, it is advisable to get a copy of the assembly listing in the . 1st file. As hinted at previously and to pique the reader's interest to read Chap. 12, although the instructions for the ADDACC macro are both add instructions, it is interesting to note that the machine code in hexadecimal is different. An add eax, 3 in machine language is 83 C0 02 and this is different than an add eax, three which in machine language is 03 05 00 00 00 32 . When using different arguments for the parameter, the instructions are different as a result of using name parameters which use strict substitution. Another thing to note is the conditional assembly of the MULACC macro when the operand is negative, where the additional neg instructions are inserted along with their associated comments due to their single semicolons (;). Lastly, note the somewhat unusual way in which the .while directive is implemented with the comparison at the bottom. Although it appears to be implemented as a post-test loop where it seems possible that the loop will execute at least once, be sure to note the jmp instruction at the beginning of the loop that prevents that from happening.

The reader is encouraged to use this program to experiment with, by using the macros given. It also forms the skeleton to add the other macros listed previously which are a part of the exercises at the end of this chapter. The reader can also add other macros and procedures as requested by the instructor or can experiment on their own initiative.

1.30 Summary

- Procedures create only one copy of the code, whereas macros create a new copy of the code every time they are invoked.
- Procedures often save and restore registers, whereas macros often do not attempt to save and restore registers.
- Procedures tend to save memory, whereas macros tend to save execution time.
- To invoke a procedure, use the call instruction followed by the name of the procedure, whereas to invoke a macro, merely put the name of the macro in the opcode field.

- Always include a ret instruction in a procedure but do not include one in a macro.
- Although more than one ret instruction can be included in a procedure, it is best to have only one and include it as the last instruction in a procedure.
- Remember to include the name of a procedure in the label field of the endp statement but do not include the name of the macro in the label field of the endm statement.
- When arguments are required when invoking a macro, use : REQ after the parameter name in the macro definition.
- When using any of the conditional assembly directives, such as if, ifb, else, and endif, do not include a period prior to the directive.
- When using the conditional assembly directive if, do not use registers or memory locations for arguments in the macro, rather use only constants. Also, when using logical operators or and and, use parentheses as in (x lt o) or (y gt o).
- Use the ifb conditional assembly directive to check for a blank matching argument and ifnb to check for a non-blank matching argument.
- Use the conditional assembly directive ifid to check if two items are identical and ifdif to check if they are different. Both of these are case sensitive, so to make them case insensitive, add the letter i to the end of the directive as in ifidi and ifdifi.

1.31 Exercises (Items Marked with an * Have Solutions in Appendix D)

1. Given the following assembly language statements and assuming memory locations and labels are properly declared, indicate whether they are syntactically correct or incorrect. If incorrect, indicate what is wrong with the statement:
 - *A. return
 - B. endm
 - C. ifdif <p>,<q>
 - *E. if eax lt o
 - *C. .ifb <parm>
 - G. call dog
 - H. endif
 - F. elseif
 - I. ifne <p1>, <p2>
2. Write a procedure to implement the factorial function as defined in the exercise section of Chap. 5.

3. Write a procedure to implement the Fibonacci numbers as defined in the exercise section of Chap. 5.
4. Write a macro to implement the factorial function as defined in the exercise section of Chap. 5.
5. Write a macro to implement the Fibonacci numbers as defined in the exercise section of Chap. 5.
6. Using conditional assembly, modify the MULACC macro defined in this chapter to not only eliminate the redundant looping but also not generate the loop instruction itself in the case that the multiplier in the operand is 0 and the answer is 0, or when it is 1 and the answer is just the value in the accumulator. (Hint: Use the power macro as an example.)
7. Implement the following instructions as macros as part of the macro calculator problem in the last section of this chapter. For the division macro, use conditional assembly to solve any problems with negative numbers. Also, when dealing with the possibility of division by 0, a -1 should be returned from the macro to indicate an error:

```
INACC  
STOREACC  
SUBACC  
DIVACC
```

1.32 Arrays

Up until this point, arrays have not been needed in the examples shown. However, this chapter will introduce the declaration of arrays, array access, indexing arrays, and how to input, process, and output arrays. Although there are many ways one can index an array, this text will present only two of them. This chapter will be concerned with the declaration of arrays of signed double words (sdword), while the declaration of an array of bytes will be introduced in the next chapter on strings. Lastly, this chapter will illustrate the use of arrays in a number of examples.

1.33 Array Declaration and Addressing

There are a couple of different ways to declare an array based on the data and the needs of the programmer. The simplest way to declare an array is to list memory location after memory location. In fact, it is entirely possible to address the next memory location after any other memory location by adding a constant to the address. For example, given the following two memory locations, it is possible to address the memory location result when referring to memory location number:

```
.data
number sdword 2
result sdword 7
```

In other words, instead of an instruction such as `mov eax, result`, which would move the integer 7 into the `eax` register, `mov eax, number +4` could also be written which would accomplish that exact same thing. Note that the `+4` does not add 4 to the contents of `number`, rather an `add` instruction would be needed to accomplish that task. Instead, it adds a 4 to the address of `number`. However, one might ask isn't result only one memory location away from `number`? That

number = 100	
result = 104	00000002

Figure 1.10: Signed double words

would be true if `number` was declared as a byte, but note that `number` is declared as an `sdword` which takes up 4 bytes as shown in Fig. 8.1.

Although the above form of addressing memory is allowable and useful in specific situations, trying to address memory locations by another variable name other than the variable name assigned to it is known as aliasing. It is not considered to be very good programming practice because it can make programs very difficult to debug and maintain, where a program that uses two different variable names for the same memory location can cause difficulty when trying to make updates to a program. To help illustrate the sort of problems that might be encountered, what if one added another variable between the variables `number` and `result` above, such as demonstrated below?

```
.data
number sdword 2
answer sdword 5
result sdword 7
```

The problem would be that when the previous instruction `mov eax, number +4` is used, the number 5 would be moved into the `eax` register instead of the number 7. Although in small programs, it might be relatively easy to find all such references, it would be very difficult in large programs. This is the reason why this method should be avoided when addressing individually declared variables.

However, it is also possible to create an array as follows, where it should be noticed that the subsequent memory locations do not have variable names attached to them:

```
numary sdword 2
    sdword 5
    sdword 7
```

Although using offsets should be avoided when addressing individual labeled memory locations, it is necessary when dealing with arrays. In the case above, the programmer has no choice but to use the variable name `numary` to access the subsequent memory locations in the array and since there are no other variable names, each of the subsequent memory locations would not be referred to via an alias. However, what if there were many entries in the above array? It could take up quite a few lines of code to create the array. Luckily, MASM has an easier way to declare the above on just one line, where the directive `sdword` need to appear only

```
numary    = 100
          = 104
          00000002
          = 108
```

Figure 1.11: Array of signed double word

once, and each of the entries would appear on the same line, each separated by a comma as follows:

```
numary sdword 2,5,7
```

In both cases, the array would appear in memory as shown in Fig. 8.2. Continuing, what if each element of an array were to be initialized to the same number, such as 0, or what if each memory location in the array did not need to be initialized? Each of the following could be used, respectively:

```
zeroary sdword 0,0,0
empary sdword ?,?,?
```

Although the above works okay for small arrays, what if there were hundreds of elements needed to be declared? Then clearly the above method would be cumbersome. Instead, the `dup` operator is very convenient, where the above would be rewritten as follows:

```
zeroary sdword 3 dup(0)
empary sdword 3 dup(?)
```

With only three elements the previous method of individually listing each element is sufficient, but as the number of elements increases, using the `dup` operator is obviously more convenient.

Given the above, how does one access individual elements of an array? For example, assume that the last element of the previously declared `numary` needed to be moved to the first element of `numary`. Remembering that arrays in C start with the zeroth element, the C equivalent of this operation would be `numary[0] = numary[2];` and the equivalent assembly code would be as follows:

```
mov eax,numary+8 ; load eax with third element
mov numary+0, eax ; store eax in first element
```

and would appear in memory as shown in Fig. 8.3.

Of course, regardless of whether one is dealing with a single memory location or with the elements of an array, memory to memory transfer needs to go through a register. As with C, the first element in an array is the zeroth element. Also, since `numary` is two memory locations away, the offset needs to be multiplied by 4 to determine the correct address because as mentioned previously, each memory




Figure 1.12: Copying individual elements of an array

location is 4 bytes long. Lastly, `numary` in the second `mov` statement technically does not need to have the `+0`, since by default the offset of a stand-alone memory location is `+0`. However, including the `+0` is a nice way of indicating to others who might read the program that the memory location specified is part of an array instead of a stand-alone memory location.

1.34 Indexing Using the Base Register

Although occasionally only access to a single element of an array is needed, more often than not access to many elements of the array is necessary and simply accessing them one at a time would prove to be inefficient. As mentioned at the outset of this chapter, there are two major ways of indexing an array. The first one is similar to indexing an array using subscripts in a high-level language,

whereas the second one is similar to using pointers and is helpful when trying to process strings as discussed in the next chapter.

Although high-level languages typically use a variable as an index when indexing arrays, in assembly language indexing is accomplished using registers. Recall from Chap. 1 that the `ebx` register is known as the base register and is very useful when indexing arrays. Although it is a register, it is used much like an index variable in the C programming language.

How would one use the `ebx` register to index an array? As an example, what if one wanted to sum all the elements of an array? To make things simpler at first, assume that the array already contains values as introduced in the last section and as shown below:

<code>numary</code>	<code>sdword</code>	<code>2, 5, 7</code>
<code>sum</code>	<code>sdword</code>	<code>?</code>

In C, the variable `sum` would first need to be initialized to 0. Further, since there are a fixed number of items to be summed, a `for` loop would be the best choice. Lastly, for each iteration of the `for` loop, the i th element of the `numary` would need to be added to `sum` using `sum=sum+numary[i]`; Optionally, the summation could be done using the shorthand notation `sum+=numary[i]`; as shown below:

```
sum = 0;
for(i=0; i<3; i++)
    sum += numary[i];
```

Of course since the `for` loop is the best choice as used above, the equivalent code in assembly language would be the use of the `.repeat-.untilcxz` directive. Unlike C code, where the variable i is being used both as a loop counter and as an index, in assembly language, two separate registers need to be used. So in addition to using the `ebx` register for indexing, the `ecx` register would be used for loop control. Lastly, since `sum` was not initialized in the data section, it must be initialized to 0 during execution time. The resulting code is as follows:

```
mov sum,0 ; initialize sum to 0
mov ecx,3 ; initialize ecx to 3
mov ebx,0 ; initialize ebx to 0
.repeat
mov eax,numary[ebx] ; load eax with element of numary
add sum,eax ; add eax to sum
add ebx,4 ; increment ebx by 4
.untilcxz
```

As discussed in the previous section, note that a 4 needs to be added to the `ebx` register to access the next signed double word. Also, do not forget that the `ecx` register should not be altered, since the `.untilcxz` directive decrements the `ecx` register by 1 automatically. After walking through the above code segment,

the contents of the registers and memory locations would be as shown in Fig. 8.4, where all values are in hexadecimal.

Why does the `ebx` register have the hexadecimal number `0000000C` (decimal 12) in it? Could it not end up addressing the memory location `sum`? Prior to the loop, the `ebx` register is initialized to 0 to access the first element in the array. After accessing and summing the current element in the array, the value in `ebx` is incremented by 4 in anticipation of accessing the next element in the array the next time through the loop. To answer the first question, during the last iteration of the loop and after the third memory location has been accessed, `ebx` is incremented by 4 to `0000000C` in anticipation of accessing the next element in the array. The answer to the second question would then be yes, where the memory location `sum`




Figure 1.13: Using `ebx` for array processing

could be accessed via `ebx` upon completion of the loop. Would this have caused an execution or run-time error if the contents of `sum` were accessed via `ebx`? As discussed in the previous section, accessing another memory location like this is possible, and no, it would not cause an error. However, it is not advisable and the result is that extra care must be taken when addressing arrays in assembly language. In all of the previous examples, the array already contains data. How can one input data into an array? To add a small twist, how could the array also be output in reverse order? Obviously this problem will require the I/O capabilities learned in Chap. 2 and will require two non-nested loops. Furthermore, assume that the user will need to first be prompted and will then enter the number of integers to be input. The code in C is as follows:

```
int array[20], n, i;
printf("\n%s", "Enter the number of integers to be input: ");
scanf("%d", &n);
if (n > 0) {
    for (i = 0; i < n; i++) {
        printf("\n%s", "Enter an integer: ");
```

```

        scanf("%d",&array[i]);
    }
    printf("\n%s\n\n","Reversed");
    for (i=n-1;i>=0;i--)
        printf(" %d\n\n",array[i]);
}
else
    printf("\n%s\n\n","No data entered.");

```

The user is first prompted to enter the number of integers to be input followed by prompts to enter the integers themselves. After the integers have been placed into the array, a loop then outputs the array in reverse order. In the event that a 0 or a negative number is entered for the first prompt, a message stating that no data was entered is displayed. The above C code can then be implemented as follows in the partial .data and .code segments below:

```

.data
msg1fmt byte 0Ah,"%s",0
msg2fmt byte 0Ah,"%s",0Ah,0Ah,0
msg3fmt byte " %d",0Ah,0Ah,0
in1fmt byte "%d",0
msg1 byte "Enter the number of integers to be input: ",0
msg2 byte "Enter an integer: ",0
msg3 byte "Reversed",0
msg4 byte "No data entered."
n sdword ?
array sdword 20 dup(?)
.code
INVOKE printf, ADDR msg1fmt,ADDR msg1
INVOKE scanf, ADDR in1fmt, ADDR n
mov ecx,n ; initialize ecx to n
mov ebx,0 ; initialize ebx to 0
.if ecx>0
.repeat
push ecx ; save ecx
INVOKE printf, ADDR msg1fmt, ADDR msg2
INVOKE scanf, ADDR in1fmt, ADDR array[ebx]
pop ecx ; restore ecx
add ebx,4 ; increment ebx by 4
.untilcxz
INVOKE printf, ADDR msg2fmt, ADDR msg3
mov ecx,n ; initialize ecx to n
sub ebx,4 ; subtract 4 from ebx
.repeat
push ecx ; save ecx
INVOKE printf, ADDR msg3fmt,array[ebx]
pop ecx ; restore ecx

```



```

sub ebx,4 ; decrement ebx by 4
.untilcxz
.else
INVOKE printf, ADDR msg2fmt, ADDR msg4
.endif

```

Notice in the assembly code segment that the use of the `.if` directive to check whether n is greater than 0 is not only helpful in outputting a message that no data was entered but also necessary to help insure that the value of the `ecx` register does not start off at 0 or a negative number for the `.repeat-untilcxz` directives that follow. Also, note that the `ecx` register is saved and restored before and after the `INVOKE` directives in the body of the loop so that the count is not destroyed. Should the contents of the `ebx` register also be saved and restored in the loop? No, because it is the only register that is not altered by the `INVOKE` directive.

1.35 Searching

As another example of using the `ebx` register for indexing, there are two main searches that the reader has probably heard about in a previous high-level programming course: the sequential search and the binary search. The former works with either unordered or ordered data, whereas the latter works with only ordered data. Since the binary search is the more complicated of the two, it is probably best to leave that to be implemented in a high-level language and this text will examine only the sequential search. Assuming that the data has already been entered into an array, that the number of elements in the array is known, and that there are no duplicates in the array, the first thing that should be done is to request from the user what data needs to be found. Then a flag must be initially cleared indicating the data has not yet been found. Next the program should loop through the array to determine whether the data being searched for is in the array. If the data is found, the flag must be set to indicate that it was found, the index set to the location of the data, and the rest of the array need not be searched. The following C program is one way of solving this problem:

```

int array[20],n=20,i,number,found;
printf("\n%s","Enter the integer to be found: ");
scanf("%d",&number);
i=0;
found=0;
while(i<n && !found)
    if(number== array[i])
        found=-1;
    else
        i++;
if (found)

```

```

    printf("\n%s\n\n", "The integer was found");
else
    printf("\n%s\n\n", "The integer was not found");

```

Note that the code was written using a while loop instead of a for loop. The use of a for loop would require the code to branch out of the middle of the loop and in C this would require the use of a break statement. This would be the equivalent of a "goto" or a jump statement in assembly language and could cause unstructured code to be written which this text has been trying to avoid. Assuming that the PROTO statements have already been written correctly, the following partial .data and .code segments implement the preceding C program:

```

.data
msg1fmt byte "%s",0
msg2fmt byte 0Ah,"%s",0Ah,0Ah,0
in1fmt byte "%d",0
msg1 byte "Enter the integer to be found: ",0
msg2 byte "The integer was found",0
msg3 byte "The integer was not found",0
array sdword 20 dup(?)
n sdword 20
number sdword ?
found sbyte ?

.code

.data
msg1fmt byte "%s",0
msg2fmt byte 0Ah,"%s",0Ah,0Ah,0
in1fmt byte "%d",0
msg1 byte "Enter the integer to be found: ",0
msg2 byte "The integer was found",0
msg3 byte "The integer was not found",0
array sdword 20 dup(?)
n sdword 20
number sdword ?
found sbyte ?

.code
INVOKE printf, ADDR msg1fmt,ADDR msg1
INVOKE scanf, ADDR in1fmt, ADDR number
mov ebx,0 ; initialize ebx to 0
mov ecx,0 ; initialize ecx to 0
mov edx,number ; load edx with number
mov found,0 ; initialize found to 0
.while(ecx<n && !found)
.if(edx==array[ebx])
mov found,-1 ; set found to -1

```

```
.else
add ebx, 4 ; increment ebx by 4
.endif
inc ecx ; increment ecx by 1
.endw
.if(found)
INVOKE printf, ADDR msg2fmt, ADDR msg2
.else
INVOKE printf, ADDR msg2fmt, ADDR msg3
.endif
```

In the above code segment, note that since only single byte is needed instead of 4 bytes to create a flag, so the found flag is only a signed byte (sbyte) instead of a signed double word (sdword). Also, notice that the ecx register is used instead of the memory location i. Although i could be used, it could not be easily used in the .while because memory to memory comparisons are not allowed. Since it would need to be transferred to a register anyway and ecx is also known as the counter register, it did not hurt to somewhat mimic the .repeat-.untilcxz directives and use the ecx register. The primary difference is that the count is going forward instead of backward and the increment of the ecx register should not be forgotten at the bottom of the loop. The other difference is the .repeat-.untilcxz directive is a post-test loop structure and the .while directive is a pre-test loop structure, and if *n* was equal to 0, the .while loop would not loop at all.

1.36 Indexing Using the esi and edi Registers

Although the use of the base register ebx is fairly straightforward, being limited to a single register can prove to be somewhat restrictive. To help with this matter, there are two index registers esi and edi, where the esi register is known as the source index register and edi is known as the destination index register. The subtle difference between the use of the ebx register and the esi and edi registers is that the former is used with the name of the array as an index and the latter are used more like pointers. In this latter case, the address of the array is first loaded into the register and the name of the array is not subsequently used. Also, these two registers are very useful when manipulating strings, as will be discussed in the next chapter. To demonstrate the different way of addressing using these registers, it is easier to examine the addressing of only a single element of an array first. For example, if only the second element of an array needed to be moved into the eax register, the instruction `mov eax, numary+4` could be written as done in the first section of this chapter. However, could the contents of `numary+4` be moved into `eax` without using `+4` attached to the name of the array? Yes, the same thing could be accomplished by using a register as an index. The following code segment would work the same assuming the existence of the previous array, `numary`:

```
mov ebx,4
mov eax,numary[ebx]
```

Although a bit cumbersome compared to using `numary + 4`, this code segment should make sense given the previous section, where `ebx` contains a 4, and `numary` indexed by `ebx` would be the address of the second element of `numary`. The dashed arrows in Fig. 8.5 show how the address of `numary`, which is 100, is added to or indexed by the 4 in the `ebx` register to create the address 104, which is the address of the second element of the array `numary`. The solid arrow in Fig. 8.5 shows the contents of memory location 104 being copied into the `eax` register.

Using the `esi` register could accomplish the same task in the following code segment which introduces the new offset operator:

```
mov esi,offset numary+4
mov eax,[esi]
```




Figure 1.14: Using `ebx` register to access a single element




Figure 1.15: Using `esi` to access a single element

Instead of causing the ``contents'' of `numary+4` to be loaded into the `esi` register, the word `offset` causes the ``address'' of `numary+4` to be loaded into the `esi` register. In the second statement, notice that there is no reference to

numary because the address of numary +4 was loaded into the esi in the first line. Now the address in esi can be used like a pointer, indicated by the square brackets, to point to the data in the second element in the array, where its contents are transferred to the eax register. The dashed arrow in Fig. 8.6 shows the esi register pointing to the second element of numary at memory location 104 and the solid arrow in Fig. 8.6 shows the contents of memory location 104 being copied into the eax register. Care must be taken when using this form of indexing. If the word offset was not included on the first line, then the "contents" of memory location numary +4, which is a 5, would be loaded into the esi register instead of the "address" of numary +4. This would have caused havoc on the second line of code because the unknown contents of memory location 5 would be loaded into the eax register. Further, if the square brackets were accidentally omitted on the second line of code, then the 104 in the esi register would be transferred to the eax register, and this is clearly not what is intended in this example.

In addition to using the offset operator, there is an alternative method. Instead of using the mov instruction with the word offset, the lea instruction, which stands for "load effective address," can be used as follows:

```
lea esi, memory+4
mov eax, [esi]
```

Although in these two instances, the end results are the same, there is a subtle difference between these two ways of getting an address. In a sense, using offset is static and lea is dynamic. When using offset, the address is calculated at assembly time, and with lea, the address is calculated at run-time. The only time the latter would need to be used is when there is a register in the second operand, where the register value could change during the course of the execution of the program and the address would need to be recalculated. Since in the above example there is no register as part of the second operand and a recalculation of the address is not necessary, either the mov and offset or the lea can be used. Further, since the use of registers in a second operand will not be used in this text, either method is acceptable. This text will use both methods interchangeably to allow readers to get use to both methods. Although the offset operator can be used with the lea instruction it is not needed, nor is it recommended, in order to keep these two methods separate and distinct.

The above shows the use of the esi register, but what of the edi register? Actually, the edi register could have been used equally well in the above example. However, as will be seen in the next chapter, some string instructions have specific uses for each register and substituting one for the other would not work. In the above case; however, if both could have been used equally well, why should one be used over the other? The answer can be found in the names of the registers mentioned above, where esi is the source index register and edi is the destination index register. As a general rule, when retrieving data from memory, it is best to use the esi register because memory is the source from which

the data is coming. When storing data back into memory, the edi register indicates the destination where the data will be placed. This use of these registers corresponds to how the string instructions work and it provides a common way of using the registers to help programmers who might look at the code in the future.

As should be recalled from Chap. 1, data can be moved from one memory location to another by simply moving the contents of one memory location to a register and then from that register to the other memory location as follows:

```
mov eax,num1
mov num2, eax
```

Although the following example is clearly less efficient than the above, it helps demonstrate how the esi and edi registers can be used to transfer data between memory locations and will be helpful in subsequent examples of array processing:

```
lea esi,num1 ; load the address of num1 into esi
lea edi,num2 ; load the address of num2 into edi
mov eax,[esi] ; move contents of where esi is pointing into eax
mov [edi],eax ; move contents of eax to where edi is pointing
```

Instead of having the memory locations as part of the two latter mov statements, they appear in the first two lea instructions, where the addresses of num1 and num2 are transferred into the esi and edi registers, respectively. Then instead of moving the contents of num1 into eax, the contents of where esi is pointing to, which is num1, are moved to eax, and then the contents of eax are moved to where edi is pointing, which is num2 as shown in Fig. 8.7.

As discussed previously, note that esi points to the source of the transfer, num1, and edi points to the destination of the transfer, num2. Could one have saved an extra instruction and just coded `mov [edi], [esi]`? At first, the answer may appear to be yes because it looks like a simple register-to-register mov instruction. However, if one thinks about it for a moment, this is not a simple register transfer, rather those registers are pointers to memory locations, and as learned in Chap. 1, memory to memory mov instructions are not allowed. As always with this form of addressing, care must be taken. For example, if the brackets around esi and edi were left off, then a 100 would be moved into eax

and the 100 would then be moved into the edi register, which is not the intended goal.

Clearly the first example is simpler, but the code above accomplishes the same task and can be expanded to transferring elements of an array. Given how the esi and edi registers work, can they be used to implement some of the array functions introduced previously? The answer is yes. As an example, consider the previous C segment which summed the elements of an array,

```
sum = 0;
```




Figure 1.16: Fig. 8.7 Using esi and edi to move data from memory to memory

```
for(i=0; i<3; i++)
    sum += numary[i];
```

which can be implemented in assembly language using the esi register as follows:

```
mov sum,0 ; initialize sum to zero
mov ecx,3 ; initialize ecx to 3
lea esi,numary+0 ; load the address of numary into esi
.repeat
mov eax,[esi] ; move contents of where esi pointing to eax
add sum,eax ; add eax to sum
add esi,4 ; increment esi by 4 to next element
.untilcxz
```

In comparison to the previous version, notice that instead of initializing the ebx register to 0, the esi register is initialized to the address of numary+0. As mentioned previously, the +0 alerts other programmers that numary is not just a simple variable, but rather is an array. The other change is that instead of accessing numary indexed by [ebx], only [esi] is used. Lastly, esi instead of ebx is incremented, but regardless of which index is used, it is incremented by 4 to access the next sdword in the array.

What if one wanted to reverse the contents of an array? Instead of trying to just write in assembly language, first think through the problem in a high-level language. Also, it is a good idea to expand the array from 3 to 5 elements so that more of a pattern can be seen:

n	sdword	5
numary	sdword	2, 4, 7, 9, 12

One of the common mistakes to be made here is an assumption that a fixed-iteration loop structure that loops five times is needed. What often happens in such circumstances is that the array is returned back to its original order. Instead, by drawing arrows to indicate which elements need to be swapped as

shown in Fig. 8.8, it should be noticed below that only approximately half the elements in the array need to be swapped. Since in this example there are an odd number of items, the middle element does not need to be swapped. Integer division can then be used to determine the number of items that need to be swapped, where in this case 5 divided by 2 is 2. The result is that the loop should iterate a fixed number of times as $n/2$.

Although this algorithm can be implemented with just one index, the i and j indexes are helpful in this example and will lend themselves to the use of the `esi` and `edi` registers in the subsequent assembly language program:

```
j=n-1;
for(i=0; i<n/2; i++) {
    temp=numary[i];
    numary[i]=numary[j];
    numary[j]=temp;
    j--;
}
```

However, the only major change that will probably be needed is to remember that `esi` and `edi` are initially loaded with the appropriate addresses of the array, as can be seen prior to the `.repeat-.untilcxz` loop below. The `esi` register is loaded with the address of the first element of the array, `numary+0`. To calculate the address of the last element of the array, the number of elements in the array is first decremented by 1 and then multiplied by 4. As seen in Fig. 8.8, 5 minus 1 is 4, times 4 is 16, which when added to 104 is 120 in decimal or 114 in hexadecimal:

```
mov ecx,n ; load ecx with contents of n
sar ecx,1 ; divide ecx by 2, number of times to loop
lea esi,numary+0 ; load address of numary into esi
mov edi,esi ; move contents of esi to edi
mov eax,n ; load eax with contents of n
dec eax ; decrement eax by one
sal eax,2 ; multiply eax by 4
add edi,eax ; add eax to edi for ending address of array
.repeat
mov eax,[esi] ; move contents where esi is pointing to eax
xchg eax,[edi] ; exchange eax and where edi is pointing
mov [esi],eax ; move eax to where edi is pointing
add esi,4 ; add four to esi for next element
sub edi,4 ; subtract four from edi for next element
.untilcxz
```

Note that the `idiv` instruction is not used in the code segment above. Since the n needs to be divided by 2, arithmetically shifting `ecx` one bit position to the right accomplishes the same task. Likewise, when calculating the ending

<u>Variable</u>	<u>Address</u>	<u>Contents</u>
n =	100	00000005
numary =	104	00000002
	108	00000004
	10C	00000007
	110	00000009
	114	0000000C

Figure 1.17: Fig. 8.8 Swap routine

address of the array, after subtracting 1, multiplication by 4 is needed, so a 2-bit arithmetic shift to the left is a simple solution (see Chap. 6). Of course, both esi and edi need to be adjusted by 4 each time through the loop. Lastly, would `xchg [esi], [edi]` have worked in the above segment? Although it would be nice if it did, remember that the `xchg` instruction cannot work between two memory locations as mentioned in Chap. 7. Also as discussed earlier in this chapter, just because it seems that two registers are being exchanged, [esi] and [edi] are actually pointing to two memory locations and thus at most only one of the operands in the exchange instruction can reference a memory location.

1.37 Lengthof and Sizeof Operators

In the previous example, the following declarations were used to declare the array and indicate its length:

```
n sdword 5
numary sdword 2,4,7,9,12
```

Putting the number of elements in a variable is clearly a better choice than to leave the number of elements as an immediate value in an instruction. The advantage of using a variable is that any time the number of elements in the array is changed, the programmer can easily change the variable *n*, which is hopefully declared in close proximity to the declaration of the array.

Although this method works and is better, it is still rather clumsy. Consider if the number 15 was added to the end of numary and the number of elements in the array changed from 5 to 6 . In addition to adding the extra element to the array, the value for n would also need to be changed as follows:

```

mov ecx,n ; load ecx with contents of n
sar ecx,1 ; divide ecx by 2, number of times to loop
lea esi,numary+0 ; load address of numary into esi
mov edi,esi ; move contents of esi to edi
mov eax,n ; load eax with contents of n
dec eax ; decrement eax by one
sal eax,2 ; multiply eax by 4
add edi,eax ; add eax to edi for ending address of array
.repeat
mov eax,[esi] ; move contents where esi is pointing to eax
xchg eax,[edi] ; exchange eax and where edi is pointing
mov [esi],eax ; move eax to where edi is pointing
add esi,4 ; add four to esi for next element
sub edi,4 ; subtract four from edi for next element
.untilcxz

```

Note that the idiv instruction is not used in the code segment above. Since the n needs to be divided by 2, arithmetically shifting ecx one bit position to the right accomplishes the same task. Likewise, when calculating the ending address of the array, after subtracting 1, multiplication by 4 is needed, so a 2-bit arithmetic shift to the left is a simple solution (see Chap. 6). Of course, both esi and edi need to be adjusted by 4 each time through the loop. Lastly, would `xchg [esi],[edi]` have worked in the above segment? Although it would be nice if it did, remember that the `xchg` instruction cannot work between two memory locations as mentioned in Chap. 7. Also as discussed earlier in this chapter, just because it seems that two registers are being exchanged, $[esi]$ and $[edi]$ are actually pointing to two memory locations and thus at most only one of the operands in the exchange instruction can reference a memory location.

```

n sdword 6
numary sdword 2,4,7,9,12,15

```

What would happen if one forgot to update the value of n from 5 to 6? In this case, the previous swapping program would process only the first five elements of the array. Now consider instead of adding an additional element to the original array, an element was removed from the array as in the following:

```

n sdword 4
numary sdword 2,4,7,9

```

What would happen if one forgot to decrease the value of n in this case? The previous swapping program code segment would still attempt to process

five elements, and whatever memory location was declared after numary would also be involved in the swap routine, which is clearly incorrect.

The solution to this problem is to neither leave the length of the array as an immediate value in an instruction nor leave it in a variable, but rather declare it using the lengthof operator. This operator instructs the assembler to calculate the length of the array at assembly time. For example, in the previous code segment from the end of the last section that reverses an array, the instruction `mov ecx, n` could be replaced with `mov ecx, lengthof numary`. Then every time the length of the array is changed, the assembler would recalculate the length of the array.

Whereas the lengthof operator indicates how many elements there are in an array, the sizeof operator indicates how many bytes there are in an array. So if with a five-element array of `sdword` the lengthof operator would return a 5, what would the sizeof operator return? Since there are 4 bytes in a `sdword` and there are 5 elements in the array, the answer would be 20.

Although the `mov eax, n` instruction in the program in the previous section could also be replaced with a `mov eax, lengthof numary` instruction, there is an even better way. Since the purpose of that section of the code segment was to calculate the number of bytes to determine the address of the last element in the array to be stored in the `edi` register, the sizeof operator could be utilized. The result is that the following code segment

```
mov eax,n ; load eax with contents of n
dec eax ; decrement eax by one
sal eax,2 ; multiply eax by 4
```

could be replaced with

```
mov eax,sizeof numary ; load eax with the size of numary
sub eax,4 ; decrement eax by four
```

In the first case, `eax` is decremented by 1, where the 4 times 4 bytes per `sdword` would be 16. Then the 16 needs to be added to the beginning address of the array to calculate the address of the last element of the array. In the second example, the size of `numary` is 20, where 4 is subtracted to get 16, which again is used to calculate the address of the last element of the array. Using both the `lengthof` and `sizeof` operators, the previous code segment could be rewritten as follows:

```
mov ecx,lengthof numary ; load ecx with length of numary
sar ecx,1 ; divide ecx by 2, of times to loop
lea esi,numary+0 ; load address of numary into esi
mov edi,esi ; move contents of esi to edi
mov eax,sizeof numary ; load eax with size of numary
sub eax,4 ; decrement eax by four
add edi,eax ; add eax to edi for ending address of array
.repeat
```

```
mov eax,[esi] ; move contents where esi is pointing to eax
xchg eax,[edi] ; exchange eax and where edi is pointing
mov [esi],eax ; move eax to where edi is pointing
add esi,4 ; add four to esi for next element
sub edi,4 ; subtract four from edi for next element
.untilcxz
```

1.38 Complete Program: Implementing a Queue

A common structure that sometimes needs to be implemented in assembly language is the queue. As learned in a second semester computer science course, a queue is known as a FIFO data structure, where the first item put into the queue is the first item taken out. The operation used to put an item in a queue is often called enqueue and the operation to remove an item from a queue is known as dequeue. Although the Intel processor has instructions to push and pop items from a stack, it does not have instructions to enqueue and dequeue items from a queue.

A queue is especially helpful when there is a faster process trying to communicate with a slower process, where data can be placed into a queue by a faster process and the slower process can then remove that data from the queue at a later time. A common example is when a fast processor needs to send something to a slow printer and data is placed in a print queue. In another example, when an interrupt occurs in a processor, it must be attended to immediately in what is known as the foreground environment. However, the foreground environment may not have time to completely process the interrupt, so the information is placed in a queue, sometimes called a foreground/background queue. Then when there are no pending interrupts, the background environment will complete the processing of the information that was previously placed in the foreground/background queue. Although the implementation of interrupts is beyond the scope of this text, the implementation of a queue provides an excellent opportunity to demonstrate the use of an array and indexing. Before looking at assembly code, it might be helpful to be reminded how queues can be implemented in a high-level language. Although ideally it would be best to use parameters, the following code again uses global variables to mimic the subsequent assembly language program. The main program uses a sentinel-controlled loop to continue to iterate until the letter *s* is input, which stands for stop. Then for each iteration of the loop, it checks for the letter *e* for enqueue or *d* for dequeue, otherwise an appropriate error message is output.

The enqueue routine should look somewhat familiar to those who have taken a second semester computer science course in data structures. Although there can be some complex ways to determine whether a queue is full, the simplest method is to use a counter such as *count*, which is checked to see if it is less than the length of the queue, *n* in this case. If so, number is placed in the rear of the queue. The variable *rear* is incremented by 1. The mod function (%) is used to cause *rear* to be reset to 0 should it exceed the length of the queue (*n*).

For example, if rear is equal to 3, the $3\%3$ is equal to 0. The dequeue routine is similar and it is left for the reader to walk through the procedure:

```
#include <stdio.h>
const int n=3;
int queue[3], number, front=0, rear=0, count=0;
char command;
int main() {
    void enqueue();
    void dequeue();
    printf("\n%s", "Enter a command, e, d, or s: ");
    scanf("%s", &command);
    while (command != 's'){
        if (command=='e'){
            printf("\n%s", "Enter a positive integer: ");
            scanf("%d", &number);
            enqueue();
        }

        else
            if (command=='d'){
                dequeue();
                if (number>0)
                    printf("\n%s%d\n", "The integer is: ", number);
            }
        else
            printf("\n%s", "Invalid entry, try again");
        printf("\n%s", "Enter a command, e, d, or s: ");
        scanf("%s", &command);
    }
    printf("\n");
    return 0;
}

void enqueue() {
    if (count<n){
        count++;
        queue[rear]=number;
        rear=(rear+1)%n;
    }
    else
        printf("\n%s\n", "Error: Queue is full");
}

void dequeue() {
    if (count>0){
        count--;
        number=queue[front];
```

```

        front=(front+1) %n;
    }
    else{
        printf("\n%s\n","Error: Queue is empty");
        number =-1;
    }
}

```

Since there needs to be a pointer for both the front and the rear of a queue, the use of the esi and edi registers, respectively, makes an excellent choice. Although procedures are often used for the enqueue and dequeue operations, the use of macros will mimic the push and pop instructions and also provide another opportunity to reinforce the concept of macros. Although not very versatile, global variables are used to communicate between the main program and the macros as with the previous C program:

```

        . }68
        .model flat,c
        .stack 100h
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
printf PROTO arg1:Ptr Byte, printlist:VARARG
        .data
in1fmt byte "%s",0
in2fmt byte "%d",0
msg1fmt byte 0Ah,"%s",0
msg3fmt byte 0Ah,"%s%d",0Ah,0
msg4fmt byte 0Ah,0
errfmt byte 0Ah,"%s",0Ah,0
msg1 byte "Enter a command, e, d, or s: ",0
msg2 byte "Enter a positive integer: ",0
msg3 byte "The integer is: ",0
errmsg1 byte "Error: Invalid entry, try again",0
errmsg2 byte "Error: Queue is full",0
errmsg3 byte "Error: Queue is empty",0
queue sdword 3 dup(?)
command sdword ?
number sdword ?
count sdword 0
        . code
enqueue macro
        .if count < lengthof queue
        inc count ; increment count
        mov eax, number ; load eax with number
        mov [edi],eax ; store eax in rear
        mov eax,edi ; copy edi (rear) to eax
        sub eax,offset queue ; subtract address of queue

```

```

    add eax,4 ; increment eax by 4
    cdq ; convert double to quad
    mov ecx,sizeof queue ; get size of queue (bytes)
    idiv ecx ; divide
    mov edi,offset queue ; load address in rear
    add edi,edx ; add remainder to rear
    .else
    INVOKE printf, ADDR errfmt, ADDR errmsg2
    .endif
    endm
dequeue macro
    .if count > 0

        dec count ; decrement count
        mov eax,[esi] ; load eax from front
        mov number,eax ; store eax in number
        mov eax,esi ; copy esi (front) to eax
        sub eax,offset queue ; subtract address of queue
        add eax,4 ; increment eax by 4
        cdq ; convert double to quad
        mov ecx, sizeof queue ; get size of queue (bytes)
        idiv ecx ; divide
        mov esi,offset queue ; load address in front
        add esi,edx ; add remainder to front
        .else
        INVOKE printf, ADDR errfmt, ADDR errmsg3
        mov number,-1 ; store -1 (flag) in number
        .endif
    endm
main proc
    mov edi,offset queue+0 ; use edi as front of queue
    mov esi,offset queue+0 ; use esi as rear of queue
    INVOKE printf, ADDR msglfmt, ADDR msg1 ; priming
    INVOKE scanf, ADDR in1fmt, ADDR command ; read
    .while command != "s" ; while not stop
    .if command=="e" ; enqueue?
        INVOKE printf, ADDR msg1fmt, ADDR msg2
    INVOKE scanf, ADDR in2fmt, ADDR number
    enqueue ; enqueue number
    .elseif command=="d" ; dequeue?
    dequeue ; deque number
    .if number >0 ; not -1 (flag)?
    INVOKE printf, ADDR msg3fmt, ADDR msg3, number
    .endif
    .else
    INVOKE printf, ADDR errfmt, ADDR errmsg1

```

```
.endif  
INVOKE printf, ADDR msg1fmt, ADDR msg1  
INVOKE scanf, ADDR in1fmt, ADDR command  
.endw  
INVOKE printf, ADDR msg4fmt  
ret  
main endp  
end
```

Although much of the assembly code is similar to its C counterpart, there are a few sections that might need some explanation, such as the enqueue routine. As previously indicated, the `edi` register is used to point to the rear of the queue. However, instead of merely adding a 1 to rear as with the previous C code, a 4

Although much of the assembly code is similar to its C counterpart, there are a few sections that might need some explanation, such as the enqueue routine. As previously indicated, the edi register is used to point to the rear of the queue. However, instead of merely adding a 1 to rear as with the previous C code, a 4 needs to be added for the sdword elements of queue. Also note that offset queue is subtracted from edi prior to the division and then offset queue is added back to edi after the division. This is because edi does not contain a simple index within queue as in the previous C program but rather edi is acting as a pointer to queue and the address of queue needs to be removed before the division can take place and then added back afterward.

1.39 Complete Program: Implementing the Selection Sort

Although sorting could be done more easily in a high-level language, it provides an excellent opportunity to examine nested loops, ifs, and again the use of the esi and edi registers. As should be recalled from a first-year sequence in computer science, there are a number of sorts called in-place sorts, which can sort the contents of an array. These sorts typically use two loops and on average if the array contains n elements, they have a time complexity of $O(n^2)$. Compared to other sorts, these are relatively slow but are acceptable with smaller sets of data and are relatively easy to learn. Three common $O(n^2)$ sorts are the selection sort, the bubble sort, and the insertion sort, of which the bubble sort is one of the more popular. Although all three sorts have a number of similarities, the selection sort tends to be a little easier to understand and a little simpler to implement, especially when trying to implement a sort in assembly language. Since many students may have already implemented a bubble sort previously in a high-level language, the implementation of the bubble sort in assembly language is left as an exercise at the end of this chapter.

The implementation of the selection sort (and the bubble sort for that matter) can be implemented in two different ways. The first method is what could be called the simplified method, where only the two loops, an if structure, and a swap routine need to be written. In this form, the sort is very inefficient but the technique of accessing various elements in the array is made quite clear. The second method of implementing the sort could be called the modified method, where the sort is modified to be more efficient and avoid any unnecessary swapping of elements or passes through the data in the array. The simplified method of the selection sort will be presented first and then it will be subsequently modified. Should the reader already be familiar with the selection sort, the following paragraph could be skipped, but it could also serve as a quick refresher as to how the sort works.

The simplified way of implementing the selection sort is to perform $n - 1$ passes through the array. During the first pass, the first element of the array will be compared with each of the subsequent elements in the array. If the ar-

ray is to be

sorted in ascending order (smallest to the largest) and if a subsequent element is smaller than the first element, a swap occurs. Comparison continues on through all of the subsequent elements. After comparing the first element with all the subsequent elements on the first pass, the smallest element will be in the first position. The process then continues on with the second element being compared to all the subsequent elements, the third element being compared to all subsequent elements, and so on until the second to the last element is compared to the last element, where the entire array will be sorted. Notice that the number of elements that need to be compared in each pass decreases by one with each subsequent pass.

The simplified selection sort can be implemented as follows in C:

```
// number of passes
for(i=0; i<n-1; i++)
    //number of comparisons
    for (j=i+1; j<n; j++)
        // compare the ith and jth element
        if (array[j]<array[i]) {
            // swap the elements
            temp=array[j];
            array[j]=array[i];
            array[i]=temp;
        }
```

Note that the outer for loop iterates $n - 1$ times for the number of passes needed through the array and that the inner for loop starts at the $i + 1$ position to compare to the subsequent number of elements. The if statement compares the subsequent element to the current element, and if it is smaller, it swaps the subsequent element with the current element as shown in Fig. 8.9.

However, the above algorithm is somewhat inefficient because it keeps swapping each time it finds a smaller number. Wouldn't it be easier to only swap once? The

answer is yes. First, the index of the starting element would be copied into a variable called `smallest`. Second, should a smaller element be found, its index is copied into `smallest`. Third, at the end of the pass the element that contains the smallest number indicated by the index in `smallest` is swapped with the element in the starting element. Now instead of a potential swap with every comparison, there is only one swap at the end of each pass as illustrated in Fig. 8.10.

The C code implementing the above algorithm can be found below:

```
// number of passes
for(i=0; i<n-1; i++) {
    // save index of first element of pass in smallest
    smallest= i;
    //number of comparisons
```




Figure 1.18: Simplified selection sort




Figure 1.19: Modified selection sort

```

for(j=i+1; j<n; j++)
    // compare jth element with smallest
    if(array[j]<array[smallest])
        // save new smallest element
        smallest=j;
    // swap first element of pass with smallest
    temp=array[i];
    array[i]=array[smallest];
    array[smallest]=temp;
}

```

The use of the two indexes i and j makes it easy to implement these two loops in assembly language using the `esi` and `edi` registers. Conversely, the use of i and j as loop counters makes it a little more difficult because only the `ecx` register can be used as a counter in nested `.repeat-` `untilcxz` loops. One solution is to use the equivalent of two while structures and thus the memory locations

i and j could easily be used as loop counters. However, the alternative is to use two `.repeat-.untilcxz` loops and to be careful to save and restore the contents of the `ecx` register at the beginning and end of the outer loop structure. In the code below, the `esi` register in a sense does double duty. At the beginning and the end of the sort, it points to the starting element, but in the middle part of the sort, it points to the smallest element. To be consistent, the input and the output also use the `esi` and `edi` registers:

```

        . }68
        .model flat,c
        .stack 100h
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
printf PROTO arg1:Ptr Byte, printlist:VARARG
        .data
msg1fmt byte 0Ah,"%s",0
msg2fmt byte "%s",0
msg3fmt byte 0Ah,"%s",0Ah,0Ah,0
msg4fmt byte " %d",0Ah,0
msg5fmt byte 0Ah,0
in1fmt byte "%d",0
msg 1 byte "Enter the number of integers to be input: ",0
msg2 byte "Enter an integer: ",0
msg3 byte "Sorted",0
n sdword ?
array sdword 20 dup(?)
temp sdword ?
        . code
main proc
INVOKE printf,ADDR msg1fmt,ADDR msg1
INVOKE scanf,ADDR in1fmt,ADDR n
INVOKE printf,ADDR msg5fmt
.if n>0 ; if n <= 0, don't continue
mov ecx,n ; load ecx with n
mov edi,offset array+0 ; load address of array into edi
.repeat
push ecx ; save ecx
INVOKE printf,ADDR msg2fmt,ADDR msg2
INVOKE scanf,ADDR in1fmt,ADDR [edi]
add edi,4 ; increment edi by 4
pop ecx ; restore ecx
.untilcxz
.if n>1 ; check >1 elements in array
mov ecx,n ; load ecx with n
dec ecx ; loop n-1 times
mov esi,offset array+0 ; load esi with address of array
.repeat

```

1.39. COMPLETE PROGRAM: IMPLEMENTING THE SELECTION SORT¹¹⁷

```
    push ecx ; save ecx
    push esi ; save address, esi now smallest
    mov edi,esi ; load address of esi in edi
    add edi,4 ; move edi to the next element
    .repeat
    mov eax,[esi] ; move smallest to eax to compare
    .if [edi]<eax ; compare smallest to next
    mov esi,edi ; save the new smallest in esi
    .endif
    add edi,4 ; move to next element to compare
    .untilcxz
    mov edi,esi ; edi points to smallest element
    pop esi ; esi points to the start element
    mov eax,[esi] ; move start element to temp
    xchg eax,[edi] ; exchange start and smallest
    mov [esi],eax ; move smallest back to start
    add esi,4 ; move start index to next
    pop ecx ; restore ecx to be decremented
    .untilcxz
    .endif
    INVOKE printf, ADDR msg3fmt, ADDR msg3
    mov ecx, n ; load ecx with n
    mov esi,offset arry+0 ; load esi with address of arry
    .repeat
    push ecx ; save ecx
    mov eax,[esi] ; load eax with element from arry
    mov temp,eax ; store eax in temp for output
    INVOKE printf, ADDR msg4fmt, temp
    add esi,4 ; increment esi to next element
    pop ecx ; restore ecx
    .untilcxz
    INVOKE printf, ADDR msg5fmt
    .endif
    ret
main endp
end
```

Notice that edi is used on input because the array is the destination of the data and esi is used on output because the source of the data is the array. Again, esi does double duty, where at the beginning and the end of the sort, it points to the starting element, but in the middle part of the sort, it points to the smallest element. The push and pop instructions are carefully used when making the transition between these two tasks. Lastly, note that a temporary memory location is used to help output the array.

1.40 Summary

- The `dup` operator allows for the declaration of large initialized or uninitialized arrays.
- The `ebx` register can be used as an index for an array, much like a variable such as i in a high-level language.
- The `esi` and `edi` registers are known as the source index register and destination index register, respectively. They work like pointers and are especially useful with strings.
- When dealing with arrays of `sdword`, remember to increment by 4 instead of 1, because a signed double word takes up 4 bytes.
- The `mov` instruction and offset operator, or the `lea` instruction, allows for getting the address of a variable, where the former is static and the latter dynamic.
- Use square brackets `[]` around the `ebx`, `esi`, and `edi` registers, not to get the contents of the register, but rather to get the contents of the memory location to which they are indexing or pointing.
- The `lengthof` operator returns how many elements are in an array, whereas the `sizeof` operator returns how many bytes there are in an array.

1.41 Exercises (Items Marked with an * Have Solutions in Appendix D)

- Given the following assembly language statements, indicate whether they are syntactically correct or incorrect. If incorrect, indicate what is wrong with the statement:
 - *A. `x sdword ?,?,?`
 - B. `y sdword 3 dup(o)`
 - *C. `mov eax,x+8`
 - D. `mov eax,y[ebx]`
 - *E. `mov esi,edi`
 - F. `mov [esi], [edi]`
- Given the contents of the following memory location, what is stored in the `eax` register at the end of each segment?

```
temp = 200    00000005
= 204        00000007
```

- *A. `mov eax,temp`
- B. `mov eax,offset temp`

```
*C. lea eax,temp
    D. mov eax, offset temp+4
*E. mov esi,offset temp
    F. mov edi,offset temp
       mov eax,[edi]
       mov eax,esi
    H. mov esi,offset temp+4
       mov eax,2
    G. lea esi,temp
       lea edi,temp+4
       mov eax,[esi]
       imul [esi]
       add eax,[edi]
```

3. Implement the following C instructions using assembly language. Assume all variables are declared as `sdword`:

```
*A. num[0] = 1;
    B. x[1] = x[2];
*C. y[i+1] = y[i];
    D. z[i] = z[j];
```

4. Given the declarations below, indicate what would be stored in the `eax` register for each of the following instructions. Note that `oarray` is of type `sword`, not `sdword` (hint: see Chap. 1):

```
narray sdword 1,2,3,4,5
marray sdword 10 dup(?)
oarray sword 15,20,25
```

```
*A. mov eax,lengthof narray
*B. mov eax,sizeof narray
    C. mov eax,lengthof marray
    D. mov eax,sizeof marray
*E. mov eax,lengthof oarray
    F. mov eax,sizeof oarray
```

5. Write both the C code and the assembly code to transfer the contents of a 20-element array of integers to a second 20-element array of integers.
6. Just as there is a simple and modified version of the selection sort, so is there both a simple version and a modified version of the bubble sort. The simple version in C is the same length as the simplified version of the selection sort presented in Sect. 8.6. a. Write both the C code and the assembly code to implement the simplified version of the bubble sort which compares every element of every pass through the array whether there was a swap on the previous pass or not.
- b. First write the C code for the modified version of the bubble sort and then write the modified version in assembly language. With the modified version, if there is not a swap on the previous pass through the array, the array is in order and there is no need to make any subsequent passes through the array.

1.42 Introduction

This chapter concerns string processing. Specifically it examines various string processing instructions that are available in MASM. Continuing on with the last chapter, it also examines the manipulation of arrays of strings.

In its simplest case, a string is nothing more than an array of bytes as opposed to an array of signed double words. So, it is possible to use all of the techniques for arrays introduced in the last chapter with strings. For example, what if one wanted to copy the contents of one string to another string? As with arrays, the `ebx` register could be used an index for both strings. However, there are a couple of subtle but important changes in the following code:

```

                .data
string1 byte "Hello World!",0
string2 byte 12 dup(?),0
                .code
mov ecx,12 ;load ecx with 12
mov ebx,0 ;load ebx with 0
.repeat
mov al,string1[ebx] ;load al with string1[ebx]
mov string2[ebx],al ;store al in string2[ebx]
inc ebx ;increment ebx by 1
.untilcxz

```

First, note that the strings are declared as byte instead of `sdword`. As a result, notice that the `mov` instructions are using only a 1-byte register, `al`, instead of a 32-bit register `eax`, where it should be recalled from Chap. 1 that `al` is the rightmost byte of the `eax` register. Lastly, instead of incrementing `ebx` by four, it is incremented only by one to account for the size of byte as opposed to the size of a `sdword`.

Just as arrays can be indexed using the `esi` and `edi` registers, so too can strings be indexed using these registers. The `esi` register can be used for `string1` as the source of the transfer, and the `edi` can be used for the destination of the transfer to `string2`:

```

                .data
string1 byte "Hello World!"
string2 byte 12 dup(?)
                .code
mov ecx,12 ; load ecx with 12
lea esi,string1 ; load esi with address of string1
lea edi,string2 ; load edi with address of string2
.repeat
mov al,[esi] ; load al with [esi]
mov [edi],al ; store al in [edi]

```



```
inc esi ; increment esi by 1
inc edi ; increment edi by 1
.untilcxz
```

As with the `ebx` register previously, `esi` and `edi` are incremented by one instead of four. As seen in the previous chapter, arrays are typically drawn vertically, and in representing how strings are processed, they are typically drawn horizontally. This is especially helpful when trying to represent arrays of strings later in this chapter. The above code could be illustrated just prior to the `mov al, [esi]` instruction the second time through the loop as shown in Fig. 9.1, where `b` represents a blank or a space.

First, the letter *H* has been transferred from the first byte in `name1` to the first byte in `name2`. Also, the `ecx` register has been decremented by one from 12 to 11, 0000000B in hexadecimal. Lastly, the `esi` and `edi` registers have been incremented by one and are pointing to the next byte in each of the strings.




Figure 1.20: Using `esi` and `edi` to move a string

1.43 String Instructions: Moving Strings (movsb)

Since there are a number of functions that need to be performed on strings, many high-level languages include specialized libraries with instructions that perform many of the unique functions that help in string processing. Although assemblers usually do not come with a library of string functions, that does not prevent users from creating their own libraries. Whether or not a library of string functions is created, or code for string processing is written on an as-needed basis, creating the code needed using the same instructions that can be used for arrays can be tedious and cumbersome. Luckily, there are some instructions provided in the Intel architecture that help to make the task a little easier.

Before looking at these instructions, it should be pointed out that some of the instructions are designed to help with array processing as well. But since the needs when processing an array tend to be different, programmers often times use the mechanisms described in the previous chapter and use the following instructions primarily for string processing. As will be shown later, having the ability to use both techniques can be helpful in some instances. As with previous instructions, there are many different options available, but only the basic and most useful ones will be discussed here.

Instruction	Meaning
movsb	Move string byte
cmps	Compare string byte
scasd	Scan string byte
stosb	Store string byte
lodsb	Load string byte

Although the above instructions have their word and double word counterparts, which can be created by substituting the letter b in the instructions with the letters w and d, respectively, only the byte instructions are listed above. One of the most useful instructions listed above to be discussed in this section is the movsb instruction which is used to move a string of bytes. This instruction is not a simple instruction because it does more than one thing. In particular, the movsb instruction does two things. First, it moves the contents of the byte pointed at by the esi register to the byte in memory pointed at by the edi register. Then it decrements the ecx register and either increments or decrements the esi and edi registers by 1. So, for example, if one wanted to move only a single byte, then whether the esi and edi registers are incremented or decremented would not matter and one could write the following code segment:

```
.data
letter1 byte 'a'
letter2 byte ?
.code
lea esi,letter1 ; load esi with address of letter1
lea edi,letter2 ; load edi with address of letter2
```

`movsb ; move string byte from [esi] to [edi]`

However, obviously the above code segment could be written much easier by simply writing the following:

```
mov al,letter1 ; load al with letter1
mov letter2,al ; store al in letter2
```

As can be seen, by itself the `movsb` instruction is not terribly useful, but when used in conjunction with some other instructions, it becomes a fairly powerful instruction. The fact that the `movsb` instruction can alter `esi` and `edi` can be very useful when moving a number of bytes. The way to determine which way `esi` and `edi` will be altered is based on the direction flag (mentioned previously in Chap. 4). The direction flag can be cleared or set by using the `cld` or `std` instructions, which stand for clear direction flag or set direction flag, respectively. If the `esi` and `edi` registers need to be incremented, the direction flag then needs to be cleared (`cld`), otherwise the direction flag needs to set (`std`) to cause the registers to be decremented. Using this information, the `movsb` instruction could be used to simplify the loop used in the previous section and relisted below on the left, where the modified code using the `movsb` instruction is listed on the right:

```
mov ecx,12
mov esi,offset string1
mov edi,offset string2
.repeat
mov al,[esi]
mov [edi],al
inc esi
inc edi
.untilcxz

mov ecx,12
mov esi,offset string1
mov edi,offset string2
cld
.repeat
movsb
.untilcxz
```

Note that `ecx`, `esi`, and `edi` all need to be initialized as before, but notice that the body of the loop is much smaller because of the power of the `movsb` instruction. An intermediate general purpose register is not needed to move from one memory location to another. Also note that it is not necessary for the programmer to increment the `esi` and `edi` registers, because it is automatically done by the `movsb` instruction. The only thing that needs to be done is to clear the direction flag using the `cld` instruction to cause the two registers to be incremented as opposed to being decremented, which only needs to be done once prior to the loop.

To help with making string processing even simpler, the above code can be further simplified by using a prefix. There are three prefixes that are useful here as listed below:

Prefix	Meaning
rep	Repeat
repe	Repeat while equal
repne	Repeat while not equal

The rep prefix works just like the .repeat-.untilcxz directives, where it decrements the ecx register until it reaches 0. Unlike the .repeat-.untilcxz directives which can have any number of instructions in the body of the loop, the rep prefix works only in conjunction with instructions like movsb. Thus the above code on the right and shown again below on the left can be further simplified as follows on the right:

```
mov ecx,12
lea esi,string1
lea edi,string2
cld
.repeat
movsb
.untilcxz
```

```
mov ecx,12
lea esi,string1
lea edi,string2
cld
rep movsb
```

As mentioned previously, the movsb instruction is not very useful when used as a stand-alone instruction, but when used in conjunction with other instructions, its power is evident. However, as with any instruction that has a lot of power, it loses some of its versatility. If each character needed to be processed as it was moved from one string to another, then the movsb instruction is not very useful. Since all it does is move bytes, which it does very well, it cannot do much of anything else. This is not to disparage the power of the movsb, but rather to show its limitations. For example, if each letter moved from one string to the other needed to be changed from lowercase to uppercase, the movsb would not be as useful and the following code segment would instead accomplish this task. For the sake of convenience in the example below, it is assumed that each character in the variable name1 is a letter of the alphabet:

```
.data
name1 byte "MaryJo"
name2 byte 6 dup(?)
```

```

.code
mov ecx,lengthof name1 ; load ecx with length

mov ecx,lengthof name1      ; load ecx with length
lea esi,name1               ; load esi with address of name1
lea edi,name2               ; load edi with address of name2
.repeat
mov al,[esi]                ; load al with [esi]
and al,11011111b            ; convert lower to upper case
mov [edi],al                ; store al in [edi]
inc esi                     ; increment esi by 1
inc edi                     ; increment edi by 1
.untilcxz

```

1.44 String Instructions: Scanning (scasb), Storing (stosb), and Loading (lods b)

The scasb instruction is a fairly useful instruction when used with either the repe or repne prefixes. For example, what if one wanted to scan a string of bytes to find whether there was a particular character in a string, such as a blank? First, the al register is loaded with the character that needs to be found. Then the edi register is loaded with address of the string to be scanned. And lastly, the ecx register needs to be loaded with the number of characters to be scanned in the string. Then each time a character is scanned in the string, the ecx register is decremented by 1 and the edi register is incremented by 1. Using the repne prefix, the string will be scanned until the character is found or until the ecx is equal to 0:

```

.data
name1 byte "Abe Lincoln"
.code
mov al,' ' ; load al with a space
mov ecx,lengthof name1 ; load ecx with length
lea edi,name1 ; load address of name1
repne scasb

```




Fig. 9.2 Using repne scasb

In the above code segment, `ecx` is initially loaded with the length of `name1`, which is 11 (B in hex). After scanning for a space, or in other words, a blank as indicated by a lowercase `b` in Fig. 9.2, the `ecx` register will have been decremented to 7, which is how many characters are left to be scanned in `name1`, and the `edi` register would be pointing to the next character in the string as shown.

The `stosb` instruction is useful to store the contents of the `al` register at the location in a string pointed at by the `edi` register. The `lodsb` instruction loads the `al` register with the contents of the string pointed at by the `esi` register. In both cases the `edi` or the `esi` register is incremented after their respective operations. The semantics of each instruction is shown below using equivalent `mov` and `inc` instructions.

String instruction	Equivalent
<code>stosb</code>	<code>mov al, [edi]</code> <code>inc edi</code>
<code>lodsb</code>	<code>mov [esi], al</code> <code>inc esi</code>

As can be seen, the string instructions are a little cleaner and have the option of using the `rep` prefix. Although both instructions could be used with a `rep` prefix, only the `stosb` instruction would benefit the most by its use when initializing a string to blanks or some other character. As an example of how many of the above instructions might be used, consider the task of taking someone's name in the normal first name followed by last name order and then reversing it so that the last name is first, followed by a comma, followed by a space, and then followed by the last name:

```
.data
name1 byte "Abe Lincoln"
name2 byte 12 dup (?)
.code
mov al, ' ' ; load al with space
mov ecx, lengthof name1 ; load length of name1
lea edi, name1 ; load address of name1
repne scasb ; find space in name1
push ecx ; save ecx
mov esi, edi ; move edi to esi
lea edi, name2 ; load address of name2
rep movsb ; copy last name to name2
mov al, ',' ; load al with comma
stosb ; store comma in name2
mov al, ' ' ; load al with comma
stosb ; store space in name2
mov ecx, lengthof name2 ; store length of name2
pop eax ; restore ecx into eax
sub ecx, eax ; sub length of last name
dec ecx ; decrement ecx for space
```

```
lea esi,name1 ; load address of name1
rep movsb ; copy first name to name2
```

In the above program, first the space needs to be found between the first and the last name in name1 using the scasb instruction. After ecx is saved for subsequent processing, the last name from name1 would be copied to name 2 using the movsb instruction. A comma and a space then need to be inserted into name 2 using the stosb instruction. Then using the previously saved value of ecx to determine the length of the first name, the first name can be copied from name1 to name 2 using the movsb instruction.

1.45 Array of Strings

What if one wanted to move an array of strings to another array of strings? It is possible that one could just treat the entire array of strings as one giant string and use a single rep movsb instruction. Given the following array of strings,

names1 byte "Abby","Fred","John","Kent","Mary"

it could be viewed as shown in Fig. 9.3.

1.46 names1

A	b	b	y	F	r	e	d	J	o	h	n	K	e	n	t	M	a	r	y
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Fig. 9.3 Strings viewed as a single string

names1			
A	b	b	y
F	r	e	d
J	o	h	n
K	e	n	t
M	a	r	y

Figure 1.21: Strings viewed as an array of strings

Although using only a rep movsb would work, the disadvantage of this approach is, what if one wanted to process each string individually? Instead it is helpful to not view names1 as just a single string but rather to view it as an array of strings as shown in Fig. 9.4.

Although the following solution to the problem does not address the issue of processing each string separately, it lays the groundwork for that in a future problem and illustrates a more versatile way to process an array of strings. What are needed are two loops: one to control the array of strings and the other for each character in the string:

```

.data
names1 byte "Abby","Fred","John","Kent","Mary"
names2 byte 20 dup(?)
.code
main proc
mov ecx,5      ; load ecx with 5
lea esi, names1 ; load esi with address of names1
lea edi, names2 ; load edi with address of names2
cld           ; clear direction flag
.repeat
push ecx      ; save ecx
mov ecx,4     ; load ecx with 4
rep movsb    ; move string from names1 to names2
pop ecx      ; restore ecx
.untilcxz

```

Note that the value of ecx needs to be pushed and popped prior to and after the rep movsb instruction since the rep prefix uses the ecx register too. Further, the cld could have been placed in the loop prior to the rep movsb instructions, but since the value of the direction flag does not change, there is no need to clear it each time and it can be placed just prior to the outer loop. Again, although the above code could have been done with only a rep movsb instruction, it lays the foundation for a problem in a subsequent section.

1.47 String Instructions: Comparing Strings (cmpsb)

What if one wanted to compare two strings to see if they are equal? As might be suspected, this can be done with a loop and an if structure, but the code for this is rather ungainly. Instead of showing this option, it is much easier to go straight to the instruction designed for this task, which is the cmpsb instruction used to compare a string of bytes. It is very similar to the movsb instruction, where it performs two major tasks. First, it compares the two bytes pointed to by the esi and edi registers and sets the appropriate flags such as the zero and sign flags. It then decrements the ecx register, and increments or decrements the esi and edi registers as indicated by the direction flag, just like with the movsb instruction. Also, similar to the rep operator with the movsb, a repe prefix can be added to allow loop control. The repe repeats, while the two characters being compared are equal. For example, given the following declaration with equal length strings, what would happen in the following code segment?

```

.data
name1 byte "James"
name2 byte "James"
.code
mov ecx, lengthof name1 ; load ecx with length

```



```
lea esi,name1      ; load address of name1
lea edi,name2      ; load address of name2
cld                ; clear direction flag
repe cmpsb         ; repeat while equal
```

First, upon completion of the segment, `ecx` would be 0. Assuming that `name1` was at memory location 100 and `name2` was at memory location 105 as in the drawing in Fig. 9.5, the final values of `esi` and `edi` would be 105 and 10A (in hex), respectively.

But what if `name1` and `name2` contained James and Jamie, respectively?

Fig. 9.5 Using `repe cmpsb` with identical strings

```
.data
name1 byte "James"
name2 byte "Jamie"
```

Since the `repe cmpsb` instructions stop after a difference is found, the value of `ecx` would be 1, after being decremented four times for the four characters that were compared. The value of `esi` and `edi` would be 104 and 109, respectively, as shown in Fig. 9.6.

At first, it would seem that the final two values of the `ecx` register would be useful in determining whether two strings are equal or not. However, what if the difference between the two strings was only in the last position as in the following?

```
.data
name1 byte "Marci "
name2 byte "Marcy "
```

Again, since `repe cmpsb` stops when a difference is found, the value of `ecx` would be 0 after being decremented five times, the value of `esi` would be 105, and `edi` would be 10A, as shown in Fig. 9.7.

If the difference between the two strings is in the last position of the strings, then the results in the registers are no different than if the two strings were identical as in Fig. 9.5. The simple solution to this problem is to make sure that the strings are at least 1 byte longer than the data they contain and there is a blank in the last position as in the following:

```
.data
name1 byte "Marci "
name2 byte "Marcy "
```

Then the previous code can be used to check whether two strings are equal or not. For example, if `ecx` is 0, then it is known that the `repe cmpsb` made it throughout the entire string, including the blanks at the end, and that they are equal. However, if `ecx` is not 0, then it is known that the two strings were not equal as shown in Fig. 9.8.

Fig. 9.8 Using `repe cmpsb` with extra blanks $\square$

```
ecx 00000001
```

So after the `repe cmpsb` instructions in the code segment above, an `.if` directive could be added to test the contents of the `ecx` register and act accordingly. However, what if adding extra bytes to the string is not a possibility? What would happen if the difference occurred in the last element of the string and `ecx` was a 0? Returning to the previous example which did not contain the extra spaces,

```
.data
name1 byte "Marci"
name2 byte "Marcy"
```

the code would still need to check to see if `ecx` was 0, but there would need to be an additional check to determine whether there was a difference between the last two characters. Since the values of `esi` and `edi` would be one beyond the last character as they are in Fig. 9.7, they would need to be "backed up" or decremented by 1 to determine whether there was a difference in the last character as illustrated in Fig. 9.9.

Modifying the previous code segment above to include input and output, the following complete program illustrates how this would work. Should `ecx` be greater than 0, then the two strings are different. However, if `ecx` is equal to zero, then `esi` and `edi` would need to be decremented to determine whether the strings are same. Note that this program assumes that equal length strings will be input:

```
. 386
.model flat,c
.stack 100h
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
printf PROTO arg1:Ptr Byte, printlist:VARARG
.data
msg1fmt byte "%s",0
msg2fmt byte 0Ah,"%s",0Ah,0Ah,0
in1fmt byte "%s",0
msg1 byte "Enter a first name: ",0
msg2 byte "Enter another first name: ",0
msg3 byte "The names are not the same.",0
msg4 byte "The names are the same.",0
name1 byte 6 dup(" ")
name2 byte 6 dup(" ")
.code
main proc
    INVOKE printf, ADDR msg1fmt, ADDR msg1
    INVOKE scanf, ADDR in1fmt, ADDR name1
    INVOKE printf, ADDR msg1fmt, ADDR msg2
```

```

    INVOKE scanf, ADDR in1fmt, ADDR name2
    mov ecx,lengthof name1 ; load ecx with length
    lea esi,name1 ; load address of name1
    lea edi,name2 ; load address of name2
    cld ; clear direction flag
    repe cmpsb ; compare while equal
    .if ecx > 0 ; check if ecx > 0
    INVOKE printf, ADDR msg2fmt, ADDR msg3
    .else
    dec esi ; back up esi one position
    dec edi ; back up edi one position
    mov al,[esi] ; load al with [esi]
    .if al != [edi] ; if not equal
    INVOKE printf, ADDR msg2fmt, ADDR msg3
    .else
    INVOKE printf, ADDR msg2fmt, ADDR msg4
    .endif
    .endif
    ret
main endp
end

```

Looking carefully at the declaration of name1 and name2, it should be noted that they have been declared to be 6 bytes long instead of 5 bytes. Is the reason for adding the extra byte to make the comparisons easier as discussed previously? The

answer is no, because the data is not a pre-defined string, but rather it is being input via an INVOKE scanf instruction. When the user keys in a string, the last thing they do is press the Enter or Return key at the end of the string. The Enter key appears at the end of the string as a binary zero, not unlike the 0 that appears at the end of a string such as msg 3 and msg 4 in the .data section in the above program. The result is that space needs to be made for the binary zero, otherwise it could spill over into other memory locations and could cause various logic errors.

What if one wanted to determine if the content of one string is less than or greater than another string? For example, what if one wanted to put the names in alphabetical order? Unfortunately, an instruction does not exist for this, but looking back at the previous program, there is a hint of a possible solution to the problem. Previously, when ecx was 0, esi and edi needed to be backed up to determine whether or not the strings were equal. What would need to happen if ecx was not 0 in the previous case as shown below?

```

.data
name1 byte "James"
name2 byte "Jamie"

```

The repe cmpsb instruction stopped after a difference was found and in this

case the value of `ecx` was 1 after being decremented four times for the four comparisons. However, to determine whether the character is less than or greater than the other, the values of the `esi` and `edi` registers would again need to be backed up in this case to 103 and 108, respectively, as shown in Fig. 9.10.

In other words, regardless of whether `ecx` is 0 or not, `esi` and `edi` need to be backed up so that the last two characters can be compared to determine whether they are the same, or whether one is larger or smaller. Using some of the same code from the previous example, where it is assumed that equal length strings for `name1` and `name2` have already been prompted for and input, it can be modified as follows:

```

; *** assume previous INVOKEs, prompts, and formats
msg3 byte "The names are the same.",0
msg4 byte "The first name is less than the second.",0
msg5 byte "The first name is greater than the second.",0
name1 byte 6 dup(" ")
name2 byte 6 dup(" ")
.code
main proc
;*** assume previous prompts and input
mov ecx,lengthof name1 ; load ecx with length
lea esi,name1 ; load address of name1
lea edi,name2 ; load address of name2
cld ; clear direction flag
repe cmpsb ; compare while equal
dec esi ; back up esi one position
dec edi ; back up edi one position
mov al,[esi] ; load al with [esi]
.if al == [edi] ; if equal
INVOKE printf, ADDR msg2fmt, ADDR msg3
.else
.if al < [edi] ; if less than
INVOKE printf, ADDR msg2fmt, ADDR msg4
.else
INVOKE printf, ADDR msg2fmt, ADDR msg5
.endif
.endif
.endif

```

1.48 Complete Program: Searching an Array of Strings

As mentioned previously, although an array of strings can be moved by simply using a single loop or by using merely a `rep movsb` instruction, the disadvantage is that the individual strings cannot be processed. In order to accomplish individual processing of each string in the array, two loops should be employed. To

illustrate this, consider the problem of searching an array of fixed length strings sequentially. As before, an outer loop is needed to iterate through each string in the array and an inner loop is needed to compare each character in the string. For this latter task, the `repe cmpsb` instructions learned in the previous section is the obvious choice. Note that since `ecx` is used for the loop control variable of the outer loop, its value must be pushed prior to the `repe cmpsb` instruction and then popped prior to the end of the outer loop:

```

        . }68
        .model flat,c
        .stack 100h
Scanf PROTO arg2:Ptr Byte, inputlist:VARARG
printf PROTO arg1:Ptr Byte, printlist:VARARG
        .data
msg1fmt byte 0Ah,"%s",0
msg2fmt byte 0Ah,"%s",0Ah,0Ah,0
in1fmt byte "%s",0
msg1 byte "Enter the state to be found: ",0
msg2 byte "The state was found.",0
msg3 byte "The state was not found.",0
arrystr byte "Illinois ","Michigan ","Iowa ",
             "Missouri ","Arkansas ","Tennessee",
             "Louisiana","Arizona ","Montana ",
             "Ohio "
n sdword 10
string byte 10 dup(?)
found sdword ?
        .code
main proc
    INVOKE printf, ADDR msg1fmt,ADDR msg1
    INVOKE scanf, ADDR in1fmt, ADDR string
    mov ecx,0 ; initialize ecx to 0
    mov found,0 ; initialize found to 0
    lea edi,arrystr+0 ; load edi with address
    .while(ecx<n && found != -1)
    push ecx ; save ecx
    lea esi,string+0 ; load address of string
    cld ; clear direction flag
    mov ecx,lengthof string ; load length of string
    repe cmpsb ; compare while equal
    dec esi ; decrement esi
    dec edi ; decrement edi
    mov al,[esi] ; load al with [esi]
    mov ah,[edi] ; load ah with [edi]
    .if (al==0)&&(ah==" ") ; compare for 0 and space
    mov found,-1 ; if yes, found

```

```

    .endif
    inc edi ; increment edi back

        add edi,ecx ; adjust edi to next string
    pop ecx ; restore ecx
    inc ecx ; increment ecx
    .endw
    .if (found == -1)
    INVOKE printf, ADDR msg2fmt, ADDR msg2
    .else
    INVOKE printf, ADDR msg2fmt, ADDR msg3
    .endif
    ret
main endp
end

```

Although each string in the array of strings in the above program is of equal length, each string entered is of variable length. Again these strings are terminated by the pressing of the Enter key which appears at the end of the string as a binary 0. In the case where the input string is equal to the fixed length string in the array, `repe cmpsb` will stop because the binary 0 will not equal the space in the array. For example, "Tennessee", 0 will not be equal to "Tennessee". So, the `esi` and `edi` registers are backed up to check for the 0 and the space, and if they are there, the string has been found. However, in the case that the strings are different, such as "Tennessea", 0, the difference will occur before the binary 0, and thus when backing up the `esi` register, the binary 0 will not be found. In the case where the string that is input is longer, such as "Ioway", 0 as compared to "Iowa", again the difference would occur before the binary 0 is encountered. However, if all these cases are caught by checking for the binary 0, why check for the space? If the presence of the space was not checked for, then "Iow", 0 would end up being equal to "Iowa", which is clearly incorrect.

Also notice that when the string is not found in the array, the value `edi` will be incorrect, so it is necessary to adjust the values to point to the beginning of the next string in the array. This is accomplished by first incrementing `edi` back to its original location to account for the decrement discussed in the previous paragraph. Then, the number of characters left in the string needs to be added, which is the value in the `ecx` register, to the `edi` register.

1.49 Summary

- The `movsb` instruction moves a string of bytes from where the `esi` register is pointing to where the `edi` register is pointing. The registers are then incremented or decremented based on the direction flag.
- The `cmpsb` instruction compares a byte in a string pointed to by the `esi` and `edi` registers. The registers are then incremented or decremented based

on the direction flag.

- Do not forget to clear the direction flag using `cld` to increment `esi` and `edi` or set the direction flag using `std` to decrement `esi` and `edi` prior to either the `cmps` or `movs` instructions.
- The `rep` prefix prior to `movsb` loops the number of times indicated by the `ecx` register, decrements the `ecx` register by 1, and continues until `ecx` is 0.
- The `repne` prefix prior to `cmps` works like the `rep` prefix in that it will quit looping when `ecx` equals 0, but it will also quit looping when the 2 bytes pointed to by `esi` and `edi` are not equal. The `repne` prefix is like `repe` but will instead quit looping when the 2 bytes are equal.
- The `scasd` instruction will scan a string for the character located in the `al` register and the `edi` register will point 1 byte after the location it is found. The `stosb` instruction will store the character in the `al` register in a string at the location pointed at by the `edi` register. And the `lods` instruction will load the character in a string pointed to by the `esi` register into the `al` register.

1.50 Exercises (Items Marked with an * Have Solutions in Appendix D)

1. Given the following assembly language statements, indicate whether they are syntactically correct or incorrect. If incorrect, indicate what is wrong with the statement:
 - *A. `movesb`
 - B. `cmps`
 - *C. `scasb`
 - D. `stosb`
 - *E. `rept strsb`
 - F. `loads`
2. Given the following declarations, walk through the following code segments and indicate the contents of the `ecx`, `esi`, `edi`, and `al` registers upon completion of each segment. You may assume that `string1` starts at memory location 100 and `string2` at memory location 105. With problem D, in addition to the registers, what would be the contents of `string2`?

`string1` byte "abcde"

`string2` byte "abcyz"

*A. `mov ecx,5`

`mov al,"c"`

`mov edi,offset string1`

`rep scasb`

```

B. mov esi,offset string1+3
   lodsb
*C. mov ecx,5
   mov esi,offset string1
   mov edi,offset string2
   repe cmpsb
D. mov ecx,5
   mov esi,offset string1
   mov edi,offset string2
   repne cmpsb
E. lea edi,string2
   mov al,"d"
   stosb

```

3. Using the esi and edi registers and a repeat-untilcxz loop, determine whether the word in a string is a palindrome. For the sake of convenience, assume that the string is 10 elements long and all the words in the string are also 10 characters long. Do not use a stack.

1.50.1 Floating-Point Instructions

This chapter introduces the fundamentals of floating-point instructions. This includes registers, memory storage, input/output, the instructions needed to perform arithmetic, and basic control structures. In order to learn the concepts in this chapter one should understand how floating-point numbers are represented in memory from a computer organization class or text. Alternatively, one can review or learn the basic concepts needed as presented in Appendix B. 7 of this text. Also, one should have knowledge of stacks usually discussed in a second semester computer science course or text such as Guide to Data Structures [3].

1.51 Memory Storage

In addition to integer, string, and Boolean data, programs can be written to store and process floating-point data. For example, in the C programming languages, floating-point data can be stored as a float which takes up 32 bits, as a double which uses 64 bits, or a long double which uses 80 bits. Using more bits allows for a larger range and more precision, but takes up more memory. In assembly language these three types can be declared as `real4`, `real8`, and `real10`, respectively, where the number indicates the number of bytes. Examples of declaring each are as follows:

```

x  real 4    ?
y  real8    3.14
z  real10   -27.1

```


1.52 Floating-Point Register Stack

As with integers, data must be loaded into a register from memory, processed, and then stored back into memory. However, instead of using the four general purpose registers `eax`, `ebx`, `ecx`, and `edx`, an eight-element floating-point register stack is used. The registers in the stack are called $ST(0)$ through $ST(7)$, where $ST(0)$ is the top of the stack as illustrated below:

This stack is different from the stack discussed in Chap. 6, where that stack is created in memory, can be declared to be different lengths, and to hold different types of data. The floating-point register stack is not declared in memory but is rather composed of registers in the CPU. Further it is used to store floating-point numbers of type `real4`, `real8`, and `real10`, and the stack has a fixed length of eight elements.

1.53 Pushing and Popping

Note that the instructions used to push and pop floating-point numbers cannot use immediate operands, so to illustrate these instructions, assume that the following have been declared:

```
num1 real4 2.71
num2 real4 3.14
num3 real4 ?
num4 real4 ?
```

First, it should be noted that all instructions that will be used to process floating-point data begin with the letter `f`. Initially, this may seem a little strange, but it helps distinguish the floating-point instructions from the instructions used with integers. In order to push the contents of a memory location onto the stack, a floating-point load instruction, `fld`, must be used. There are many variations on this instruction, but for the purposes of this text only this simple one will be used. To demonstrate how to take a copy of the contents of memory location `num1` and load, or push, it onto the stack the following would be used:

```
fld num1
```

which would cause the stack to appear as follows with the change indicated in green.

If `num 2` needed to be pushed onto the stack, the following would be used,

```
fld num2
```

and the stack would look as follows:

Note that the `2.71` is pushed further down on the stack and the `3.14` is now at the top of the stack. At first, the use of a floating-point load instruction may seem a little confusing because it might seem that a push instruction should be

used instead. But if one thinks of the process as loading a register which just happens to be the first position on a stack of registers, then with time it becomes a little more natural. What if one wanted to copy or pop a number off the top of a stack? Again, there are many variations, but two instructions will be shown here. The first is a floating-point store instruction which makes a copy of the number at the top of the stack and stores it in a memory location. For example, to make a copy of the number 3.14 at the top of the stack and store in num3, the following would be used,

```
fst num3
```

and the stack would look as follows:

Notice that there is no change on the stack and the number 3.14 is still at the top. This is because the fst instruction is not really a true pop instruction but is rather like a peek instruction that one might have learned about in a second semester computer science course or text. The number is not popped from the stack, but rather a copy is made of the item at the top of the stack and stored in the specified memory location. This is convenient at times, but can cause some confusion if one isn't careful.

Should one want to actually pop an item from the stack, the fstp must be used, where the p stands for pop. So, if the following instruction is executed:

```
fstp num3
```

the results would appear as follows:

In this example, note that the 3.14 is no longer on the stack and the 2.71 has moved up to the top. Lastly, executing the following, fstp num4 would cause the stack to be empty as shown below:

This is the typical result that one often expects when removing items from a stack and is also the instruction that will be used most often in this chapter.

1.54 Simple Arithmetic Expressions

In order to perform addition, the floating-point add instruction is used. As with the fld and fstp there are a number of variations but only the simplest ones will be shown here. The first is fadd which will take a number from memory and add it to the number at the top of the stack.

Given the previously declared variables with new values and an additional variable called answer as declared below,

```
num1 real4 2.5  
num2 real4 3.5  
num3 real4 4.0  
num4 real4 5.5  
answer real4 ?
```

how could the following be implemented?

```
answer = num1 + num2;
```

First, num1 would need to be pushed onto the stack. Could num2 be pushed instead? Yes, but it is much more convenient to keep the order from left to right to help with readability and consistency. Then using the fadd instruction, num2 would be added to the number at the top of the stack. Lastly, the sum would be popped off the stack and stored in answer. The resulting code would be as follows:

```
fld num1 ; push num1 onto the stack
fadd num2 ; add num2 to the first element on the stack
fstp answer ; pop the sum off the stack into answer
```

What about the following expression using subtraction?

```
answer = num1 - num2;
```

As might be suspected the corresponding instruction is fsub. However, does the order of operands matter in this case? In one sense yes, but in another no. Note that there exists an instruction which allows the reversal of the operands call fsubr where the *r* stands for reverse. However, again it helps consistency and readability to use the standard order of left to right and use the fsub instruction. The resulting code is as follows:

```
fld num1 ; push num1 onto the stack
fsub num2 ; subtract num2 from the first element on the stack
fstp answer ; pop the difference off the stack into answer
```

Multiplication works as addition but instead uses the fmul instruction. So, the following expression,

```
answer = num1 * num2;
```

would result in the following code segment:

```
fld num1 ; push num1 onto the stack
fmul num2 ; multiply num2 to the first element on the stack
fstp answer ; pop the product off the stack into answer
```

Likewise, division is similar to subtraction but instead the fdiv instruction is used. However, note that there is no reverse instruction, so care must be taken with the order of the operands and the following expression:

```
answer = num1 / num2;
```

would be implemented as follows:

```
fld num1 ; push num1 onto the stack
fdiv num2 ; divide num2 into the first element on the stack
fstp answer ; pop the quotient off the stack into answer
```

1.55 Complex Arithmetic Expressions

As done in the previous examples, one can take each equation encountered and proceed directly to trying to create the necessary assembly language code. With simple examples this might be acceptable, but with more complicated expressions, the chances of making mistakes increase dramatically. A far more efficient method is to first convert the expression to postfix as might have been discussed when learning about stacks from a second semester course or text. Then the postfix expression can be easily written in the necessary floating-point instructions.

At first the extra step of converting to postfix might seem to be slowing down the process since code is not being written immediately. However, the time needed to produce the code is quicker and much less prone to time-consuming debugging. Using the following variables and values,

w	real4	2.5
x	real4	3.5
y	real4	4.0
z	real4	5.25
answer	real4	?

consider the following C-like statement:

```
answer = w + x + y;
```

Recall that to do the conversion, first the variables are placed in the same order that they appeared in the original expression. Then since the addition on the left is done first, it is placed immediately after the two variables w and x. Then the second addition is placed after the variable y to add the sum of w and x and the variable y. The postfix equivalent of the right side of the expression is as follows:

```
w x + y +
```

Now to write the code of the corresponding expression going from left to right, the first step is to push any operands onto the stack. In this case the first two variables are pushed onto the stack as follows:

When encountering an operator, the operation should be performed on the top two numbers on the stack and the results pushed back onto the stack. To accomplish this task, the fadd instruction without an operand can be used. It takes the item on the top of the stack in ST(0) and adds it to the second item on the stack in ST(1). Then the first item is popped off the stack and the sum now occupies ST(0).

A question that might be asked at this point is what happened to the 3.5 that was popped off the stack? Where did it go? Since there is no operand in which to put the item, it simply is discarded.

Continuing, next the value in y is pushed onto the stack

Then as before when the fadd instruction is executed, the item at the top of the stack in ST(0) is added to the item in the second element in the stack in ST(1). Then the item at the top of the stack in ST(0) is popped off the stack and the second item in the stack in ST(1) moves up to the position in ST(0).

Lastly, the item on the top of the stack is popped off the stack and put into answer as follows:

Putting all the above together results in the following assembly language code segment:

```
fld w ; push w onto the stack
fld x ; push x onto the stack
fadd ; add the top two items leaving the sum at the top
fld y ; push y onto the stack
fadd ; add the top two items leaving the sum at the top
fstp answer ; pop the result off the stack into answer
```

In another more complex example that uses other operations, consider the following:

$\text{answer} = w / x - y * z;$

Remembering that multiplication and division have a higher priority than addition and subtraction and that in a tie the order is from left to right, the division is performed first, then the multiplication, and lastly the subtraction. The resulting postfix expression has the division after w and x, the multiplication after the y and z, and the subtraction is last after the quotient and product of the previous two operations as follows:

$w x / y z * -$

The resulting assembly code derived from the postfix expression is given below:

```
fld w ; push w onto the stack
fld x ; push x onto the stack
fdiv ; divide leaving quotient at top of the stack
fld y ; push y onto the stack
fld z ; push z onto the stack
fmul ; multiply leaving product at top of the stack
fsub ; subtract leaving difference at top of the stack
fstp answer ; pop the result off the stack into answer
```

As should be seen, the conversion from the postfix expression to the assembly language code is rather straight forward. However, if one is having difficulty converting an expression to postfix, be sure to go back to the text from a second semester course to review the process. Also, there are various practice exercises at the end of this chapter to help reinforce the process.

1.56 Mixing Floating-Point and Integers

Although it might be easier to not usually mix integers with floating-point numbers, on occasion it might be necessary. This section is not a comprehensive discussion of the matter, but gives a basic introduction so that one can be aware of some of the potential pitfalls when mixing these two types.

Since a floating-point number cannot be used with the general-purpose registers such as `eax`, it is instead necessary to use integers with the floating-point stack. In order to do so, the letter *i* (for integer) must be used after the letter *f* in a floating-point instruction. For example, instead of *fld*, the *fild* instruction would be used, and instead of *fstp*, the *fistp* instruction would be used, and so on. To illustrate how these instructions work, assume the existence of the following declarations:

```
one sdword 1
one2 real4 1.2
one5 real4 1.5
one7 real4 1.7
two5 real4 2.5
three sdword 3
ansint sdword ?
ansflt real4 ?
```

So, to transfer the integer contents of the variable `one` to the integer variable `ansint` using the floating-point stack, the following instructions would work:

```
fild one
fistp ansint
```

Although it would be much easier to use the `eax` register and `mov` instructions to transfer an integer, the above illustrates how the `fild` and `fistp` instructions work. Continuing, how could one move an integer to a floating-point variable? As one might suspect, the *i* is removed from the `fistp` instruction as follows:

```
fild one
fstp ansdbl
```

When moving a floating-point number to an integer, one needs to be careful with respect to rounding. Consider the following two sets of instructions from left to right that move the floating-point numbers 1.2 and 1.7, respectively.

```
fild one2 fild one7
fistp ansint fistp ansint
```

As one might suspect, the rounded results from left to right would be 1 and 2, respectively. What about moving the number 1.5 as follows?

```
fld one5
fistp ansint
```

As many have been taught, if the number is a 1.5, it should be rounded upwards. The result in the above example is 2, but not necessarily for the reason one might expect. Continuing, consider moving the number 2.5 in the following:

```
fld two5
fistp ansint
```

It would seem that the result in ansint would be 3, but in reality, it would be 2. Why? The reason is that the floating-point default mode for rounding a 0.5 is to round up or down to the nearest even number. In the case above the nearest even number to 2.5 is 2 and in the previous example that the nearest even number to 1.5 is 2. Although this may seem strange to some, it is a commonly used way of rounding numbers to avoid the accumulation of rounding errors and is known as bankers rounding. For example, adding the following numbers:

$$1.5 + 2.5 + 3.5 + 4.5 + 5.5$$

the answer would be 17.5. But what if all the above numbers were rounded up? The numbers would then be as follows:

$$2 + 3 + 4 + 5 + 6$$

The answer would be 20 which is significantly different than 17.5. But what if bankers rounding was used? The 1.5 would be rounded to 2 and the 2.5 would be rounded to 2, and so on, as follows:

$$2 + 2 + 4 + 4 + 6$$

The sum would be 18, which is much closer to the original answer of 17.5. Although this is a short list of numbers, the difference between using numbers that are rounded-up and the original numbers could be quite significant and the bankers rounding method helps mitigate this difference.

What if one didn't want to use rounding, but rather just truncate all numbers? This can be done by altering the floating-point control word, but this is beyond the scope of this text. Another solution is to use instructions that truncate, and this is accomplished by inserting the letter *t* just prior to the letter *p* in the fistp instruction so that it becomes the fisttp instruction. Taking the four real4 numbers previously declared above and using them from left to right with the fisttp instruction results in the following:

```
fld one2
fld one7
fld one5
fld two5
```

```
fisttp ansint
fisttp ansint
fisttp ansint
fisttp ansint
```

Instead of resulting in the answers 1, 2, 2, and 2 as would have happened with the `fistp` instruction, the above `fisttp` instruction would result in 1, 1, 1, and 2, respectively. Which instruction should be used? It depends on the application; however, rounding using the `fistp` instruction is probably going to be more useful.

How does all this affect arithmetic? Consider the following three code segments:

```
fld two5
fiadd three
fstp ansflt
```

```
fld two5
fiadd three
fistp ansint
```

```
fld one5
fiadd three
fistp ansint
```

The first one on the left adds 2.5 and 3, stores the result using `fstp`, and the floating-point answer is 5.5. The second example uses the same numbers, uses the integer `fistp` instruction instead, and the answer is 6 which is the closest even number to 5.5. However, the answer to the third one on the right adding 1.5 and 3 also using the integer `fistp` instruction rounds the 4.5 to 4 since it the closest even integer.

The result is that when mixing integers and floating-point numbers, it must be done carefully. If one is just placing some integer values into floating-point variables, then there probably isn't as much of a problem, but if floating-point results are being placed into integer variables, then extra care must be taken. The problem can only be exacerbated in more complex arithmetic expressions with intermediate results stored in integer variables. Instead of just coding and expecting the answers to be correct, it is important to walk through the code with key examples to ensure that the answers are what are expected. Of course, the simplest way to avoid many of these problems is to use just floating-point numbers with floating-point numbers, but when integers need to be used along with floating-point numbers, careful coding and testing is necessary.

1.57 Input/Output

Before continuing with some more complex code examples that utilize control structures and a complete program, the topic of input and output should be

discussed. As with integer input and output, it is helpful to first look at how floating-point input and output work in a C program.

1.57.1 float and real4

First, consider a simple program that inputs, transfers, and outputs a number of type float:

```
#include <stdio.h>
int main() {
    float x,y;
    printf("\n%s", "Enter a float for x: ");
    scanf_s("%f", &x);
    y = x;
    printf("\n%s%6.4f\n\n", "The float in y is: ", y);
    return 0;
}
```

If the number 3.14159 were input, the output would be as follows:

The float in y is: 3.1416

Note that in the `scanf_s` the `%f` indicates that a number of type float will be input. Also recall that just as with integers discussed in Chap. 2, an ampersand, `&`, must proceed the variable `x` on input to indicate the address where the number being input will be stored. On output the first number in the format prior to the decimal point indicates the total number of positions allocated to output the number, in this case 6. The second number after the decimal point indicates the number of positions that will be after the decimal point, in this case 4. Notice that since there are only four positions available for the five digits after the decimal point, the fourth position is rounded on output.

Although in assembly `real4` numbers can be input, they need to be moved to a `real8` variable before they can be output, as can be seen in the corresponding assembly code that follows:

```
. }68
    .model flat,c
    .stack 100h
printf PROTO arg1:Ptr Byte, printlist:VARARG
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
in1fmt byte "%f",0
msgofmt byte 0Ah,"%s",0
msg1fmt byte 0Ah,"%s%6.4f",0Ah,0Ah,0
msg0 byte "Enter a float for x: ",0
msg1 byte "The float in y is: ",0
x real4 ?
y real8 ?
```

```
.code
main proc
    INVOKE printf, ADDR msgofmt, ADDR msg0
    INVOKE scanf, ADDR in1fmt, ADDR x
    fld x
    fstp y
    INVOKE printf, ADDR msg1fmt, ADDR msg1, Y
    ret
main endp
end
```

The above is all rather similar to integer input and output. Again, as a reminder, notice the *ADDR* prior to the variable *x* in the *scanf* statement. Also, in the *msg 1 fmt* note the *%6.4f* formatting for the output. Further, notice that *x* is declared as real 4 and *y* as real 8 and the transfer of *x* to *y* via the stack to allow for the output of the floating-point number.

1.57.2 double and real8

Given this need to move from a real4 to a real8 for output, it might be more convenient to just use real8 numbers for both input and output unless there is a concern for conserving memory, especially with arrays. Consider the following C program using type double:

```
#include <stdio.h>
int main() {
    double x,y;
    printf("\n%s", "Enter a double for x: ");
    scanf_s("%lf", &x);
    y = x;
    printf("\n%s%6.4f\n\n", "The double in y is: ", y);
    return 0;
}
```

Notice above that to input a number of type double, *lf* is used. However, *lf* is not needed for output and a simple *f* is sufficient. Although the transfer from *x* to *y* is not really needed and the contents of *x* could be output, the assignment statement remains for consistency between the preceding and subsequent programs. The equivalent code in assembly language is as follows:

```
.686
        .model flat,c
        .stack 100h
printf PROTO arg1:Ptr Byte, printlist:VARARG
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
.data
in1fmt byte "%f",0
```

```

msgofmt byte 0Ah,"%s",0
msg1fmt byte 0Ah,"%s%6.4f",0Ah,0Ah,0
msg0 byte "Enter a float for x: ",0
msg1 byte "The float in y is: ",0
x real4 ?
y real8 ?
.code
main proc
    INVOKE printf, ADDR msgofmt, ADDR msg0
    INVOKE scanf, ADDR in1fmt, ADDR x
    fld x
    fstp y
    INVOKE printf, ADDR msg1fmt, ADDR msg1, y
    ret
main endp
end

```

The above is all rather similar to integer input and output. Again, as a reminder, notice the ADDR prior to the variable x in the scanf statement. Also, in the msg1fmt note the %6.4f formatting for the output. Further, notice that x is declared as real4 and y as real8 and the transfer of x to y via the stack to allow for the output of the floating-point number.

```

.686
.model flat,c
.stack 100h
printf PROTO arg1:Ptr Byte, printlist:VARARG
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
.data
in1fmt byte "%lf",0
msgofmt byte 0Ah,"%s",0
msg1fmt byte 0Ah,"%s%6.4f",0Ah,0Ah,0
msg0 byte "Enter a double for x: ",0
msg1 byte "The double in y is: ",0
msg2fmt byte 0Ah,"%s",0Ah,0Ah,0
x real8 ?
y real8 ?
.code
main proc
    INVOKE printf, ADDR msgofmt, ADDR msg0
    INVOKE scanf, ADDR in1fmt, ADDR x
    fld x
    fstp y
    INVOKE printf, ADDR msg1fmt, ADDR msg1, y
    ret
main endp
end

```

Again, as with the C program, the important thing to notice is that *lf* is used for input and *f* is used for output.

1.57.3 long double and real10

Although *real8* numbers should be sufficient for most processing, the input and output of *real10* numbers are discussed in passing should they be needed. As done previously, first consider a C program using the long double type below:

```
#include <stdio.h>
int main() {
    long double x,y;
    printf("\n%s", "Enter a long double for x: ");
    scanf_s("%Lf", &x);
    y = x;
    printf("\n%s%6.4Lf\n\n", "The long double in y is: ", y);
    return 0;
}
```

Probably the most important thing to notice in the above C program is that for both input and output the *Lf* format must be used. Also note that it uses an upper-case *L* and not a lower-case *l*. Since the code for *real10* in assembly language would be very similar to the previous two examples, it is left as an exercise at the end of this chapter. Again, since type *real8* is more convenient to use in output than type *real4* and has close to the same precision as *real10*, the type *real8* will be used in subsequent examples.

1.57.4 Inline Assembly

As will be seen in the next chapter, unfortunately inline assembly does not work with 64-bit integers. However, it can be used with all three floating point numbers: float, double, and long double. For example, the program from Sect. 10.7.2 can be rewritten using inline assembly for the transfer of *x* to *y* as follows:

```
#include <stdio.h>
int main() {
    double x,y;
    printf("\n%s", "Enter a double for x: ");
    scanf_s("%lf", &x);
    __asm {
        fld x
        fstp y
    }
    printf("\n%s%6.4f\n\n", "The double in y is: ", y);
    return 0;
}
```

As can be seen, the assembly code pushes the value of x onto the floating-point stack and then the value is popped off the stack into y . Although high-level control structures cannot be used with inline assembly, the advantage of inline assembly is that floating-point instructions and formatting can be tested without having quite the intricacy of assembly language input and output. Once they have been tested and debugged, they can be transferred and converted to an assembly language program.

1.58 Comparisons and Selection Structures

When comparing floating-point numbers using the 586 and earlier processors, there were various instructions, including the `fcom`, `fcomp`, `fcompp` instructions, which were rather cumbersome to use. However, starting with the 686 processors the additional less cumbersome `fcomi` and `fcomip` instructions were introduced. Whereas the original instructions needed to have the condition codes `C0`, `C2`, and `C3` copied into the Carry, Parity, and Zero flags, respectively, the new instructions set the flags automatically. Due to this advantage, these new instructions are the ones will be discussed in this text.

Starting with a simple comparison using an if-then structure, consider the following C program segment:

```
if(x > y)
    printf("\n%s\n", "x is greater than y");
```

The segment simply compares to floating-point numbers, x and y , and if x is greater than y it outputs the corresponding message. First, it should be noted that unfortunately the previous high-level control structures used with integers introduced in Chap. 4 cannot be used with floating-point numbers. This is because the high-level structures use the `cmp` instruction which uses implied integer subtraction. However, this should not pose too much of a problem because the `fcomi` and

`fcomip` instructions can be used with the label names such as `ifo1`: that were also introduced in Chap. 4. Another consideration is that although floating-point values are signed, instead of using signed conditional jumps such as `jl`, `jge`, `jle`, and `jbe`, unsigned conditional jumps such as `ja`, `jae`, `jb`, and `jbe` are used which branch based on the contents of the flags set by the `fcomi` and `fcomip` instructions. So, how do these instructions work? Looking first at the `fcomi` instruction, the two values that need to be compared are loaded onto the floating-point register stack. For example, using the C program from above the values x and y are loaded as follows:

```
fld y
fld x
```

Does it matter which value is loaded first? The answer is again yes and no. If the value on the right side of the conditional in the if statement is pushed on first, then the conditional jump statement that follows the `fcomi` statement will

need to be reversed as was done back in Chap. 4 when dealing with integers. However, if the order in which the two values are pushed onto the stack is opposite from the above order, then writing the unconditional jump statement will unfortunately tend to be more complicated. So, although reversing the order might initially seem to be more confusing, the resulting code is simpler and will be more consistent with the previous technique used with integers. Assuming x and y contained a 5.0 and 7.0, respectively, then the stack would look as follows:

Next the `fcomi` has the following form:

```
fcomi st(0),st(i)
```

It compares the zeroth item at the top of the stack with the i th item on the stack. Although any item on the stack can be used for the second parameter, this text will use the second item on the stack, which is `st(1)` as shown below:

```
fcomi st(0),st(1)
```

However, one problem with the above code is that although it would work correctly, it does not pop the values off the stack. This might not be a problem with a small program with only a few numbers on the stack, but if there are more than 8 items loaded onto the stack an exception can occur. So, it makes more sense to get into the habit of popping the values off the stack using the `fcomip` instruction as follows:

```
fcomip st(0),st(1)
```

Then since the original code in the C program was,

```
if(x > y)
```

then reversing the condition as has been done previously with integers, the corresponding conditional jump statement is as follows:

```
jbe endifo1
```

Putting all of the above pieces together along with the corresponding labels and output statement, the assembly code segment would be as follows:

```
    ; if x > y
ifo1: fld y
      fld x
      fcomip st(0), st(1)
      jbe endifo1
theno1: INVOKE printf, ADDR msg2fmt, ADDR msg2
endifo1: nop
```

In another example, consider the nested if-then-else-if structure in the following C code segment:

```

if (x > y)
    printf("\n%s\n", "x is greater than y");
else
    if(x < y)
        printf("\n%s\n", "x is less than y");

```

In this example, if x is greater than y or x is less than y , the corresponding message is output, but no message is output when they are equal. The first part of the assembly code segment for the outer if will be the same as before, but instead of branching to `endifo1`, it will branch to `elseo1`. Instead of the then section falling through to the `endifo1`, it will need to branch around the `elseo1` section to the `endifo1`.

```

        ; if x > y
ifo1: fld y
        fld x
        fcomip st(0), st(1)
        jbe elseo1
theno1: INVOKE printf, ADDR msg2fmt, ADDR msg2
        jmp endifo1
elseo1: nop
        ; nested if to be inserted here
endifo1: nop

```

In the else section there will be the nested if-then structure as follows:

```

; if x < Y
ifo2: jae endifo2
theno2: INVOKE printf, ADDR msg2fmt, ADDR msg3
endifo2: nop

```

Does there need to be another comparison in the nested if? And if so since the values were popped of the stack, do they need to be pushed back on the stack? In this case the answer to both questions is no because the flags are still set the same from the previous comparison. However, if the flags are altered by other instructions prior to the second if in the outer else section, then the answer would be yes to both questions. When in doubt, it doesn't hurt to reload and compare the values as part of the second if structure and this is done in the following segment:

```

                                ; if x > y
ifo1: fld y
        fld x
        fcomip st(0), st(1)
        jbe elseo1
theno1: INVOKE printf, ADDR msg2fmt, ADDR msg2
        jmp endifo1

```

```

else01: nop
      ;if x < y
ifo2: fld y
      fld x
      fcomip st(0), st(1)
      jae endifo2
then02: INVOKE printf, ADDR msg2fmt, ADDR msg3
endifo2: nop
endifo1: nop

```

1.59 Complete Program: Implementing an Iteration Structure

Having looked at a couple of selection structures, it is time to examine a loop structure using floating-point numbers as part of a complete program. First, are there any changes needed with respect to loops that use a counter for loop control such as the `repeat-untilcxz`, `repeat-until`, or `while-endw`? The answer is not really because typically floating-point numbers are not used as loop control variables, but rather integers are used. However, one might use floating-point numbers with a sentinel-controlled loop, but a `while-endw` cannot be used because it is implemented using a `cmp` instruction which uses implicit integer subtraction. Again, the while loop will need to be implemented using labels and the `fcomip` instruction. Starting with a C program, consider the following which continues to sum values until a zero or negative number is entered:

```

#include <stdio.h>
int main() {
    double x, sum;
    sum = 0.0;
    printf("\n%s", "Enter a positive double for x:");
    scanf_s("%lf", &x);
    while (x > 0.0) {
        sum = sum + x;
        printf("\n%s", "Enter a positive double for x: ");
        scanf_s("%lf", &x);
    }
    printf("\n%s%6.1f\n\n", "The sum is: ", sum);
    return 0;
}

```

First looking at just the while loop, the values to be compared must be loaded into the floating-point register stack. Again, the second value is loaded first and the conditional jump is reversed. Also, there needs to be an unconditional jump back to the beginning of the loop.

1.59. COMPLETE PROGRAM: IMPLEMENTING AN ITERATION STRUCTURE 153

```

        ;while x > 0.0
while01: fld zero
        fld x
        fcomip st(0), st(1)
        jbe endw01
        ; body of loop
        jmp while01
endw01: nop

```

Of course, without a change in the loop control variable, there would be an infinite loop. As discussed in Sect. 5.4 previously and shown in the C program above, there will be a priming read prior to the loop and it will be repeated as the last instruction in the body of the loop. Putting all the pieces together results in the following complete program:

```

        . }68
        .model flat,c
        .stack 100h
printf PROTO arg1:Ptr Byte, printlist:VARARG
scanf PROTO arg2:Ptr Byte, inputlist:VARARG
in1fmt byte "%lf",0
msgofmt byte 0Ah,"%s",0
msg1fmt byte 0Ah,"%s%6.1f",0Ah,0
msg0 byte "Enter a positive double for x: ",0
msg1 byte "The sum is: ",0
x real8 ?
sum real8 ?
zero real8 0.0
        .code
main proc
        ;sum = 0.0
        fld zero
        fstp sum
        INVOKE printf, ADDR msgofmt, ADDR msg0
        INVOKE scanf, ADDR in1fmt, ADDR x
        ;while x > 0.0
while01: fld zero
        fld x
        fcomip st(0), st(1)
        jbe endw01
        ; sum = sum + x
        fld sum
        fadd x
        fstp sum
        INVOKE printf, ADDR msgofmt, ADDR msg0
        INVOKE scanf, ADDR in1fmt, ADDR x

```

```

    jmp while01
endw01: nop
    INVOKE printf, ADDR msg1fmt, ADDR msg1,
sum
    ret
main endp
end

```

1.60 Complete Program: Implementing an Array

Just as with integers, it is possible to create an array of floating-point numbers. To keep the complete program simple, assume the existence of the following five-element array:

```
array real8 1.0, 2.3, 3.7, 4.9, 5.1
```

The array could have been declared as follows:

```
array real8 5 dup(?)
```

and then the data could be input, but this is left as an exercise at the end of the chapter. In either case, in order to output the contents of the array, the number of elements could be stored in a variable n and since there is a fixed number of elements a post-test .repeat-.untilcxz loop could be used. Using the esi register, the address of the array would need to be placed into that register prior to the loop. Also, instead of incrementing the esi register by 4, it would need to be incremented by 8 because each element of the array is eight bytes long. The basic loop structure is as follows:

```

lea esi,array
mov ecx,n
.repeat
; body of loop
add esi,8
.untilcxz

```

An INVOKE statement can be used to output the contents of each element. However, one cannot just output the contents of [esi]. Instead the INVOKE printf statement needs to know the size of each element so real8 ptr [esi] must be used as follows:

```
INVOKE printf, ADDR msg2fmt, real8 ptr [esi]
```

Lastly, recall that the INVOKE statement destroys a number of registers including the ecx register, so it must be saved before and restored after the call. A temporary variable could be used or it could be pushed onto the stack which is what is done in the instance. The code segment is as follows:

```

lea esi,array
mov ecx,n
.repeat
push ecx
INVOKE printf, ADDR msg2fmt, real8 ptr [esi]
add esi,8
pop ecx
.untilcxz

```

Could the same be done with the ebx register? Yes, where the loop control is the same, but instead ebx is initialized to a zero and the INVOKE statement uses array[ebx] instead. As before ebx is incremented by eight and the ecx still needs to be pushed and popped as shown in the following segment:

```

mov ebx,0
mov ecx,n
.repeat
push ecx
INVOKE printf, ADDR msg2fmt, array[ebx]
pop ecx
add ebx,8
.untilcxz

```

Of the two methods the second one is probably the simpler due to the less complicated INVOKE statement and mimicking the more common indexing method used in high-level languages. The complete program outputting the array twice using the two different registers and including headings prior to the two sets of output is shown below:

```

.686
.model flat, c
.stack 100h
printf PROTO arg1:Ptr Byte, printlist:VARARG
.data
msg1fmt byte 0Ah,"%s",0Ah,0Ah,0
msg2fmt byte " %3.1f",0Ah,0
msg1 byte "Array",0
n sdword 5
array real8 1.0, 2.3, 3.7, 4.9, 5.1
.code
main proc
INVOKE printf, ADDR msg1fmt, ADDR msg1
lea esi,array
mov ecx,n
.repeat
push ecx
INVOKE printf, ADDR msg2fmt, real8 ptr [esi]

```

```
pop ecx
add esi,8
.untilcxz

INVOKE printf, ADDR msg1fmt, ADDR msg1
mov ebx,0
mov ecx,n
.repeat
push ecx
INVOKE printf, ADDR msg2fmt, array[ebx]
pop ecx
add ebx,8
.untilcxz
ret
main endp
end
```

1.61 Summary

- Use `f` for outputting Real4 and Real8 numbers, and use `Lf` for outputting Real10 numbers.
- Just like with integers, use an `&` prior to the variable when inputting floating-point numbers.
- Remember to use `f` for inputting Real4 numbers, `lf` for Real8 numbers, and `Lf` for Real10 numbers.
- Remember when using the `fistp` instruction the answer is rounded to the nearest even number and when using the `fisttp` instruction the answers are truncated.
- Although the `fsubr` instruction could be used when performing subtraction with numbers in reverse order, it is easier to push the numbers onto the floating-point stack from left to right and use the `fsub` instruction instead.
- Note that floating-point variables cannot be used to control high-level structures because the `cmp` instruction uses implied integer subtraction, so the floating-point compare (`fcompi`) and unsigned conditional jumps (`ja`, `jb`, etc.) must be used to create the control structures.
- Although the order of pushing numbers onto the floating-point stack is not really important when using the `fcompi` instruction, the resulting code is easier to write when the second number being compared is pushed on first and the conditional jump is reversed as done previously with integers.

- Depending on which is used to declare an array, `real4`, `real8`, `real10`, be sure to increment the index by 4, 8, or 10, respectively.

1.62 Exercises (Items Marked with an * Have Solutions in Appendix D)

1. Given the following variables, what are the results in the variable z for each of the following code segments?

1.63 Exercises (Items Marked with an * Have Solutions in Appendix D)

1. Given the following variables, what are the results in the variable z for each of the following code segments?

```
w real8 2.0
x real8 5.5
y real8 6.5
z sdword ?
```

*a. `fld w`

b. `fld y`

c. `fld y`

d. `fld w`

`fld y`

`fld w`

`fld w`

`fld y`

`fadd`

`fdiv`

`fsub`

`fadd`

`fistp z`

`fld x`

`fistp z`

`fist z`

`fadd`

`fld w`

`fisttp z`

`fadd`

`fistp z`

2. Convert the following C-like arithmetic instructions into post-fix form and then write the corresponding assembly language instructions. Assume that all variables are of type real8.

```
*a. answer = x - y + z;
b. result = (w + x) / (y - z);
c. info  = a / b * c - d;
d. data  = i * j + (k / (m - n));
```

3. Can pre-fix form of an arithmetic statement be used to create the corresponding assembly language instructions? If not, why not? If so, how? For either answer, show how it can or cannot be done by giving an example of a conversion to pre-fix and the corresponding assembly instructions.
4. Write the equivalent assembly language code segment for the C program in Sect. 10.7.3 (which uses type long double).
5. Given the code using inline assembly in Sect. 10.7.4, rewrite it to work with float and long double types (Hint: For type float, see Sect. 10.7.1).
6. Alter the if-then code segment at the beginning of Sect. 10.8 to implement an if-then-else structure to output the message “ x is less than or equal to y ” in the else section.
7. Alter the if-then-else-if code segment in Sect. 10.8 to add an else section to the nested if statement to output the message “ x and y are equal”.
8. Change the complete program in Sect. 10.9 to implement a do-while structure (post-test loop) instead of while structure (pre-test loop). Make sure that it works properly for 0.0 and negative numbers.
9. Change the complete program in Sect. 10.10 to instead prompt for and input the five numbers.

1.64 -Bit Processing

To store larger 64-bit numbers in a high-level language such as Java, variables are declared as long, or in the C programming language long long is used. In assembly language a quad word, qword, or signed quad word, sqword, is used. Just as a 32-bit double word is twice as big as a regular 16-bit word, a 64-bit quad word is two times larger than a 32-bit double word or four times larger than a 16-bit word, thus the name quad word. In comparison with double words, Table 11.1 indicates the range of numbers that are capable of being stored in quad words. Note that in comparison with a 32-bit double word which can store a signed positive number of approximately two billion, a 64-bit quad word can store a signed positive number of approximately 9 quintillion. Of course, with the advantage of being able to store larger numbers, there is the disadvantage

of using more memory. Although using a single variable of 8 bytes instead of 4 bytes does not seem to be much of a disadvantage, this increase in memory is especially noticeable when using large arrays.

There are two other disadvantages of using 64-bit registers. First is that they cannot be used in in-line assembly. Although this presents some inconvenience, the focus of this text has been to use stand-alone assembly. A greater inconvenience is that the high-level directives such as `.if`, `.while`, and so on, cannot be used in 64-bit mode. However, as discussed in Chap. 4, it is possible to mimic high-level language structures and some examples will be given in Sect. 11.6.

1.65 Four General Purpose Registers

Of course, without registers to store 64-bit numbers, 64-bit processing would be difficult. Just as the letter *e* was added to indicate the 32-bit version of the *ax* register, the letter *r* is added to the *ax* register to indicate the 64-bit version. The

Table 11.1 Types, number of bits, and range of values

Type	Number of bits	Range (inclusive)
sqword sdword	64	−9,223,372,036,854,775,808 to +9,223,372,036,854,775,807
qword	64	0 to +18,446,744,073,709,551,615
sdword	32	−2,147,483,648 to +2,147,483,647
dword	32	0 to +4,294,967,295

following is a variation of Fig. 1.4 to show the relationship between the 8-bit *al* register, 16-bit *ax* register, 32-bit *eax* register, and 64-bit *rax* register.

As before, the same diagram could be used to show the relationship with the *rbx*, *rcx*, and *rdx* registers by substituting the letters *b*, *c*, and *d* in place of the letter *a*.

Although the above appears fairly-straight forward, there are some interesting similarities and differences in the relationship between the different size registers. Recall that in 32-bit mode when the *al* register is altered, it does not alter the upper bits of the *ax* register and when the *ax* register is altered, the upper bits of the *eax* register are not altered. The same is true for these registers in 64-bit mode. For example, assuming that the *ax* register contained the following in hexadecimal,

ax F7E5

and the following instruction was executed,

```
mov al,00h
```

then only the lower 8 bits would be altered as follows:

ax F700

Likewise, if eax contained the following,

eax
FEDCBA98

and the following instruction was executed,

mov ax,0

the eax register would look as follows:

eax FEDC0000

However, when altering eax, the equivalent is not true in rax. If the rax register contained the following,

rax 0123456789ABCDEF

and the following instruction clearing eax to 0 was executed:

mov eax,0

the entire rax register would be cleared to 0 as follows:

rax
0000000000000000

What has happened is that zero was moved into the eax register which is the lower 32 bits of the rax register, and then the upper 32 bits are cleared to zero. Using the original number in the rax register from above, consider the following example:

mov eax,-1

Since a negative one is represented as FFFFFFFFh in eax, the contents of eax are copied into the lower 32 bits of rax and the upper 32 bits of the rax register are again cleared to zero as shown below:

rax
00000000FFFFFFFF

What if one wanted to have the above result in all F's in hex to indicate a negative one? To accomplish this task the movsxd instruction would be used. It copies from a 32-bit memory location or register to the lower 32 bits of a 64-bit register and then extends or propagates the sign-bit from the leftmost position of the lower 32 bits into the upper 32 bits of the 64-bit register. Further, instead of initially placing the data into eax as in this case, it could be placed into another register such as ebx. From there the movsxd would copy the contents of ebx into rax as follows:


```
mov ebx,-1
movsxd rax,ebx
```

This would result in the following:

```
rax FFFFFFFFFFFFFFFFFF
```

What would happen with the segment below?

```
mov ebx,0
movsxd rax,ebx
```

In this case the results would look the same as when using the `mov` instruction, although the process is different. In the case of the `mov` instruction the upper bits are cleared to zero and in the `movsxd` case the sign-bit is propagated into the upper 32 -bits. If the end result is the same, why would one use `movsxd`? The reason is that if there is unknown data in a 32-bit register or memory location then the `movsxd` would extend the sign-bit, whether a zero or a one, accordingly. For example, assume the variable `num32` contains a signed integer:

```
mov ebx, num32
movsxd rax,ebx
```

Then regardless of what is in the sign-bit in `num32`, it will be extended into the upper 32 bits of `rax`. This principle also applies to the other three general purpose registers as well: `rbx`, `rcx`, and `rdx`.

Does the default of clearing the upper 32 bits of a 64-bit register cause any difficulties? If one is using just 32 bits in a 64-bit environment, then there is really no concern what is happening in the upper 32 bits of `rax` or the other registers. And if one is using only 64 -bit registers and instructions then again this should be of no concern. Only if one is moving from the 32 -bit environment to the 64 -bit environment and is accustomed to the upper bits of a register being maintained or if one is mixing 32 -bit and 64 -bit registers and processing, then it could cause some difficulties. However, provided one is aware of the differences and is careful when choosing the correct instructions such as `movsxd`, then it should not pose much of a problem.

1.66 Other 64-Bit Registers

In addition to the four general purpose registers, other registers also have 64-bit variations. These include the index registers `esi` and `edi` which are available as `rsi` and `rdi`, and the stack registers `esp` and `ebp` as `rsp` and `rbp`, respectively. The `eflags` register also has a 64-bit counterpart as `rflags`, but although this could allow for more error flag bits in the future, at present the upper 32 bits of this register are not used. Also, the instruction pointer register `eip` has a corresponding 64-bit register `rip`.

Table 1.8: Summary of 64-bit registers

64-bit registers	Name	32-, 16-, and 8-bit sub-registers	Brief description and/or primary use
rax	Accumulator	eax, ax, al	Arithmetic and logic
rbx	Base	ebx, bx, bl	Arrays
rcx	Counter	ecx, cx, cl	Loops
rdx	Data	edx, dx, dl	Arithmetic
rsi	Source index	esi, si	Strings and arrays
rdi	Destination index	edi, di	Strings and arrays
rsp	Stack pointer	esp, sp	Top of stack

Table 11.2 Summary of 64-bit registers

Lastly, in addition to the 64-bit variations of the previous 32-bit registers, there are also eight new numbered 64-bit registers available. They are $r8$ through $r15$, which each have corresponding sub-registers. These registers include the 32-bit double word registers $r8d$ through $r15d$ where the d stands for double word, the 16-bit word registers $r8w$ through $r15w$ where the w stands for word, and the 8-bit byte registers $r8b$ through $r15b$ where the b stands for byte.

Note that whereas the upper 8 bits of the original 16-bit registers are directly accessible, such as ah in the ax register, there is no name for the upper 8 bits in the corresponding $r8w$ through $r15w$ registers. Given this, how does one access these bits if there is no corresponding name? It is no different than when there is no name for the upper 16 bits of a 32-bit register or no name for the upper 32 bits of a 64-bit register. The solution is to use logic and shifting instructions.

Also, note that although registers such as $r8b$ can be used with registers such as al , they cannot be used in the same instruction with registers such as ah . Then, how does one move information from ah to $r8b$? Again, the answer is logic and shifting instruction from Chap. 6 and as will be shown in Sect. 11.5 of this chapter. For convenience, all the registers discussed in this section are summarized in Table 11.2.

1.66.1 -Bit Integer Output

As was mentioned earlier, 64-bit processing cannot be used with in-line assembly. This is unfortunate because it would make it easy to use C input/output instructions to test various 64-bit code segments. Further, the technique for inputting and outputting 64-bit integers is also different than the one used previously for 32-bit integers and this section gives examples of simple I/O to facilitate testing of code presented in this chapter.

As before, a simple program to output a 64-bit integer in C is given first followed by the equivalent program in assembly language. Consider the following C program which is similar to the program at the end of Sect. 2.3 except the variables are declared as type `long long` instead of type `int` and formatting using `%11d` instead of `%d`, respectively:

```
#include <stdio.h> int main() {
```

```

long long num1=5, num2=7;
printf("\n%lld%s%lld\n\n",num1," is not equal to ",num2);
return 0;
}

```

In order to implement the above program using 64-bit output, there have been some changes made compared to 32-bit processing. Along with other important information in the directions for using Visual Studio in Appendix A.3, be sure to change the Solutions Platforms in the ribbon in from *x86* to *x64*.

In the program itself, note that the following statements are no longer needed and can simply be deleted from the program.

```

. }68
.model flat, c
.stack 100h

```

As will be discussed shortly, input and output is no longer done via parameters, so it is not necessary to include the parameter list after the `PROTO` statements as done previously and the new statements are as shown below:

```

printf PROTO
scanf PROTO

```

The data segment is similar to what has been used previously and appears as follows:

```

msglfmt byte 0Ah,"%lld%s%lld",0Ah,0Ah,0
msg1 byte " is not equal to ",0
num1 sqword 5
num2 sqword 7

```

The only change to notice is that instead of the variables being defined as `sdword`, they are defined as `sqword`. Also note the change in formatting from `%d` to `%lld`.

If parameters are not used as mentioned above, then how is information communicated to the input and output procedures? The information is sent either through registers or the stack, where it is more convenient to use registers. This is because the stack needs to be adjusted before and after calling a `printf` or `scanf` to allow for what is known as shadow space. This space is used for storing parameters and should be at least 32 bytes in length although larger amounts to ensure sufficient space can be used such as the 64 bytes in this text. This is accomplished by adjusting the stack pointer by subtracting 64(40h) from the stack pointer prior to the call and adding back 64(40h) after the call. Also note that the `INVOKE` statement is not used but rather the `CALL` statement is used instead. So, a call to a `printf` statement will appear as follows:

```

sub rsp,40h
CALL printf
add rsp, 40h

```

To communicate with the `printf` routine using registers, the information is sent via the `rcx`, `rdx`, `r8`, and `r9` registers. Since these registers are being used for communications, any information in these registers can be lost during the I/O process. Further, the contents of the `rax` and `r9` through `r11` are altered by the `printf` statement. So if these registers are being used, it is important to save their contents prior to any I/O and have them returned to the registers afterwards. However, recall back in Sect. 1.6 that unless speed is a very important concern, it was recommended not to keep data in the registers in the first place and that it be returned to memory to avoid such complications. Again, it is easier to optimize a correctly implemented program than to get a supposedly optimized program to work correctly.

Prior to calling the `printf` statement, the registers need to be loaded with the format and data in the same order that was used previously with parameters. Instead of sending the address of the format statement or strings using `ADDR`, the address is loaded into the registers using `mov` and `offset` or using the `lea` instruction. In this first example `mov offset` is used but in later examples `lea` is used. So, the format statement can be loaded first into `rax` followed in order with the data in the subsequent registers. The resulting code would look as follows:

```
mov rcx,offset msglfmt
mov rdx,num1
mov r8,offset msg1
mov r9, num2
```

Putting everything together results in the following program:

```
printf PROTO
scanf PROTO
.data
msglfmt byte 0Ah,"%lld%s%lld",0Ah,0Ah,0
msg1 byte " is not equal to ",0
num1 sqword 5
num2 sqword 7
main proc
mov rcx,offset msglfmt
mov rdx, num1
mov r8,offset msg1
mov r9, num2
sub rsp, 40h
CALL printf
add rsp, 40h
ret
main endp
end
```

But what if more items need to be output in one line and only four registers are being used? For example, consider the following modification of the preceding C program code segment which outputs three numbers instead and uses "!=" to save space:

```
#include <stdio.h> int main() {
    long long num1 = 5, num2 = 7, num3 = 9;
    printf("\n%lld%s%lld%s%lld\n\n",num1," != ",num2," != ",num3);
    return 0;
}
```

The problem is that in addition to the format there are five additional items to output on one line and there are only four registers. Although the stack could be used, it is easier to move the last two items and their formatting along with the "" to a second printf statement so that all of the output stays on one line as follows:

```
#include <stdio.h>
int main(){
    long long num1 = 5, num2 = 7, num3 = 9;
    printf("\n%lld%s%lld",num1," != ",num2);
    printf("%s%lld\n\n", " != ",num3);
    return 0;
}
```

The corresponding assembly code would be as follows:

```
printf PROTO
scanf PROTO
.data
msglfmt byte 0Ah,"%lld%s%lld",0
msg2fmt byte "%s%lld",0Ah,0Ah,0
msg1 byte " != ",0
num1 sqword 5
num2 sqword 7
num3 sqword }
.code
main proc
    mov rcx,offset msglfmt
    mov rdx, num1
    mov r8,offset msg1
    mov r9, num2
    sub rsp, 40h
    CALL printf
    add rsp, 40h
    mov rcx,offset msg2fmt
    mov rdx,offset msg1
```

```

    mov r8,num3
    sub rsp, 40h
    CALL printf
    add rsp, 40h
    ret
main endp
end

```

Although this takes up a little more code, it is relatively simple to implement and avoids any complications that might occur when using the stack. As an aside, all of the other types of formats in C can be used by attaching the letters 11 after the % and before the format letter. For example, should one want to output the contents of a 64-bit memory location as unsigned, %llu should be used, and to output in hexadecimal, %lx should be used.

1.66.2 -Bit Integer Input

Input is similar to output and using a C program that is similar to the one at the end of Sect. 2.4, consider the following program that inputs number, adds 1 to it, and then outputs number:

```

#include <stdio.h>
int main() {
    long long number;
    printf("\n%s", "Enter an integer: ");
    scanf("%lld",&number);
    number++;
    printf("\n%s%lld\n\n", "The integer is: ", number);
    return 0;
}

```

As before, there needs to be a prompt prior to the input and also the number input needs to be output, both of which are relatively simple given in the last section. The new part of interest here is the scanf. First, the address of the input format is loaded into rcx and then the address of the location number loaded into the rdx register. In both cases, the mov and offset or the lea instruction should be used. Notice that the stack has to be adjusted before and after the call just as it was with the printf statement all as shown below:

```

lea rcx, in1fmt
lea rdx, number
sub rsp, 40h
CALL scanf
add rsp, 40h

```

The instruction to increment number by 1 is the same as with the 32-bit instruction as follows:

inc number

In fact, many of the instructions will be very similar between 32-bit and 64-bit processing as will be discussed in the next section. Further, it should be noted that the scanf alters the same registers, rax, rcx, rdx, and r8 through r11, as the printf. So again, care must be taken by the programmer to not leave any relevant data in these registers. The program to implement the above C program can be found below:

```
printf PROTO
scanf PROTO
.data
in1fmt byte "%lld",0
msgofmt byte 0Ah,"%s",0
msg1fmt byte 0Ah,"%s%lld",0Ah,0Ah,0
msg0 byte "Enter an integer: ",0
msg1 byte "The integer is: ",0
number sqword ?
.code
main proc
lea rcx, msgofmt
lea rdx, msg0
sub rsp,40h
CALL printf
add rsp,40h
lea rcx, in1fmt
lea rdx, number
sub rsp,40h
CALL scanf
add rsp, 40h
inc number
lea rcx, msg1fmt
lea rdx, msg1
mov r8,number
sub rsp, 40h
CALL printf
add rsp, 40h
ret
main endp
end
```

1.67 Logic and Arithmetic Applications

Although as seen in the previous section, much of the processing using 64-bit registers and memory locations are the same as with 32-bit registers, it does

not hurt to give some examples of code to illustrate this similarity. Also, of particular interest is to show some applications and the interaction between 32-bit and 64-bit processing.

1.67.1 Shift and Rotate

As discussed in Sect. 11.2, instructions cannot access both a register such as `ah` in the `rax` register and `r8b` in the `r8` register. In other words, the instruction `mov r8b, ah` is illegal. Although it would be nice to have such an instruction for the

relatively rare occurrences that this type a transfer is necessary, it is really just a minor inconvenience. Since such an instruction does not exist, how would one accomplish it? Remembering various instructions from Chap. 6, shifting or rotating could be used. First the data in `ah`, which are bits 8 through 15 of the `rax` register, would need to be moved to the `al` register. A simple `mov al, ah` could be used, but this would destroy the contents of the `al` register. If that data is no longer needed then this method would not be a problem.

However, what if that data needed to be kept? It could be stored in another register or a memory location, but that might destroy other data or require the use of memory. So, how could it be done without having to use any other registers or memory locations? An alternative is to shift the contents of the `rax` register 8 bits to the right. But again, the data that was previously in `al` would be destroyed. The solution is to rotate the contents of the `rax` register to the right 8 bits, which would cause the contents of the lower 8 bits to be rotated to the upper 8 bits in bit positions 56 through 63 of `rax`. For example, if `rax` originally contained the following:

And the instruction `ror rax, 8` was executed, then `rax` would look as follows:

EF0123456789ABCD

Then, the contents of `al` could be transferred to `r8b`. Further, if the contents of `rax` could be restored back by rotating it to the left 8 bits. This set of instructions is as follows:

```
ror rax,8
mov r8b,al
rol rax,8
```

But why rotate all 64 bits of the `rax` register? Alternatively, only the `ax` subset of the `rax` register could be rotated by 8 bits as shown below:

Then the contents of `ax` could be rotated back and the code for this is solution is as follows:

```
ror ax,8
mov r8b,al
rol ax,8
```


1.67.2 Logic

Also as discussed in Sect. 11.1, two alternatives of moving 32 bits into a 64-bit register resulted in either the upper- 32 bits becoming all zeros or extending the sign bit into the upper 32 bits. What if one did not want to alter the upper 32 bits of a 64-bit register? The way this can be accomplished is by using various logic instructions. This would involve clearing out the lower 32 bits of the 64-bit register to zero, creating a mask to not alter the upper 32 bits, and using an and instruction to put the new data into the lower 32 bits. Again, assume the data in the rax register is as follows:

```
rax
0123456789ABCDEF
```

The mask necessary to clear out the lower 32 bits yet maintain the upper 32 bits is as follows:

```
FFFFFFFF00000000
```

At first, one might be tempted to put the above in an operand of an and instruction, but it is too large. Instead it can be placed into a memory location using an unsigned qword. Remember when using hexadecimal number that the left most position of the number must be a digit and not a letter, so note the leading 0 as follows:

```
mask64 qword 0FFFFFFFF00000000h
```

Then an and instruction can be used,

```
and rax,mask64
```

which results in the following:

Now if the new data to be put into the lower 32 bits was AABBBCCDDh, it could be loaded into a 32-bit register such as ebx,

```
mov ebx,0AABBBCCDDh
```

which would automatically clear out the upper bits of rbx as follows:

And lastly, the data in rbx could be placed into the lower 32 bits of rax using an or instruction as follows:

```
or rax,rbx
```

which would result in the new data in the lower 32 bits while the original data in the upper 32 bits would remain undisturbed as shown below:

```
rax 01234567AABBBCCDD
```

The completed code segment is as follows:

```
and rax,mask64
mov ebx,0AABBCCDDh
or rax,rbx
```

The result is that if a particular instruction does not exist to do precisely what needs to be done, it is usually not a problem because given the vast number of other instructions available, the appropriate task can be accomplished. If necessary, the resulting group of instructions can be placed into a macro using parameters as introduced in Chap. 7 and reviewed in Sect. 11.8. Further, this is left as an exercise at the end of the chapter.

1.67.3 Arithmetic

The use of 64-bit arithmetic is very similar to 32-bit arithmetic except much larger numbers can be used. Whereas there is very little difference with addition and subtraction, there are a few differences with multiplication and division.

Recall from Chap. 3 that when two 32-bit numbers are multiplied the result occupied 64 bits in the `edx: eax` register pair. A similar result occurs when two 64-bit numbers are multiplied where the results are placed in the 128 bits of the `rdx: rax` register pair. Consider the following code example:

```
mov rax,2147483647
mov rbx,2
imul rbx
```

The above number loaded into the `rax` number is the largest signed positive number that can be stored in 32 bits. If 32-bit multiplication were to be used the results would have extended into the `edx` register. However, when using 64-bit registers, the resulting product can be contained in the 64-bit `rax` register and only the sign-bit would be propagated into the `rdx` register. If the number were much larger than this, then it too would extend into the `rdx` register, but in this text, 64 bits should be more than sufficient.

Similar to 32-bit division where the `edx: eax` register pair contains the 64-bit dividend, 64-bit division needs to have the dividend initially in the 128-bit `rdx: rax` register pair. Instead of using the `cdq` (convert double to quad) instruction prior to the 64-bit division, one must use the `cqo` (convert quad to oct, where oct refers to eight words) instruction. This extends the sign bit of `rax` into `rdx` prior to 64-bit division as illustrated in the following code segment:

```
mov rax,7fffffffffffffffh
cqo
mov rbx,2
idiv rbx
```

A hexadecimal number is used to show that the largest possible positive number is being stored in the `rax` register where bit 63 is a 0. Similar to the previous 32-bit division, the quotient is stored in `rax` and the remainder is stored in `rdx`.

1.68 Control Structures

As mentioned at the beginning of this chapter, unfortunately high-level directives do not work in 64-bit mode. Fortunately, most of the ways that control structures can be created using comparisons and branches in 32-bit mode can also be done in 64-bit mode. However, there are some potential areas that could cause problems when mixing 32-bit and 64-bit data. For example, consider the following code segment which loads a number into `eax` and then compares a number in `rax` register with the number 2. It then stores in `rdx` the address of `msg3` indicating that it is equal or `msg4` indicating that it is not equal for later output:

```
; **** Caution: possibly incorrectly implemented code ****
mov eax,2
; if rax == 2
ifo1: cmp rax,2
jne else01
then01: lea rdx, msg3
jmp endifo1
else01: lea rdx,msg4
endifo1: nop
```

As before, the condition is reversed using `jne` to branch to `else01` if not equal and falls through to `then01` if it is equal. Yes, it is probably not a good idea to initially load `eax` and then compare using `rax`, but assume for the time being that the data to be compared was originally in the `eax` register. The above code segment works okay in this particular instance, but what is a potential problem with the code segment? Instead, consider if `eax` contained a -2 and was compared to a -2 instead as shown below:

```
; **** Caution: incorrectly implemented code ****
mov eax,-2
; if rax == -2
ifo1: cmp rax,2
jne else01
then01: lea rdx, msg3
jmp endifo1
else01: lea rdx,msg4
endifo1: nop
```

Although the previous code segment worked, would the above code work? The answer is no because recall from Sect. 11.1 that when data is moved into the `eax` register or related registers, the contents of the upper 32 bits of `rax` is set to zeros. The result is that the negative number in `eax` is a positive number in `rax`. To correct this problem, the `movsxd` instruction should be inserted in the second line to extend the sign bit into the upper 32 bits as follows:

```

mov eax,-2
movsxd rax,eax
; if rax == -2
ifo1: cmp rax,2
jne else01
then01: lea rdx, msg3
      jmp endif01
else01: lea rdx,msg4
endif01: nop

```

Although it was technically not needed in the initial code using positive numbers, it should probably be inserted there as well to ensure that if a negative number is input instead of assigned, the code would work correctly.

In another example, what if one wants to use the loop instruction with 64-bits? In the following segment the rax register would be loaded with the appropriate number.

```

      mov rcx, 1000000000h
for01: nop
      ;body of loop
      loop for1
endfor01: nop

```

In the above case a hexadecimal number is being used to show that the number in rcx is larger than 32 bits. However, note that the loop will iterate over four billion times and it might be slow depending on the speed of the computer being used. Further, even larger numbers can be used with 64 bits in the above code segment and the loop could take a significant amount of time to execute.

As with preceding examples using an if statement and mixing 32-bit and 64-bit registers, care must be taken. Although one would not want to load a negative number into the rcx register using a loop instruction, a similar error could occur with other forms of loop control. Consider for example the following loop that starts with a negative number in ecx, adds one to it until it reaches zero, and the count of the number of iterations is kept in rax.

```

; **** Caution: incorrectly implemented code ****
mov rax,0
mov ecx, -3
; while ecx < 0
while01: cmp rcx,0
      jge wendo1
      inc rax
      inc ecx
      jmp while01
wendo1: nop

```

Although ecx is a negative number there are all zeros in the upper 32 bits of the rcx register. The result is that rcx is a positive number and the count of iterations would be zero. If the movsxd instruction is added, note that there is yet

another problem with the above code segment. Notice that `ecx` is being incremented in the body of the loop instead of `rcx` which again causes the upper 32 bits to be zero and the loop only iterates once. The corrected code is as follows:

```
mov rax,0
mov ecx, -3
movsxd rcx,ecx
; while ecx < 0
while01: cmp rcx,0
        jge wend01
inc rax
inc rcx
jmp while01
wend01: nop
```

The simple solution to such problems is to avoid mixing 32-bit and 64-bit registers. However, if one must use both, be careful that the code works the way one intended by running sufficient tests.

1.69 Arrays

Indexing 64-bit arrays is similar to indexing 32-bit arrays. Instead of declaring the array with `sdword`, they are declared using `sqword`. The code is very similar except that 64-bit registers are used. A simple example that uses only a five-element array, sums the contents of the array in `rax`, uses `rsi` for the index, and uses the loop instruction is as follows:

```
mov rax,0
lea rsi,numarray
mov rcx,5
for2: add rax,[rsi]
add rsi,8
loop for2
endfor1: nop
```

As before, the `rsi` register needs to be initialized with the address of the array. Notice that it is similar to an array of 64-bit `real8` numbers. Since eight-byte quad words are being used instead of four-byte double words, `rsi` is incremented by 8 instead of 4.

Likewise, an example using the `rbx` register in conjunction with the name of the array is as shown in the code segment below:

```
mov rax,0
mov rbx,0
mov rcx,4
for2: nop
```

```
    add rax,numarray[rbx]
    add rbx,8
    loop for2
endfor2: nop
```

Again, `ebx` should be incremented by 8 instead of 4. Both ways of addressing an array are similar to the 32-bit methods and either can be chosen based upon the task at hand or whichever is more convenient.

1.70 Procedures and Macros

As initially discussed in Chap. 7, procedures and macros are extremely helpful for taking common code and not having to rewrite it over and over again. It is suggested to review that chapter if necessary before proceeding with this section.

1.70.1 Calling 64-Bit Procedures

Having called the `printf` and `scanf` procedures, much has already been learned on how to use procedures. Simple procedures can be called using memory locations and various registers as done with 32-bit procedures, but then their use would be restricted to only those memory locations and registers. Instead, the same conventions that are used to call procedures like `printf` and `scanf` can be used which would allow for consistency in the passing of information via `rcx`, `rdx`, `r8`, and `r9` as necessary and a value could be returned via `rax`. Recall that there is a complete program in Sect. 7.2 to implement the power function. A limitation to this program is that when restricted to just the 32 bits of the `eax` register, the largest signed number that could be calculated would be slightly over two billion and the largest non-signed number is slightly over four billion. When using 64-bits the largest signed number that can get calculated, as mentioned at the beginning of this chapter, is in excess of nine quintillion and the largest unsigned number is over eighteen quintillion.

Without rewriting the entire program, which is left as an exercise at the end of this chapter, the following only implements the calculation portion using only positive integers to illustrate the sending and returning of values. First, the calling program segment is shown below:

```
mov rcx,x
mov rdx,n
call power
mov answer, rax
```

Recall that values are passed in the registers `rcx`, `rdx`, `r8`, and `r9`, and that a value is returned in `rax`, so that `x` and `n` are passed via `rcx` and `rdx`, and `answer` can be returned via the `rax` register. The procedure is implemented as follows:

```

power proc
    push rbx
    push rdx
    push r8
    mov r8,rdx
    mov rbx,1
    mov rax,1
    ; while rbx <= r8
while01: cmp rbx,r8
    jg endw01
    imul rcx
    inc rbx
    jmp while01
endw01: nop
    pop r8
    pop rdx
    pop rbx
    ret
power endp

```

Notice that the value in *rdx* is copied to the *r8* register in the procedure; why is this? The reason is that the product from the *imul* instruction is stored in the *rdx* : *rax* register pair, so that the exponent *n* passed in *rdx* would get destroyed the first time through the loop. As a result, the value is moved to *r8* for comparison in the loop. Since *rbx* and *r8* are destroyed in the implementation of the loop, and *rdx* is destroyed via the multiplication, they are pushed and popped off the stack.

1.70.2 Using a Macro to Call printf

The problem when using input and output is that there are a number of instructions that need to be written before and after the calling of *printf* and *scanf*. Further, if one is not careful, there could be cases where information that was left unintentionally in registers could be lost. Is there a way that this can be simplified? One way is to write a procedure that could contain all the necessary code. But the problem would be that a procedure would be written that would call another procedure and this could get a little confusing.

As an alternative, a macro could be written to contain all the required information that in turn would call the appropriate input or output procedure. As should be recalled from Chap. 7, the disadvantage of this would be that each time the macro is invoked a copy of all the code in the macro is inserted into the program. However, this would be no more code than what would have to have been written in the first place. As discussed previously, there are a number of alternatives than can be used when writing macros. To keep it somewhat simple here only the common case where a string is followed by an integer will be considered. For example, the following typical code segment,

```

lea rcx, msg1fmt
lea rdx, msg1
mov r8, num64
sub rsp, 40
CALL printf
add rsp, 40

```

could be placed into a macro as follows,

```

output64 macro
    lea rcx, msg1fmt
    lea rdx, msg1
    mov r8, num64
    sub rsp, 40
    CALL printf
    add rsp, 40
endm

```

and could be invoked by just typing:

```
output64
```

Although this cleans up the main program considerably, it only allows the user to output the contents of `num64` with `msg1` according to the format in `msg1fmt`. To allow more versatility, parameters should be used. To again keep things simple and to make sure that exactly three parameters are used, the `:REQ` directive is used in the following macro definition:

```

output64 macro format:REQ, message:REQ, number:REQ
    lea rcx, format
    lea rdx, message
    mov r8, number
    sub rsp, 40
    CALL printf
    add rsp, 40
endm

```

To invoke it would be as follows:

```
output64 msg1fmt,msg1,num64
```

Of course, when using parameters as above, other formats, messages, and integers could be used as arguments. Though this works well, sometimes users of procedures such as `printf` forget that registers can be altered. Further, macros tend to hide any other registers that are being altered. Although programmers should be extra careful when using procedures and macros to ensure that the data in registers is saved, it is sometimes helpful to have the macro save the

registers. Though this takes up even more memory space each time they are invoked, the subsequent time saved in debugging might be worthwhile especially in the early stages of writing a program. Then later the saving and restoring of registers in the macro could be removed and the code optimized.

Of the 12 general purpose registers *rax* through *rdx* and *r8* through *r15*, seven of them have their contents altered by the macro and subsequent procedure call. Obviously *rcx*, *rdx*, and *r8* are altered by the macro, but also recall that *rax* and *r9* through *r11* are altered by *printf* as well. Only *rbx* and *r12* through *r15* are left untouched. Unfortunately, there is not a 64-bit equivalent of the 32-bit *pushad* and *popad* instructions that saves and restores all the registers on the stack. Instead the seven affected registers need to be saved and restored individually. This can be done either by using the stack, storing them in separate memory locations, or using an array. Using a different name for the macro can be helpful so that the programmer can choose between the one that saves registers and the one that does not as needed. Looking at the stack option, the macro definition for *output64s* where the *s* stands for save is as follows:

output64s macro format:REQ, message:REQ, number:REQ

```
    push rax
    push rcx
    push rdx
    push r8
    push r9
    push r10
push r11
lea rcx,format
lea rdx,message
mov r8, number
sub rsp, 40
CALL printf
add rsp, 40
pop r11
pop r10
pop r9
pop r8
pop rdx
pop rcx
pop rax
endm
```

Yes, the appearance of seven push and seven pop instructions does look rather unwieldy and takes up a number of memory locations. Again, if one is careful to always return data from registers to memory locations, then using the non-saving macro would not cause a problem. If one does not want to use the stack for saving the registers, then as mentioned previously, seven different memory locations each with the name of the saved register as part of their

names could be created to save and restore the seven registers, Also, a seven element array maybe called `saveregs` could be used. Although neither of these two options would save any instructions and would in fact use just a little more memory, they are alternatives to using the stack and are left as exercises at the end of the chapter.

Also, instead of having two separate macros, one macro could be written using conditional assembly as discussed in Sect. 7.5 and this is left as an exercise. Further, a macro could also be written for `scanf` and this is also left as an exercise at the end of the chapter.

1.71 Complete Program: Reversing an Array

To incorporate a number of the items in this chapter, a simplified version of the program at the end of Sect. 8.2 to reverse an array is rewritten using 64-bit processing and shown below:

```
printf PROTO
scanf PROTO
.data
msg1fmt byte "%s",0
msg2fmt byte "OAh,"%s",0Ah,0Ah,0
msg3fmt byte " %lld", 0Ah,0Ah,0
in1fmt byte "%lld",0
msg2 byte "Enter an integer: ",0
msg3 byte "Reversed",0
n sqword 5
array sqword 5 dup(?)
.code
main proc
    mov rcx,n ; initialize rcx to n
    mov rbx,0 ; initialize rbx to 0
for01: nop
    push rcx ; save rcx
    lea rcx, msg1fmt
    lea rdx, msg2
    sub rsp, 40
    CALL printf
    add rsp, 40
    lea rcx, in1fmt
    lea rdx,array[rbx]
    sub rsp,40
    call scanf
    add rsp,40
    pop rcx ; restore rcx
    add rbx,8 ; increment rbx by 8
```

```

    loop for01
endfor01: nop
    lea rcx, msg2fmt
    lea rdx, msg3
    sub rsp, 40
    CALL printf
    add rsp, 40
    mov rcx,n ; initialize rcx to n
    sub rbx,8 ; subtract 8 from rbx
for02: nop
    push rcx ; save rcx
    lea rcx, msg3fmt
    mov rdx, array[rbx]
    sub rsp, 40
    CALL printf

    add rsp, 40
    pop rcx ; restore rcx
    sub rbx,8 ; decrement rbx by 8
    loop for02
endfor02: nop
    ret
main endp
end

```

Notice that the value of `rcx` is pushed onto the stack at the beginning of each loop and popped off the stack at the end of each loop since `rcx` is used to communicate information to the `printf` and `scanf` procedures. Also, the simplified code only inputs a fixed number of items, but to input a variable number of items by inputting and checking that the value in `n` is non-negative is left as an exercise at the end of the chapter.

1.72 Summary

- Remember that moving data into the lower 32 bits of a 64-bit register causes the upper 32 bits to be cleared to zeros.
- The `movsxd` instruction copies the contents of a 32-bit register or memory location into a 64-bit register and then extends or propagates the sign into to the upper 32 bits of the 64-bit register.
- One cannot mix in an instruction using the upper 8 bits of the corresponding 16-bit register, such as `ah` in the `ax` register, with the lower 8 bits of the new numbered 64-bit registers such as `r8b`.
- Use `%llx` format to be able to output all the bits of a 64-bit word in hexadecimal.

- The PROTO statements do not need a parameter list.
- Both the printf and scanf statements use CALL statement instead of INVOKE.
- Use the rcx, rdx, r8 and r9 registers to send information to the printf statement.
- Remember that when invoking a macro all the code in the macro is inserted into the program.
- Remember to load the address of the input format into rcx and the address of the memory location to be input into rdx prior to calling scanf.
- When accessing the elements of an array of 64-bit integers, be sure to increment the register, rsi, rdi, or rbx, by 8 instead of 4.

1.73 Exercises (Items Marked with an * Have Solutions in Appendix D)

1. Given the following code segments, indicate the contents of the 64-bit register (show all 64 bits in hexadecimal).

```
a. mov rax, 0FFFFFFFFFFFFFFFFh
b. mov rbx, 0
   mov eax, 0
   mov ebx, -2
*c. mov rcx, 0011002200330044h
d. mov rdx, 0FFFFFFFF01234567h
   mov cx, -1
   mov edx, 789ABCDEh
```

2. Implement the following C code segments in MASM assembly language as done in Sect. 11.3.

```
a. #include <stdio.h>
int main() { long long num1 = 5, num2 = 7; printf("\n%lld\n", num1);
printf("%s\n", " is not = "); printf("%lld\n", num3); return 0 ; }
b. #include <stdio.h>
int main() {
long long num1 = 1, num2 = 2, num3 = 3, num4 = 4; printf("\n%lld%s%lld",
num1, "< ", num2); printf("%s%lld%s, "< ", num3, "< "); printf("%lld\n",
num4); return 0 ;
}
```

3. Write a macro to implement the code segment at the end of Sect. 11.5.2.
4. Rewrite the program from Sect. 7.2 to incorporate 64-bit processing and parameter passing as described in Sect. 11.8.1.

5. Rewrite the macro definition `output64s` in Sect. 11.8.2 to use each of the following and then invoke them in a program:
 - a. Use memory locations `saverax`, `savercx`, etc., to save and restore the seven affected registers.
 - b. Use a seven-element array called `saveregs` to save and restore the seven affected registers.
6. As an alternative to having two macros to output, one that saves registers and one that does not, write a single macro using conditional assembly as discussed in Sect. 7.5. A fourth parameter could be used and if the parameter is blank it does not save the registers, otherwise it does.
7. Write a macro called `input64` that incorporates a prompt via a parameter, inputs using `scanf`, and returns the value input via the `rax` register.
8. Write a macro using parameters to move data from a 32-bit memory location into the lower 32 bits of another 64-bit memory location without altering the upper 32 bits.
9. Modify the complete program in Sect. 11.9 to input the value of `n` and make sure that it is not a negative number as done in the original program in Sect. 8.2.
10. Modify the complete program in Sect. 11.9 to use `rsi` instead of `ebx`. Start by replacing the second instruction `mov ebx, 0` with `lea rsi, array`.

1.74 Introduction

The purpose of this chapter is to help draw further connections between assembly language and computer organization. The advantage of using assembly language is that one can see computer organization from a software perspective. Further, by examining machine language one can see some of the principles discussed in a computer organization text. It is especially helpful to see these principles implemented in a real machine language. As may have been learned from a computer organization text, there are many considerations that need to be taken into account and many different formats that can be used for machine language instructions. One of the first considerations is the size of the instruction. With a larger the instruction, more opcodes can be included, more registers can be referenced, and more memory locations can be addressed. Also, how a particular instruction is divided up indicates how many of each of the above can be included. For example, assume a 16-bit word is divided up as follows: bits 15-13 for the opcode, 12-11 for referencing registers, and 10-0 for addressing memory locations, as shown in Fig. 12.1.

Since there are 3 bits allocated for the opcode, there are 2^3 or 8 possible opcodes. Given 2 bits for the register, there are 2^2 or 4 possible registers, and given 11 bits allocated for addresses, there are 2^{11} or 2,048 possible addresses for memory locations.

Although there is a very formal procedure and elaborate mechanism for determining the machine code for Intel machine language, one can discover the format of many instructions by examining the machine code of particular instructions. What is interesting about this method is that one can examine the machine language of almost any given processor and see some of the instruction layouts merely by inspection.

1.75 inc and dec Instructions

By turning on the assembly language listing option as discussed in Appendix A.4, the .lst file can be generated and the corresponding assembly language can be seen in the left columns of the listing. The complete program listings used in this chapter are shown in Sects. 12.9 and 12.10. Although the programs do not do anything relevant in terms of an algorithm, their sole purpose is to list out selected machine language instructions for comparison. Of course the machine language is given in hexadecimal, but it can always be converted to binary (see Appendix B) in order to see the corresponding bit patterns. Then the bit patterns can be carefully examined to help one understand the machine language.

For example, in the following section of code taken from the complete listing in Sect. 12.9, there are listed a number of similar instructions. Again, the code listed does not actually do anything useful, because the purpose is to understand the corresponding machine language. In this first example, only the single register inc instructions are shown for the sake of simplicity, where multiple register instructions, immediate data instructions, and instructions that address memory locations will be examined later:

Address	Machine	Assembly	Address	Machine	Assembly
00000000	40	inc eax	00000004	44	inc esp
00000001	41	inc ecx	00000005	45	inc ebp
00000002	42	inc edx	00000006	46	inc esi
00000003	43	inc ebx	00000007	47	inc edi

First note that the above inc instructions are only two hex digits long, or in other words these are 1-byte instructions, where each byte has its own sequential memory address to the left. In looking at the above code, it should be noticed that the last hexadecimal digit changes as the register changes, where the hexadecimal digits 0, 1, 2, 3, 4, 5, 6, and 7 correspond to the registers eax, ecx, edx, ebx, esp, ebp, esi, and edi, respectively. The equivalent numbers in binary would be 0000, 0001, 0010, 0011, 0100, 0101, 0110, and 0111, respectively. Note as discussed above, it took three binary digits (the three low-order bits of the second hex digit) to represent the eight registers, leaving five binary digits (four from the first hex digit and the one high-order bit from the second hex digit) to represent the opcode for inc as 01000. The format for inc 32-bit register instruction is as in Fig. 12.2.

inc reg 01000XXX ,where XXX is the register.

Fig. 12.2 Format for the inc instruction

Expanding on the above, the following is a listing of the dec reg instruction:

Address	Machine	Assembly	Address	Machine	Assembly
00000008	48	dec eax	0000000C	4C	dec esp
00000009	49	dec ecx	0000000D	4D	dec ebp
0000000A	4A	dec edx	0000000E	4E	dec esi
0000000B	4B	dec ebx	0000000F	4F	dec edi

Again, it should be noticed that each instruction is only 1 byte long as before. Yes, there is a hex 4 in the first four bits, but the second hex digit is not as obvious as the previous example. However, looking at the binary for these instructions, it becomes much more obvious:

Hex	Binary	Assembly	Hex	Binary	Assembly
48	01001000	dec eax	4C	01001100	dec esp
49	01001001	dec ecx	4D	01001101	dec ebp
4A	01001010	dec edx	4E	01001110	dec esi
4B	01001011	dec ebx	4F	01001111	dec edi

In examining the binary carefully, it should be noticed that the same bit pattern appears again in the right three bit positions representing the registers eax through edi. The only difference is that bit position 3 (fourth from the right) is a 1 instead of a 0 . The result is that the only difference between the inc and dec instructions is this one bit, where a 0 tells the processor to increment the specified register by 1 , whereas a 1 tells the processor to decrement the specified register by 1 . The machine code for the dec 32-bit register instruction is as in Fig. 12.3.

For convenience, a summary of the machine code representation of the registers can be found below, where the first column shows the machine code and the second column shows the corresponding register:

Machine code	Register
000	eax
001	ecx
010	edx
011	ebx
100	esp
101	ebp
110	esi
111	edi

1.76 mov Instruction

Before looking at how memory is addressed, it is helpful to first examine the mov instructions that are used with registers to notice some similarities and differences between them and the inc reg and dec reg instructions. First, the mov reg, reg instructions will be examined, but instead of listing out all possible register combinations, only some of the possibilities are shown in the interest of saving space. Besides, the patterns should become obvious by examining only a select few of these instructions:

Address	Machine	Assembly
00000010	8B C0	mov eax, eax
00000012	8B C1	mov eax, ecx
00000014	8B C2	mov eax, edx
00000016	8B C3	mov eax, ebx
00000018	8B C4	mov eax, esp
0000001A	8B C5	mov eax, ebp
0000001C	8B C6	mov eax, esi
0000001E	8B C7	mov eax, edi
00000020	8B C8	mov ecx, eax
00000022	8B C9	mov ecx, ecx
00000024	8B CA	mov ecx, edx
00000026	8B CB	mov ecx, ebx

The first thing to notice is that the address of each instruction is incremented by 2 each time and that the machine language is now 4 hex digits long, which is 16 bits or 2 bytes long. The first byte in all the instructions is an 8B which one can probably assume is part of the opcode of the instruction indicating a register-to-register mov instruction. Although there appears to be somewhat of a pattern in the second byte of the first eight instructions, it probably would not hurt to convert these to binary to see the pattern in all of the instructions more clearly as in Table 12.1.

The suspicions of the first set of instructions appear to be correct, where eax through edi are 000 through 111, respectively, just like the inc and dec instructions previously. Then looking at the last four instructions, the same pattern appears to begin again with eax through ebx with 000 through 011, respectively.

In examining bits 3 through 5 (the 4th through 6th bits from the right), it can be noticed that for the first eight instructions, the bits are set at 000 and for the second four instructions they are set at 001, which happens to correspond to the first registers in the operand, eax and ecx, respectively. So part of the reason why this mov reg, reg instruction is larger than the inc reg instruction is that it needs to have room to reference two registers instead of one. The result is that one should be able to surmise that the mov between two registers has the format shown in Fig. 12.4.

Having looked at a few of the single and double register instructions, what

Table 1.9: Binary of second byte of the mov reg, reg instruction

Hex	Binary (2nd byte)	Assembly
8B C0	11000000	mov eax, eax
8B C1	11000001	mov eax, ecx
8B C2	11000010	mov eax, edx
8B C3	11000011	mov eax, ebx
8B C4	11000100	mov eax, esp
8B C5	11000101	mov eax, ebp
8B C6	11000110	mov eax, esi
8B C7	11000111	mov eax, edi
8B C8	11001000	mov ecx, eax
8B C9	11001001	mov ecx, ecx
8B CA	11001010	mov ecx, edx
8B CB	11001011	mov ecx, ebx

does an instruction with immediate data look like? Staying with the mov instruction is a convenient next step:

Address	Machine	Assembly
00000028	B8 00000001	mov eax,1
0000002D	B9 0000000A	mov ecx,10
00000032	BA FFFFFFFF	mov edx,-1
00000037	BB FFFFFFF6	mov ebx,-10

mov reg,reg

10001011 11XXXXYY

where XXX is the destination register

and YYY is the source register

Fig. 12.4 Format for the mov reg, reg instruction

Note that the addresses of each instruction is incremented by 5 and that each instruction is 10 hexadecimal digits long, equating to 40 bits, indicating that each instruction is 5 bytes long. It should be noticed that the last four bytes of each instruction is the hexadecimal equivalent of the immediate data in the instruction, where it should be remembered that negative numbers are represented in two's complement notation (see Appendix B). Again it helps to convert the first byte to binary to determine the machine code as follows:

Hex	Binary	Register
B8	10111000	eax
B9	10111001	ecx
BA	10111010	edx
BB	10111011	ebx

As before, the last three binary digits 000, 001, 010, 011 are the same digits as in previous instructions representing the `eax`, `ecx`, `edx`, and `ebx` registers, respectively. The format then is as illustrated in Fig. 12.5.

As can be seen, the immediate data instructions are fairly straightforward, so how would memory addresses be represented in an instruction? Just as the first instruction in the `.code` segment starts at relative address or relative memory location 00000000, the first memory location declared in the `.data` segment will also start at relative memory location 00000000 as shown below:

Address	Machine	Assembly language
00000000	00000003	<code>.data</code>
00000004	00000005	<code>num1 sdword 3</code>
00000008	00000000	<code>num2 sdword 5</code>

A relative memory location is relative to the beginning of either the `.data` or the `.code` segment, because at the time of assembly it is not known where in memory the segments will be loaded. When the machine language is subsequently loaded into memory, the relative addresses are eventually changed into the absolute addresses or absolute memory locations within RAM, where the first memory location in RAM is absolute memory location 00000000. Using the memory locations declared above, consider the following `mov` instructions that use the `eax` register and address memory locations `num1` and `num2`:

```
mov reg,imm \10111XXX
```

where XXX is the register and YYY is the immediate data

Fig. 12.5 Format for the `mov reg, imm` instruction

Address	Machine	Assembly
00000003 C	A1 00000004	<code>mov eax, num2</code>
00000041	A3 00000008	<code>mov num3, eax</code>

As with immediate data, note that these two instructions are both 5 bytes long. Unlike the immediate instructions, where the data in the rightmost four bytes was the actual number contained in the instruction, here it should be noticed that address of the memory locations `num 2` and `num 3` appears in the right four bytes. In other words, instead of the numbers 00000005 or 00000000 appearing in the instruction, the addresses of the memory location `num2` and `num3`, 00000004 and 00000008, appear in the instructions, respectively. In converting the first byte to binary, notice that there is only one bit difference between the instructions:

Hex	Binary	Assembly
A1	10100001	<code>mov eax, mem</code>
A3	10100011	<code>mov mem, eax</code>

When moving from memory to the `eax` register, bit 1 (second from the right) is set to 0, and when moving from the `eax` register to memory, note that bit 1 is set to 1. The format of these instructions is as shown in Fig. 12.6.

But what about moving to and from memory into the other registers? Consider the following instructions using the other three general purposes registers:

Address	Machine	Assembly
00000046	8B 0D 00000004	<code>mov ecx, num2</code>
0000004C	8B 15 00000004	<code>mov edx, num2</code>
00000052	8B 1D 00000004	<code>mov ebx, num2</code>
00000058	89 0D 00000008	<code>mov num3, ecx</code>
0000005E	89 15 00000008	<code>mov num3, edx</code>
00000064	89 1D 00000008	<code>mov num3, ebx</code>

The address part of the instruction in the rightmost four bytes looks okay. Also the `mov reg, mem` instructions all have an 8 B in the first byte and all the `mov mem, reg` instructions have an 89 in the first byte. Again looking at the binary:

Hex	Binary	Assembly
8B	10001011	<code>mov reg, mem</code>
89	10001001	<code>mov mem, reg</code>

Note that some of the leftmost five bits are different from the `eax` instructions in Fig. 12.6, where instead bit 3 (fourth from the right) is a 1 instead of a 0 and bit 5

(sixth from the right) is a 0 instead of a 1. Further, it is opposite from the `mov eax, mem` and `mov mem, eax` instructions, where bit 1 (second from the right) is a 0 when transferring from memory to a register and a 1 when transferring from a register to memory.

The most noticeable difference is that these instructions are 1 byte longer than their `eax` counterparts, 6 bytes instead of 5. Recall from Chap. 1 when discussing registers, it was said that the `eax` instruction is usually the preferred register, because it tends to be shorter. Further, in many of the code examples in this text, the `eax` register tends to be used more often than any other register. As can be seen in the code segment above, the use of the other registers is indeed a little less efficient in terms of memory utilization.

The first byte above is indeed the opcode, but look carefully at the binary of the second byte below. Clearly the first two bits on the left are the same. Also, the last three bits are the same and do not contain the register bits as they have in the past. However, in examining the middle bits, specifically bits 3 through 5 (fourth through sixth from the right), they are different with each instruction. As with previous instructions, 001, 010, and 011 correspond to the registers `ecx`, `edx`, and `ebx`, respectively:

Hex	Binary	Register
0D	00001101	ecx
15	00010101	edx
1D	00011101	ebx

Since this instruction format is a little more complicated than all the other ones examined thus far, in order to analyze the bits of this second byte, it is helpful to break it down and look at the machine language format of this byte. Although there are many possible bit combinations, this text only looks at a few of them. The basic format of this byte is shown in Fig. 12.7.

As seen previously, the bits in the middle (bit positions 3 through 5) indicate the register and is abbreviated as *reg*. The two bits on the left (bit positions 6 and 7) indicate the mode of the instruction and is abbreviated as *mod*. The three bits on the right (bit positions 0 through 2) indicate the register and mode, abbreviated as *r/m*. With a 00 in the *mod* field and a 101 in the *r/m* field, this means that it is displacement only addressing mode or in other words direct addressing mode.

Again, the meaning of the contents of this byte can be very complicated and will vary with other instructions that use different addressing schemes, but only a few of the simpler instructions are discussed here to serve as an introduction.

1.77 add and sub Instructions

Having looked at some of the *inc*, *dec*, and *mov* instructions, what about some of the arithmetic instructions? Instead of looking at all the registers as done in previous examples, only a few select registers will be examined, since the bit patterns of the registers have already been established. Consider the following register-to-register add and sub instructions:

Address	Machine	Assembly	Binary equivalent
0000006A	03C0	add eax, eax	00000001111000000
0000006C	03C1	add eax, ecx	00000001111000001
0000006E	03C2	add eax, edx	00000001111000010
00000070	03C3	add eax, ebx	00000001111000011
00000072	2BC8	sub ecx, eax	0010101111001000
00000074	2BC9	sub ecx, ecx	0010101111001001
00000076	2BCA	sub ecx, edx	0010101111001010
00000078	3BCB	sub ecx, ebx	0010101111001011

Note that instead of separating the binary equivalent of the opcodes on separate lines, it is now included off to the right for convenience. The first thing to notice is that the opcode for the add instruction and sub instruction differs by only 2 bits, where bit position 3 and 5 of the first byte (fourth and sixth from the right) are each a 1 instead of a 0 for the subtract instruction. Also notice that the rightmost six bits of the second byte appear to represent the registers, just like

the register-to-register mov instruction earlier. However, given the previous information on mod, reg, and r/m sections in the mov instruction, it sheds some further light on the format of this and the previous mov reg, reg instruction. The reg section is indeed the first register that appears in the operand. However, the 11 in the leftmost two bits of the second byte is actually part of the mod section and indicates that the r/m section holds the code for the second register in the operand. This is true for all three of the reg, reg instructions examined in this chapter: the mov reg, reg, add reg, reg, and sub reg, reg instructions. Given the reg mod r/m byte, there can be up to 256 bit combinations which allow for a large number of addressing modes. This is one of the reasons the Intel processor is known as a complex instruction set computer (CISC) as opposed to a computer with fewer instructions and addressing modes known as a reduced instruction set computer (RISC).

Looking at add and sub instructions referencing memory locations below, what similarities and differences can be found between them and various previous instructions?

First, one should notice that the first byte is the same as the previous add and sub register-to-register instructions. Next, the last four bytes (32 bits) are the relative addresses of the memory location num2, just like the previous mov reg,mem instructions. Lastly, the mod reg r/m byte is different than the add reg,reg and sub reg,reg instructions, but they are the same as the previous mov reg,mem instructions because they are addressing memory in the same fashion.

Address	Machine	Assembly	Binary equivalent
0000007A	03 05 00000004	add eax, num2	0000001100000101
00000080	03 0D 00000004	add ecx, num2	0000001100001101
00000086	03 15 00000004	add edx, num2	0000001100010101
0000008C	03 1D 00000004	add ebx, num2	0000001100011101
00000092	2B 05 00000004	sub eax, num2	0010101100000101
00000098	2B 0D 00000004	sub ecx, num2	0010101100001101
0000009E	2B 15 00000004	sub edx, num2	0010101100010101
000000A4	2B 1D 00000004	sub ebx, num2	0010101100011101

1.78 mov offset and lea Instructions

Having looked at the moving of the contents of a memory location to a register, what is the difference between that type of instruction and a lea instruction or mov offset instruction. First consider the following mov and lea instructions:

Address	Machine	Assembly	Binary equivalent
000000AA	8B 35 00000004	mov esi, num2	1000101100110101
000000B0	8B 3D 00000004	mov edi, num2	1000101100111101
000000B6	8D 35 00000004	lea esi, num2	1000110100110101
000000BC	8D 3D 00000004	lea edi, num2	1000110100111101

The first two instructions are the same as previous mov instructions, where in this case they use the esi and edi registers and the middle bits of the mod reg r/m byte are 110 and 111, respectively. The two lea instructions have the same format in the second byte as the mov instructions. The only difference is in the first byte, where bit positions 1 and 2 (second and third from the right) are opposite of each other. As before, the result is that when the processor encounters the mov instructions, the ``contents'' of memory location num2 are loaded into the specified register and with the lea instructions the ``address'' of memory location num2 is loaded into the specified register. In this instance in the former case, the contents 5 is loaded and in the latter case the address 4 is loaded. Remember from Chap. 8 that the mov offset and lea instructions effectively perform the same task, where they both load the address of the memory location into the specified register. The only difference is that the mov offset instruction is static and the lea instruction is dynamic, where in the former the address is determined at assembly time, and in the latter the address can be indexed by a register and is determined at run-time. In looking at the machine language for the mov offset instruction below, the address is in the last four bytes of the instruction as with many other instructions before. Further, its machine language equivalent is 1 byte shorter than its lea counterpart, where instead of the esi and edi registers being included in a mod reg r/m byte, they are in the rightmost three bits (bit positions 0 through 2) of the opcode byte as 110 and 111, respectively:

Address	Machine	Assembly	Binary equivalent
000000C2	BE 00000004	mov esi, offset num2	10111110
000000C7	BF 00000004	mov edi, offset num2	10111111

The significant difference, as discussed in Chap. 8, is that instead of the ``contents'' of the memory location being moved to the register, the ``address'' of the memory location is moved to the register. Although probably not readily apparent, the machine language for the opcode of the mov offset instruction should be familiar. In fact it is the same, 10111, as the mov reg, imm instruction discussed in Sect. 12.3. Compare the following instructions with the instructions above:

Address	Machine	Assembly	Binary equivalent
000000CC	BE 00000004	mov esi,4	10111110
000000D1	BF 00000004	mov edi,4	10111111

Note that the machine language of the mov reg, imm instructions looks the same as that of the mov reg, offset mem instructions. Does this mean that one should use the immediate instructions to try to load the address of a memory location? The answer is no, because if the contents of the register in both cases were to be output, there would be a difference between the two. Recall from Sect. 12.3 the discussion of relative and absolute addresses. Although the value in the immediate case is 4, the value in the offset case is not the relative address of num2 which is 4, but rather the absolute address of the memory location

num2 which can vary depending on where the program is loaded in memory. For example, when running a test program on one computer, the value in the register was 00404008 (in hex) and when running the same program on another computer, the value in the register was 010F4008 (in hex). Even if one could use the correct address of a particular variable, is this method of addressing memory a good idea? As might be suspected, the answer is no, because the location of the variable could change and cause potential logic or execution errors. Using immediate values to access memory instead of variable names or pointers essentially reduces an assembly language instruction back to that of a machine language instruction, thus eliminating the purpose of using an assembler in the first place.

1.79 jmp Instructions

As discussed in Chap. 4, the jmp instruction is known as an unconditional jump because it jumps regardless of the setting of the eflags register. Although the code is somewhat nonsensical, the machine language of the jmp instruction can be found in the following sample code segment:

Address	Machine	Assembly language instruction
000000D6	EB 04	jmp around
000000D8	90	above: nop
000000D9	90	nop
000000DA	EB FC	jmp above
000000DC	90	around: nop

The opcode for the jmp instruction is EB, but what is interesting to note is that the machine code for the jmp around instruction does not contain the address of the around: nop instruction. Why is this and is it referring to memory location 04? The answer to the latter part of this question is no, because the jmp instruction is referring to a location that is relative to itself. However, if a 04 is added to the address of the memory location where the jmp around instruction is located, 000000D6, the new address is 000000DA. But looking at the above code segment, 000000DA is the address of the jmp above instruction and not the around: nop instruction. If the 04 is the address relative to jmp around instruction, why does the number seem incorrect? The answer is that the number is not incorrect because the number is relative to where the instruction pointer is pointing. To those who have had or are currently taking a computer organization course, the following explanation should be somewhat familiar. As a program executes, the instruction pointer or instruction counter in the CPU points to the instruction that it is about to fetch and subsequently execute. After an instruction is fetched, the instruction pointer is then incremented to point to the next instruction that follows the current instruction. Then when the current instruction is decoded and executed, the instruction pointer is no longer pointing at the current instruction but rather to the next instruction in anticipation of fetching it. In the above case, the instruction pointer is not pointing to

memory location 000000D6, but rather 000000D8, so that when 04 is added to 000000D8, the answer is 000000DC which is the address of the correct instruction. The value 000000DC is then placed into the instruction pointer so that when the next instruction is fetched, it is not to the above: nop instruction but rather the around: nop instruction.

However, what about the jmp above instruction? Shouldn't it be a jump backward and not a jump forward? The answer is yes, where the FC is not a positive number but rather a negative number (11111100 in binary) and it is in two's complement form. Using the techniques learned in Appendix B, 11111100 in binary is equal to -4 in decimal. Remembering that the instruction pointer is not pointing to the current instruction, jmp above, but rather the instruction after it at memory location 000000DC, subtracting a 04 results in the address 000000D8, which is the correct location and instruction.

Notice that only 1 byte is used to store the relative offset address and this allows for jumps of only +127 bytes or -128 bytes away. On older 16-bit machines, this sometimes posed a problem and on occasion a code segment needed to be rewritten to accommodate jumps in excess of the above limitations. Fortunately, with the newer 32-bit processors, this does not pose a problem because should the unconditional jump be more than the above limits, a 32-bit relative offset is generated.

1.80 Instruction Timings

As has been mentioned at various points throughout this text, some instructions are faster than others. It has been implied that the faster instructions should be used over their slower counterparts when necessary, and also when it does not interfere with the readability and maintainability of the program. Why is it that some instructions are faster than others? The reasons for this can have to do with many different factors including the size of the instruction, whether the instruction references registers as opposed to memory locations, the complexity of the operation, and how the processor is designed. When a processor is said to be rated at a speed 2.0 GHz, that means that the processor can execute at two billion cycles per second, where different instructions take different number of cycles to execute. For the discussion that follows, the instruction timings are based on a 32-bit Pentium processor. Generally, instructions that do not reference memory tend to be faster than instructions that do reference memory. For example, an inc eax instruction takes only one clock cycle of the CPU's time, whereas an inc instruction that references memory such as inc number takes three clock cycles. The reason why registers are faster is that registers are internal to the CPU, whereas memory is external to the CPU and it takes time for data to be transferred from RAM to the CPU or from the CPU to RAM.

Does this mean that memory should be avoided? No, because there are only four general purpose registers in the 32-bit Intel processor and it would be impractical to restrict a program to only four registers. What it does mean is that when in a critical section of the program where speed is important, a value might

be left in a register in lieu of returning the value back to memory so that readability is sacrificed for efficiency.

In looking at a few other instructions, take the `imul` instruction for example. To multiply the contents of `eax` by the contents of another register such as the `ecx` register would take 10 clock cycles. Similarly, the `idiv ecx` instruction takes 46 clock cycles. However, if in each case the contents of the `ecx` register are a power of 2, an arithmetic shift instruction could be used instead. For example, if the contents of the `ecx` register in an `imul ecx` instruction were a 4, then a `sal eax, 2` instruction could be used instead, which would take only three clock cycles which is over three times faster. Likewise, if the contents of `ecx` in an `idiv ecx` instruction were a 32, a `sar eax, 5` instruction could be used which also takes only three cycles and would be 15 times faster. This is the reason why some programmers will use arithmetic shift instructions instead of multiplication and division instructions. Does this mean that arithmetic shifts should always be used over their arithmetic counterparts? No, not necessarily. If the code appears in a non-critical section of a program that is executed only once, then the increase in speed may not be worth the decrease in readability. However, if the arithmetic operation occurs in a time-sensitive section of a program, then the increase in execution speed far outweighs any loss of readability or maintainability.

1.81 Floating-Point and 64-Bit Instructions

Having examined 32-bit integer instructions, it is interesting to look at some of the similarities and differences in comparison with floating-point and 64-bit integer instructions. Because the instructions sets are so large, only a few of these latter instructions will be examined in this section.

1.81.1 Floating-Point Instructions

Just as there were differences in the machine language between various 32-bit integer instructions, there are also differences in the machine language of floating-point instructions. For example, assuming that the variables x , y , and z are declared as `real4`, consider the `fld`, `f st`, and `fstp` instructions as shown below:

Machine language	Assembly language	Binary of the second byte from the left
D9 05 00000000	<code>fld x</code>	000 00 101
D9 15 00000004	<code>f st y</code>	000 10 101
D9 1D 00000008	<code>fstp z</code>	000 11 101

Note that the instructions are six bytes long, where the last four bytes is the operand that contains address of the memory location. In the second byte from the left there are some differences and the binary equivalent is shown on the last column on the right. It is separated by spaces to help indicate where the